

Adequacy Score Between Declared and Real Urgency in a Digital Smart Green Supply Chain

R&D BST, Digital Smart Green Supply Chain: Planning Score



Author

John Hoarau, R&D Intern, Direction Innovation Numérique, ALLEN

Industrial Supervisors

Sébastien Perthuisot Directeur de Programme de Recherche et Expert
Qualité & Supply Chain, DIN, Groupe ALLEN,
s.perthuisot@alten.fr

Fabien Jules Chargé Innovation, ALLEN,
stage.innovation.center@alten.fr

Academic Supervisor (ESTIA)

David Gomez Professeur en Informatique, ESTIA Institut de
Technologie, Bidart, d.gomez@estia.fr

Academic Supervisor (Cranfield)

Peter Sherar Dr, Centre for Computational Engineering Sciences,
Cranfield University, p.sherar@cranfield.ac.uk

ALLEN
Innovation Department

2026
Company Confidential

Remerciements

Je tiens à remercier l'ensemble des personnes qui ont accompagné le déroulement de cette opération de R&D au sein de la Direction Innovation Numérique d'ALTEN.

Je remercie tout particulièrement **Émilienne Lardy**, **Selma Ferhat** et **Sébastien Perthuisot**.

Émilienne et Selma ont accompli un travail remarquable dans mon encadrement et mon suivi scientifique : leur exigence, leur disponibilité et la qualité de leurs relectures ont directement fait progresser la rigueur de ce mémoire.

Sébastien, Directeur de Programme de Recherche et Expert Qualité & Supply Chain à la DIN, m'a apporté son encadrement industriel et sa confiance sur l'orientation scientifique du sujet.

Je remercie mes encadrants académiques, **David Gomez**, Professeur en Informatique à l'ESTIA, et **Peter Sherar**, du Centre for Computational Engineering Sciences de Cranfield University, pour leurs conseils méthodologiques et le cadrage académique apporté à ce mémoire.

Enfin, je remercie l'ensemble de l'équipe de la Direction Innovation Numérique, ainsi que les autres stagiaires du département, pour leur accueil et les échanges qui ont nourri ce travail tout au long du stage.

Contents

Remerciements	1
List of Figures	5
List of Figures	6
List of Tables	7
List of Tables	7
Abstract	8
Résumé	9
1 R&D Operation Context	10
1.1 Internal Research Context	10
1.1.1 Smart Green Supply Chain Programme	10
1.1.2 Project: DISCO Digital Twin	10
1.1.3 Capitalisation of Prior ALTEN-DIN Research	12
1.2 Document Structure	13
2 Scientific Problem and Obstacles to Overcome	14
2.1 Relevant Context and Motivating Observations	14
2.2 Formal Problem Statement	15
2.3 Operation Objectives and Research Questions	16
2.4 Research Hypotheses	17
2.5 Technical and Scientific Challenges	19
3 State of the Art	22
3.1 Introduction: Key Definitions	22
3.2 Literature Review	23
3.2.1 Stream 1: Risk Propagation and Deadline-Sensitive Value	25
3.2.2 Stream 2: Cognitive Bias in Urgency Declaration	27
3.2.3 Closest Antecedents: Digital Twins of Network Risk and MCDM Supplier Scoring	29

3.2.4	Enabling Method: Bayesian and Probabilistic Risk Updating	29
3.2.5	Enabling Method: Digital Twin and Real-Time Decision Support	30
3.2.6	Enabling Methods: Multi-Criteria Decision Methods . . .	32
3.2.7	Analogical and Contextual Literature	33
3.2.8	Validation Methodology: Serious Games	35
3.3	Summary: Limitations of Existing Solutions	36
3.4	Research Axis and Proposed Adequacy-Propagation Structure . .	36
4	Conducted Operations	38
4.1	Methodology: The Agile Test-and-Learn Program	38
4.2	Mathematical Model	39
4.2.1	Core Model: Real and Declared Local Urgency, and Network Propagation	39
4.2.2	Core Model: Asymmetric Cognitive Adequacy (Prospect Theory)	46
4.2.3	Monte Carlo Estimate of the Time Block	47
4.2.4	Network Criticality Index	48
4.3	Application and Implementation Choices	51
4.3.1	System Architecture	51
4.3.2	Expert Criteria Weighting: The AHP Elicitation Model . . .	52
4.3.3	Stochastic Judgment Modeling: Fuzzy BWM and PROMETHEE II Network Outranking	55
4.3.4	Weekly Review Cycle and Calibrated Event Engine	56
4.3.5	Explainability, Export, and Operational Hardening	60
4.4	Use Case and Validation: The HÉLIOS Serious-Game Campaign . .	62
5	Conclusions	75
5.1	Limitations	75
5.2	Next Steps	77
5.3	Bilan Personnel et Compétences Mobilisées	77
5.3.1	Compétences techniques mobilisées	78
5.3.2	Apprentissages et difficultés rencontrées	78

List of Figures

1	The DISCO serious game: 22 suppliers, transport decisions, and the three scored dimensions.	11
2	Conceptual view of the declared vs. real urgency gap and its two pathologies. The step line is declared urgency U_d and the smooth curve is real urgency $U_r(t)$; the shaded areas mark where the two mismatch, false urgency F where U_d exceeds U_r , hidden risk H where U_r exceeds U_d	15
3	The three-layer urgency adequacy architecture investigated in this work (Challenge 5).	21
4	Map of the seven literature streams and the gap they leave open.	24
5	The agile Test-and-Learn iteration loop of the R&D operation.	39
6	Causal structure behind the six-block selection: four direct rupture mechanisms and two arbitration variables.	44
7	Bidirectional urgency propagation on the supply chain DAG.	46
8	One iteration of the systematic shock simulation behind the criticality index.	49
9	Network criticality ranking after a systematic shock simulation.	50
10	Doxygen-generated UML view of the SupplyScore architecture: single facade, gated writes, multi-tenant persistence.	52
11	Correspondence between the four AHP criteria and their computed counterparts.	53
12	AHP questionnaire page with the consistency-ratio validation gate.	54
13	The four-section weekly review cycle committed as a single signed transaction.	57
14	Weekly review page: the four sequential sections (AHP, KPIs, milestones, events).	58
15	From a declared event to a calibrated KPI update: the three calibration operators.	59
16	Criteria weighting interface (Fuzzy-BWM) feeding the explainability decomposition.	61
17	Per-node explainability: exact additive decomposition of U_d and U_r	61
18	Administration page: database sizes, migration version, and backup controls.	62
19	The HÉLIOS network with its per-arc coefficients: γ (downstream demand attenuation) and β (upstream risk amplification), and each node's tier rank. NovaFab (rank 4) is the expected bottleneck; the dashed Meridian arc is the inert backup arc of Section 4.2.	64

20	The eight-node HÉLIOS campaign network. The person icon on each node marks the consultant who plays it, feeding U_d through the weekly AHP questionnaire; the lower band shows the public data series that independently feed U_r	67
21	Real urgency U_r (solid) against synthetic declared urgency U_d (dashed) for all eight nodes across the dry-run v5 replay (T0 to T18, shaded bands mark phases P1, P3, P5). The synthetic U_d follows U_r at every node by construction, which is why H1, H2, H5, and H6 are not interpretable on this run.	70
22	Network-average adequacy score A across the dry-run v5 replay, annotated with the scenario's calibrated events. The score falls sharply at fast turning points (demand spike, Texas freeze, fire and Suez, drought), then rises through the squeeze (P3) as the saturated network is uniformly tracked by the synthetic U_d , before eroding again as the bullwhip and congestion events land.	71

List of Tables

1	Component definitions for the CO ₂ scoring equation.	11
2	DISCO technology stack and key parameters.	12
3	Core variables and domains.	15
4	Component definitions for the urgency discrepancy equation. . .	16
5	Urgency pathologies and potential consequences.	16
6	Literature stream synthesis: contribution and gap relative to this thesis.	25
7	Comparative urgency aggregation and growth models and their properties.	26
8	Temporal Motivational Theory (TMT) parameters and domains. . .	28
9	Heuristic mapping of medical triage concepts to DISCO logistics prioritisation.	34
10	The six real urgency blocks: formula and explanation.	41
11	Selection filter applied to each candidate KPI block.	43
12	Component definitions for the network propagation equations. . .	45
13	Component and parameter definitions for the asymmetric adequacy equations.	47
14	Component definitions for the AHP consistency ratio equation. . .	54
15	Component definitions for the Fuzzy BWM defuzzification.	55
16	Component definitions for the PROMETHEE II net flow.	56
17	The fourteen calibrated event types.	58
18	Component definitions for the event calibration operators.	59
19	The eight HÉLIOS campaign nodes and their real-world anchors. . .	63
20	The twelve calibrated engine events that drive the HÉLIOS crisis, with their real-world anchors.	65
21	Dry-run v5 verdicts on synthetic data: an engineering check, not a hypothesis test.	69
22	Propagated real urgency U_r by node across the dry-run v5 replay (selected turns).	71
23	Calibration confusion matrix on the dry-run v5 synthetic replay (152 node-turn observations).	71
24	LLM-pilot verdicts: eight isolated language-model agents replaying the real AHP questionnaire through the production pipeline (exploratory, not pre-registered).	72
25	Final per-node snapshot of the LLM pilot's SupplyScore project, read directly from the application after turn 18.	73
26	Documented limitations and their operational implications.	76

Abstract

Modern supply chains operating in Physical Internet networks face a fundamental signal-quality gap: the urgency declared by operators often diverges from the physically computable real urgency. This leads to over-triage (resource waste) and under-triage (hidden rupture risk). In Supply Chain 4.0 contexts marked by volatility and multi-tier dependencies, no existing digital twin architecture fully corrects this cognitive bias while propagating it consistently across a supplier network.

This report summarises a R&D internship that proposes a cognitive adequacy score $A(t) \in [0, 100]$ to measure and correct the gap between declared operator urgency and a multi-block, DAG-propagated digital twin real urgency $U_r(t)$. The approach combines pairwise elicitation (AHP, Fuzzy Best-Worst Method) for declared urgency, a noisy-OR aggregation of six operational KPI blocks for real urgency, PROMETHEE II outranking and systematic shock simulation for network-wide criticality, and a calibrated event engine for keeping the underlying KPI state current. The framework is implemented in Python and operated as a persistent, multi-tenant digital twin exercised through a weekly-cadence serious game, with a dedicated calibration service comparing the adequacy signal against recorded operational outcomes.

Résumé

Les chaînes logistiques modernes opérant dans les réseaux de l'Internet Physique font face à un écart fondamental de qualité de signal : l'urgence déclarée par les opérateurs diverge souvent de l'urgence réelle physiquement calculable. Il en résulte du sur-triage (gaspillage de ressources) et du sous-triage (risque de rupture caché). Dans des contextes Supply Chain 4.0 marqués par la volatilité et les dépendances multi-rangs, aucune architecture de jumeau numérique existante ne corrige pleinement ce biais cognitif tout en le propageant de manière cohérente sur un réseau de fournisseurs.

Ce mémoire synthétise un stage de R&D qui propose un score d'adéquation cognitive $A(t) \in [0, 100]$ pour mesurer et corriger l'écart entre l'urgence déclarée par l'opérateur et une urgence réelle $U_r(t)$ multi-blocs, calculée par le jumeau numérique et propagée sur le graphe orienté acyclique du réseau. L'approche combine une élicitation par comparaisons par paires (AHP, Fuzzy Best-Worst Method) pour l'urgence déclarée, une agrégation noisy-OR de six blocs de KPI opérationnels pour l'urgence réelle, le surclassement PROMETHEE II et une simulation systématique de chocs pour la criticité réseau, ainsi qu'un moteur d'événements calibrés pour maintenir à jour l'état des KPI sous-jacents. Le cadre est implémenté en Python et opéré comme un jumeau numérique persistant et multi-tenant, exercé à travers un serious game à cadence hebdomadaire, avec un service de calibration dédié comparant le signal d'adéquation aux résultats opérationnels enregistrés.

1 R&D Operation Context

1.1 Internal Research Context

1.1.1 Smart Green Supply Chain Programme

The Direction Innovation Numérique (DIN) of ALTEN conducts applied R&D at the intersection of digital twin technology, logistics optimisation, and sustainable transport.

The Smart Green Supply Chain programme (Ben Zakour, 2024) articulates three complementary pillars:

- **Lean:** Minimisation of waste (cost, delay, over-stock) through pull-flow logic and real-time scheduling.
- **Green:** Minimisation of carbon footprint per transport leg, aligned with the Global Logistics Emissions Council (GLEC) Framework v3.2 (Smart Freight Centre, 2023) and ISO 14083:2023 (International Organization for Standardization, 2023).
- **Predictive:** Anticipation of disruptions through probabilistic modelling and machine learning, consistent with the digital supply chain twin approach (Ivanov; Dolgui, 2020).

This programme places ALTEN-DIN within the broader Industry 4.0 movement (Kagermann et al., 2013), which uses cyber-physical systems, Internet of Things (IoT), and digital twins to transform reactive supply chains into adaptive, self-organising networks.

The Physical Internet framework (Montreuil, 2011), proposing open, mutualized logistics networks analogous to internet routing protocols, provides the theoretical backdrop for the supply chain simulation environment described below.

1.1.2 Project: DISCO Digital Twin

DISCO (Supply Chain Illustrator) is a serious game and digital twin developed by the DIN to simulate a complete supply chain for a fictitious satellite product. It involves 22 component suppliers in domains such as structural assemblies, optical systems, and embedded electronics (Baxas, 2025). Players assume the role of supply chain buyers and make decisions about supplier selection, transport routes, and factory allocation. The DISCO technology stack is summarised in Table 2.

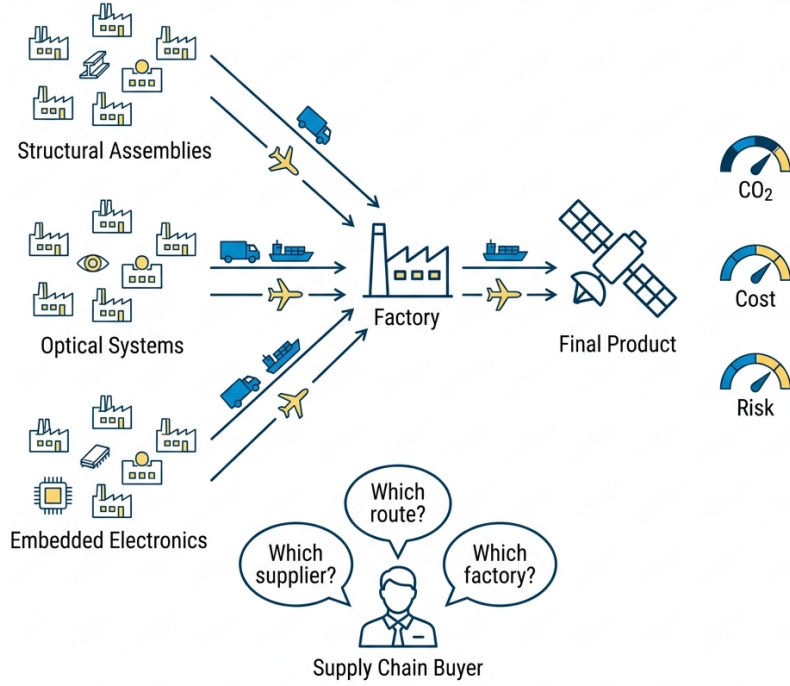


Figure 1: The DISCO serious game: 22 suppliers, transport decisions, and the three scored dimensions.

The scoring engine applies Min-Max normalisation over the values of the current game session:

$$\text{ScoreCO}_2 = 1 + 9 \times \frac{\text{realCO}_2 - \text{co}_{2,\min}}{\text{co}_{2,\max} - \text{co}_{2,\min}}, \quad (1)$$

The operational variables of the CO₂ scoring model are described in Table 1.

Table 1: Component definitions for the CO₂ scoring equation.

Component	Type	Description
ScoreCO ₂	Dependent variable	Normalized greenhouse gas score on a [1, 10] scale
realCO ₂	Input variable	Measured carbon emissions on the transport leg
co _{2,min}	System parameter	Minimum historical or target emissions threshold
co _{2,max}	System parameter	Maximum historical or target emissions threshold

with analogous formulae for cost and risk. The scale ranges from 1 to 10, where 1 corresponds to the best possible performance and 10 to the worst.

To support the demonstration and continuous improvement of mature supply chain capabilities, the DISCO digital twin incorporates several specialized modules developed in prior research cycles: The routing module (Baulain et al., 2024), the risk scorer (Kabouche, 2025), the CO₂ quantification methodology (Cordova, 2024), and the software integration framework (Auclin, 2025).

Table 2: DISCO technology stack and key parameters.

Component	Technology / Value
Backend API	Spring Boot + Java 21
Database	PostgreSQL (Azure)
Frontend	Angular (localhost:4200)
REST API (test)	https://test-disco.altenlabs.net/back/
Authentication	Player key / Azure AD SSO OAuth2
Scores computed	CO ₂ , Cost, Risk per supplier/itinerary [1–10]

1.1.3 Capitalisation of Prior ALTEN-DIN Research

This R&D operation is situated within a body of prior DIN research conducted under the Smart Green Digital Supply Chain programme. These works form the organisational and thematic context of the present operation; the system delivered here (SupplyScore, Section 4) is a standalone Python application and does not integrate with these modules yet. Their coupling is a stated perspective of the programme, not a delivered capability.

Baxas (Baxas, 2025) defines the DISCO software architecture, including the Time To Service (TTS)/Time To Recovery (TTR) formalisation.

Baulain and Dubarry (Baulain et al., 2024) developed the Physical Internet routing module of DISCO, providing transit times via the HERE Maps API (Vincenty distances, mandatory 4h30 Heavy Goods Vehicle (HGV) pause compliance) and the Itinerary Performance Index (Itinerary Performance Index (IPG)) normalised in [0, 1].

Kabouche (Kabouche, 2025) developed the supplier risk scorer combining delay history, geographic vulnerability, and financial fragility into a [1, 10] scale.

Cordova (Cordova, 2024) developed the CO₂ quantification module per route using the GLEC/VECTO methodology with log-normalised output $\hat{x}_{\text{CO}_2}$.

Mondelice (Mondelice, 2025) developed the Transport Management System interface providing real-time transit-time extraction.

1.2 Document Structure

This document is structured as follows. Section 2 formalises the scientific problem, states the operation's objectives, research questions and hypotheses, and outlines the primary research challenges. Section 3 reviews the state of the art across thematic subsections. Section 4 details the R&D iterations, the methodology, and the serious-game validation campaign. Finally, Section 5 states the documented limitations of the delivered framework and closes with a personal assessment.

2 Scientific Problem and Obstacles to Overcome

2.1 Relevant Context and Motivating Observations

Modern supply chain operators face simultaneous multi-dimensional assessments (cost, deadline, risk, carbon) where most monitoring and planning tools process one criterion at a time (Pinedo, 2012; Chakraborty et al., 2020), which is precisely why practitioners fall back on judgment-based prioritisation.

This work addresses the layer that comes *before* any planning or scheduling decision, and which the classical literature has not fully investigated: the urgency signal itself is systematically biased at the source and must be measured and corrected before it can safely inform any downstream decision (M. Zhu et al., 2018; Rusou et al., 2020).

In a Physical Internet network (Montreuil, 2011), where logistics capacity is mutualized across actors, a biased urgency signal propagates and amplifies systemic misallocations. Over-triage (declaring an urgency higher than the physical situation warrants) starves other network tenants, whereas under-triage (declaring an urgency lower than warranted) creates invisible stock rupture risk that cascades downstream (Li et al., 2020; Ivanov; Dolgui, 2020).

While digital twin architectures and real-time decision support systems have demonstrated strong capabilities in operational monitoring, disruption detection, resource reallocation, and order synchronisation (Ding et al., 2023; Leung et al., 2022; Hrušovský et al., 2021; Karam et al., 2022; X. Zhu et al., 2025), they do so primarily by computing the physical operational state of the system. To the best of the present review, existing systems do not appear to explicitly model the gap between a declared client urgency and a computed operational urgency.

The core scientific novelty of this R&D operation therefore lies in the formal comparison between these declared and computed urgencies, propagated coherently across a multi-tier network. This missing integration motivates the introduction of an urgency adequacy layer in digital operations architectures.

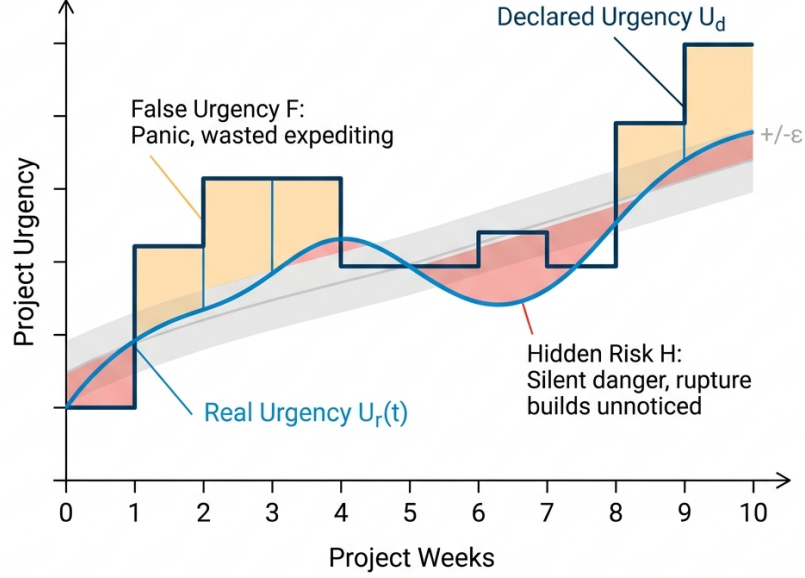


Figure 2: Conceptual view of the declared vs. real urgency gap and its two pathologies. The step line is declared urgency U_d and the smooth curve is real urgency $U_r(t)$; the shaded areas mark where the two mismatch, false urgency F where U_d exceeds U_r , hidden risk H where U_r exceeds U_d .

2.2 Formal Problem Statement

Let $U_d \in [0, 1]$ denote the declared urgency and $U_r(t) \in [0, 1]$ the real urgency computed from operational data in the digital twin. The core variables and domains are summarized in Table 3.

Table 3: Core variables and domains.

Symbol	Meaning	Domain
t	Current time (project hours, ISO week granularity)	$\mathbb{R}$
$U_r(t)$	Physically/operationally computable real urgency	$[0, 1]$
U_d	Operator-declared urgency	$[0, 1]$
ε	Tolerance threshold (cognitive noise band)	$[0, 1]$

The mismatch between both signals is defined as

$$\Delta U(t) = U_d - U_r(t) \in [-1, 1]. \quad (2)$$

The components of this mismatch equation are explained in Table 4.

Table 4: Component definitions for the urgency discrepancy equation.

Component	Type	Description
$\Delta U(t)$	Dependent variable	Raw urgency mismatch signal at time t
U_d	Input variable	Urgency level declared subjectively by the human operator
$U_r(t)$	Dynamic variable	Urgency level computed operationally by the digital twin

A tolerance band $\varepsilon \in [0, 1]$ is introduced so that small deviations are not treated as operationally meaningful. In this setting, over-declaration corresponds to $\Delta U(t) > \varepsilon$, under-declaration to $\Delta U(t) < -\varepsilon$, and approximate agreement to $|\Delta U(t)| \leq \varepsilon$. Potential consequences of these mismatch conditions are summarized in Table 5.

Table 5: Urgency pathologies and potential consequences.

Condition	Interpretation	Potential consequence
$\Delta U(t) > \varepsilon$	Over-declaration of urgency (false urgency)	Excess expediting, avoidable cost, possible delay for other flows
$\Delta U(t) < -\varepsilon$	Under-declaration of urgency (hidden risk)	Hidden lateness or stockout risk, delayed intervention, downstream disruption
$ \Delta U(t) \leq \varepsilon$	Acceptable alignment	No material urgency mismatch

An asymmetric treatment of mismatch may be introduced through weights such as $\lambda_{\text{under}} \geq \lambda_{\text{over}}$ when the application context treats under-prioritisation as more harmful than over-prioritisation. This is motivated by operations literature showing that decision errors can generate unequal downstream consequences (Kim et al., 2020; DiwasSingh et al., 2020), and by scheduling literature that models differentiated tardiness penalties rather than uniform delay costs (Baker et al., 1982; Chauhan et al., 2022). However, the degree of asymmetry is treated here as a design and calibration choice to be validated in the target use case, rather than as a universal property.

2.3 Operation Objectives and Research Questions

The R&D operation is entitled “*Adequacy Score Between Declared and Real Urgency in a Smart Green Supply Chain Digital Twin*”. Its three operational objectives are:

- O1.** Define a multi-block, DAG-propagated real urgency $U_r(t)$ and a cognitive adequacy score $A(t) \in [0, 100]$ to measure the gap between the human planner's declared urgency U_d and the digital twin's operational state.
- O2.** Develop and implement a network propagation and criticality-driven prioritisation layer, combining PROMETHEE II outranking with systematic shock simulation, to identify structurally critical nodes.
- O3.** Validate the adequacy-driven prioritisation logic empirically, through worked scenarios and a calibration service comparing the hidden-risk signal against recorded operational outcomes.

Four research questions structure the scientific approach to these objectives:

- RQ1. Multi-block real urgency ($U_r(t)$):** How can a physically and operationally grounded real urgency $U_r(t)$ be computed dynamically in a multi-tenant digital twin from heterogeneous KPI blocks (temporal, capacity, performance, risk, cost, environmental), and propagated coherently across a multi-tier supply chain graph?
- RQ2. Cognitive adequacy formulation ($A(t)$):** How should an adequacy score $A(t)$ be formulated to measure the alignment between operator declared urgency U_d and computed real urgency $U_r(t)$, while asymmetrically penalizing under-triage vs. over-triage, and can this scoring demonstrate predictive validity against real operational outcomes?
- RQ3. Network prioritisation and criticality:** How can pairwise elicitation methods (AHP, Fuzzy Best-Worst Method) be combined with outranking aggregation (PROMETHEE II) and systematic shock simulation to identify which nodes of a supply network are structurally critical and warrant prioritised attention?
- RQ4. Operational validation within a digital twin:** How can this cognitive-physical framework be operationalised and empirically validated within a persistent, multi-tenant digital twin exercised as a weekly-cadence serious game, and what predictive validity does it demonstrate against recorded operational incidents?

2.4 Research Hypotheses

These hypotheses apply to the entire study. They are formulated here, once, so that the evaluation criteria are fixed before any method is introduced, and they match the pre-registered hypothesis register of the validation campaign described in Section 4.4. Each is falsifiable: it states an observable outcome, with a numeric threshold, that the delivered system either does or does not exhibit.

- H1 (Cognitive latency):** Declared urgency follows computed real urgency with a delay of at least one weekly cycle on the majority of nodes: operators

register a degradation after the physical signal, not with it. Test: the lagged correlation between U_d and U_r peaks at a lag of one week or more for at least five of the eight campaign nodes.

H2 (Early detection of hidden risk): The hidden-risk signal $H = \max(U_r - U_d, 0)$ exceeds 0.3 at the crisis-origin node several weeks before the rupture becomes publicly visible.

H3 (Structural criticality): The systematic shock simulation ranks the true crisis origin first, by impact on the final customer, among deep-tier suppliers, on at least 80% of the informative (non-saturated) weekly rounds.

H4 (Predictive validity): At threshold $H \geq 0.5$ with a four-week horizon, the hidden-risk signal predicts recorded negative outcomes (missed milestone, critical event, node abandonment) with precision and recall both above 0.5.

H5 (Urgency contagion): On weeks when crisis news becomes public, operators of nodes that are not physically affected still raise their declared urgency, reproducing the mere-urgency effect (M. Zhu et al., 2018) at network scale.

H6 (Perception hysteresis): Once the physical situation improves, declared urgency decays at less than half the rate of real urgency: panic outlives its cause.

If H1 and H5 fail, operators need no debiasing and the adequacy layer is redundant. If H2 and H4 fail, the hidden-risk signal is a descriptive curiosity rather than a decision aid. If H3 fails, network prioritisation cannot be trusted operationally. If H6 fails, over-declaration is self-correcting and the asymmetric penalty needs recalibration.

Model Assumptions. Distinct from these testable hypotheses, the framework rests on a set of stated modelling assumptions, each verifiable in the delivered code: the network is a directed acyclic graph (no feedback loops, linear topological sort); reviews follow a strict ISO weekly cycle; declared urgency is damped as it descends toward suppliers while real risk is amplified as it ascends toward the customer, each scaled by a per-arc coefficient; the six KPI blocks are treated as independent to permit a probabilistic aggregation; task statuses drive explicit override rules; and the cognitive loss is deliberately asymmetric, with under-declaration priced at roughly twice the cost of panic. These assumptions are design commitments, not empirical claims; the limitations they induce are catalogued in Section 5.1.

2.5 Technical and Scientific Challenges

To evaluate the validity of the proposed framework, its performance is investigated against five scientific challenges identified in contemporary literature. Each challenge is stated here as an open problem; the design choices made to overcome them are in Section 4.

Challenge 1: Justifying the Urgency Aggregation

Real urgency can be modeled from a single deadline-proximity signal or from an aggregation of heterogeneous risk blocks (slack, capacity saturation, performance, failure exposure, cost drift, emissions overshoot). A single-block, purely time-driven curve is analytically convenient but discards information available in a live digital twin. Conversely, any multi-block aggregation rule must be theoretically justified against simpler linear or averaging alternatives (K. Chen et al., 1992; K. Chen et al., 1996; J. Chen et al., 2025; Mills et al., 2010): which rule preserves the visibility of a single catastrophic indicator without letting noise dominate, and on what grounds?

Challenge 2: Reliability of Declared Urgency

Eliciting the declared urgency U_d via pairwise multi-criteria methods introduces cognitive noise and ranking instability. These methods are sensitive to human inconsistency and rank reversal, meaning that declared urgency cannot be assumed as a deterministic ground truth (Afrasiabi et al., 2022; Triantaphyllou; Yanase, 2024; Liang et al., 2020). Operator judgments must therefore be treated as noisy, stochastic inputs (Stam et al., 1997): how can an elicitation protocol detect and reject inconsistent judgments before they contaminate the network state?

Challenge 3: Value-Added of the Adequacy Signal

Computing $\Delta U(t) = U_d - U_r(t)$ is necessary but not sufficient. The decisive scientific test is whether this adequacy signal improves the identification of operationally at-risk nodes beyond the raw variables already available to an operator (local KPI thresholds, unweighted urgency, or simple node degree) (Cohen et al., 2019; Bojko et al., 2021). If the score does not demonstrably outperform these simpler baselines against recorded outcomes, it serves merely as an auxiliary decision metric rather than a primary planning signal (Ahmetoglu, 2025; Turner et al., 2021).

Challenge 4: Composite Score Interpretability

Aggregating multiple criteria into a single composite score, whether the local urgency U_r across six KPI blocks or a ranking across a network of nodes, risks obscuring trade-offs and weakening decision interpretability. Composite-index literature warns that nominal weights do not always reflect real influence due to component correlations (Greco et al., 2018; Becker et al., 2017; Borgonovo et al., 2021): how can a coordinator still see *which* indicator is driving a node's urgency once everything is folded into one number?

Challenge 5: No Precedent for Urgency Adequacy

Digital twin systems routinely monitor real-time operational states, disruption responses, and order synchronization (Ding et al., 2023; Leung et al., 2022; J. Zhang et al., 2021; Y. Zhang et al., 2026; Zhou et al., 2021). However, to the best of our knowledge, no existing digital twin or decision support system simultaneously captures a client-declared urgency signal, computes a model-based operational urgency from live system state, and measures their mismatch as an active, network-propagated indicator. The absence of precedent cuts both ways: it signals an opportunity, but it also means there is no established architecture, calibration practice, or evaluation benchmark to build on (Zheng; Lu, et al., 2021; Niloofar et al., 2023; Kalaboukas et al., 2021).

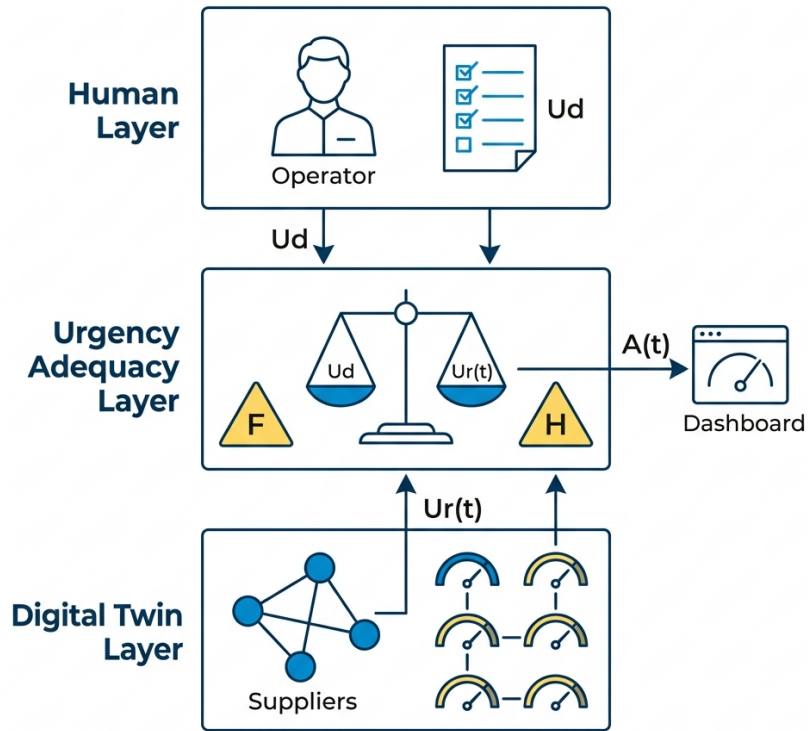


Figure 3: The three-layer urgency adequacy architecture investigated in this work (Challenge 5).

3 State of the Art

3.1 Introduction: Key Definitions

Urgency is defined here as the degree to which a supply chain task requires prompt action to avoid a material deterioration in operational performance. It differs from *importance*, which reflects strategic value, and from *priority*, which reflects the current execution order. In this work, urgency is treated as a multi-dimensional operational quantity, aggregated from heterogeneous KPI evidence rather than from a single deadline-driven signal (J. Chen et al., 2025; M. Zhu et al., 2018).

Real-time scheduling refers to the class of algorithms that assign execution order to tasks given temporal constraints (deadlines, inter-task dependencies), resource constraints (vehicle capacity, supplier availability), and optimality criteria (makespan, weighted tardiness) (Pinedo, 2012). This literature is drawn on here only as historical grounding for deadline-sensitive value functions; the present work does not deliver a scheduling algorithm.

Digital twin is a synchronized digital representation of a physical system that supports monitoring, simulation, and decision-making using real-time and historical data (Ivanov; Dolgui, 2020). In supply chains, digital twins are used to improve visibility, disruption assessment, and replanning under changing conditions (Ivanov; Dolgui, 2020; Ding et al., 2023; Hrušovský et al., 2021).

Multi-criteria decision-making (Multi-Criteria Decision-Making (MCDM)) covers methods that combine multiple, potentially conflicting performance dimensions into a decision (Saaty, 1980; Brans; Vincke, et al., 1986). In the present work, the main dimensions considered for node prioritisation are time pressure, capacity saturation, operational performance, failure risk, cost drift, and carbon overshoot; these six dimensions are named here only as vocabulary and are developed into a formal aggregation model in Section 4.2.

Physical Internet (Physical Internet (PI)) is the model of open, mutualized logistics networks that route shipments through shared infrastructure using standardised modular containers (Montreuil, 2011). In such shared logistics environments, inaccurate urgency assessment may have effects beyond a single actor by influencing the allocation of mutualized capacity. This makes urgency evaluation particularly important in Physical Internet settings (Montreuil, 2011; Leung et al., 2022).

Serious game is a simulation environment designed to train decision-makers through realistic scenarios with immediate feedback and scoring (Sterman, 1989; Prensky, 2001). In this project, DISCO functions both as a serious game and as a digital twin: users interact with simulated supply chain decisions, while the

underlying digital model evaluates operational consequences on a weekly cadence.

3.2 Literature Review

This review sets out to answer a single question: which existing methods can be assembled to measure, correct, and propagate a biased urgency signal, and what part of that assembly is missing from the literature? The relevant work is broad but fragmented across operations research, behavioural science, and digital twin engineering. Read individually, each stream solves one piece of the problem; read together, they leave a specific slot empty. No retrieved work combines, within a single supply-chain decision architecture, the five elements this thesis requires: (i) a multi-block operational urgency function, (ii) a structured client-side urgency input, (iii) an adequacy layer comparing stated and operational urgency, (iv) network propagation of that urgency across a multi-tier DAG, and (v) empirical validation of the adequacy signal against recorded outcomes (Sawik, 2013; Sawik, 2014; Sawik, 2016; Cohen et al., 2019; Hosseini; Ivanov, 2020).

The review is therefore organised by the role each stream plays in this argument, rather than by chronology or by method family. Two *foundational* streams establish why the problem exists: deadline-sensitive value and network risk propagation on one side, cognitive bias in urgency declaration on the other. Three *enabling* streams supply the computational toolkit: Bayesian belief revision, digital twin operations, and multi-criteria decision methods. One stream gathers the *closest antecedents*, whose precisely identified gap isolates what this work adds. Finally, *analogical* literature (medical triage, the Physical Internet, carbon accounting) provides vocabulary and context only, and is flagged as non-foundational wherever it is used.

Table 6 summarises the seven streams and their relationship to the thesis.

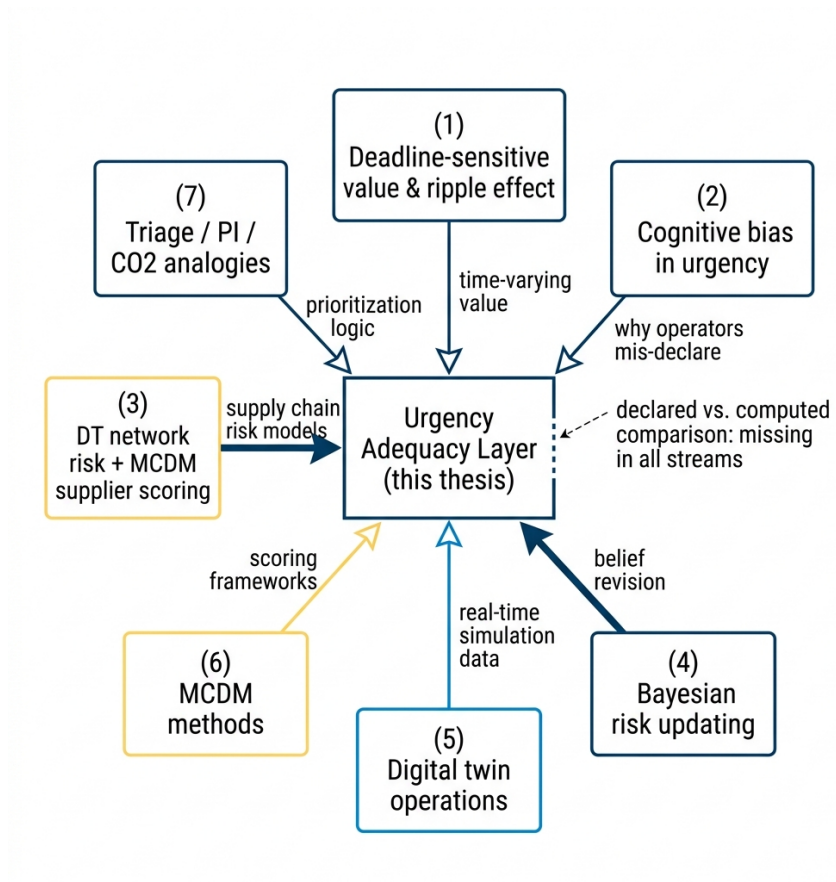


Figure 4: Map of the seven literature streams and the gap they leave open.

Table 6: Literature stream synthesis: contribution and gap relative to this thesis.

Stream	Role	Main contribution
Risk propagation and deadline-sensitive value	Foundation	Monotone time-value functions; ripple-effect propagation on networks
Cognitive bias in urgency	Foundation	Documents urgency mis-declaration patterns
Digital twins of network risk and MCDM supplier scoring	Closest antecedent	Multi-criteria supplier scoring and network disruption monitoring
Bayesian and probabilistic risk updating	Enabling method	Conjugate/probabilistic belief revision under new evidence
Digital twin operations	Enabling method	Real-time replanning and disruption response
MCDM methods	Enabling method	Criteria weighting and outranking
Medical triage / PI / CO ₂	Analogy / Context	Asymmetric urgency vocabulary; network context

3.2.1 Stream 1: Risk Propagation and Deadline-Sensitive Value

Role: Foundation.

Why: It demonstrates that operational value can be formulated as a time-dependent and network-dependent quantity, providing the primary motivation for treating urgency as a variable that is both time-varying and propagated across a multi-tier network rather than computed at a single node in isolation.

Within real-time systems, deadline-sensitive scheduling theory models task value as a function of completion or response time, an approach known as time-value or utility-accrual scheduling (K. Chen et al., 1996; J. Chen et al., 2025). Chen and Mühlethaler (K. Chen et al., 1996) show that task utility is typically constant up to a deadline and then non-increasing with delay, and a broader survey of utility-accrual scheduling confirms that deadline-sensitive literature models task value as a function of response time with bounded, structured degradation after critical thresholds (J. Chen et al., 2025). This establishes the historical precedent for treating urgency as time-varying, but it stops at a single task in isolation.

A second, complementary line of evidence comes from network risk-propagation

research. Ivanov and colleagues formalise the ripple effect, whereby a local disruption at one supplier cascades through downstream tiers and amplifies as it travels along the network (Ivanov; Dolgui, 2020; Ivanov; Dolgui, 2019; Li et al., 2020). Probabilistic risk-inference methods for supply chain systems (Zheng; L. Zhang, 2020a) further show that aggregating heterogeneous risk indicators into a single node-level score is itself an open modelling question, distinct from how that score should then be propagated.

Table 7 compares four candidate aggregation and growth models found across these two lines of literature.

Table 7: Comparative urgency aggregation and growth models and their properties.

Model	Formula	Key property	Use case
Linear (time)	$U_0 + \lambda t$	Constant pressure	Administrative tasks
Hyperbolic (time)	$K/(t_c - t)$, bounded via $\tanh$	Finite-time singularity, discounting	Synthetic scenario generation
Weighted average (blocks)	$\sum_m \omega_m u_m$	Fully compensatory: one weak block hidden by strong ones	Simple composite indices
Noisy-OR (blocks)	$1 - \prod_m (1 - u_m)^{\omega_m}$	Non-compensatory: one saturated block dominates	Multi-block operational risk

This thesis adopts the noisy-OR aggregation over the fully compensatory weighted average, on the grounds that a single catastrophic KPI block (for instance a complete capacity failure) should not be masked by acceptable scores on the remaining blocks, a property established in the probabilistic-inference literature on combining independent causal evidence (Pearl, 1988). The bounded hyperbolic discounting curve is retained separately, purely as an auxiliary tool for generating realistic synthetic urgency trajectories used in testing and calibration; a structurally similar acceleration appears in Log-Periodic Power Law (Log-Periodic Power Law (LPPL)) models of financial crash precursors (Johansen et al., 2001; Johansen et al., 2000), used here strictly as an analogical reference for that auxiliary generator, since empirical analyses by Feigenbaum (Feigenbaum, 2001) and Brée et al. (Brée et al., 2010) show that LPPL-family models are highly parameter-sensitive and unsuited to direct use as a production risk model.

Summary.

Known: Deadline-sensitive scheduling supports time-dependent task valuation (J. Chen et al., 2025); ripple-effect research supports network propagation of local disruptions (Ivanov; Dolgui, 2020; Li et al., 2020).

Assumed: A non-compensatory, noisy-OR aggregation of KPI blocks is proposed as the local urgency model, to be propagated with a leaky product rule across the network.

Open gap: Neither stream specifies how to aggregate heterogeneous, block-structured operational risk indicators at a single node before propagating them, and neither addresses user-declared urgency or the mismatch between stated and operational urgency.

3.2.2 Stream 2: Cognitive Bias in Urgency Declaration

Role: Foundation.

Why: It establishes that client-declared urgency is systematically biased and cannot be taken at face value as a planning input.

Behavioural operations research suggests that human decision-makers systematically misrepresent urgency, a finding with direct implications for any system that relies on operator-declared urgency as an input.

Two convergent lines show this from different angles. Zhu, Yang and Hsee (M. Zhu et al., 2018) report the Mere-Urgency Effect in five controlled experiments: participants consistently over-prioritised imminent low-value tasks over distant high-value tasks, and the bias persisted even when participants were informed of its existence. Steel (Steel et al., 2006) provides a complementary formal model through Temporal Motivational Theory (Temporal Motivational Theory (TMT)):

$$V(t) = \frac{U \times E}{1 + \Gamma \times D(t)^k}, \quad (3)$$

where U is objective utility, E is the expectancy of success, $D(t)$ is the delay to deadline, and Γ , k are individual discounting parameters. As the delay to deadline shrinks toward zero, the motivational value $V(t)$ converges to the full product of utility and expectancy: immediacy becomes more salient as the deadline approaches. Near-term tasks can therefore attract disproportionate attention even when their broader objective value is lower (Steel et al., 2006), which helps explain why imminent tasks may be over-declared as urgent. The operational roles of the parameters in this formulation are structured in Table 8.

Table 8: Temporal Motivational Theory (TMT) parameters and domains.

Symbol	Concept	Operational Role
$V(t)$	Temporal motivational value	Dynamic task salience
U	Objective utility	Task strategic importance
E	Expectancy of success	Probability of task completion
$D(t)$	Delay to deadline	Remaining safe execution window
Γ	Impulsiveness parameter	Operator sensitivity to delay
k	Sensitivity scaling parameter	Non-linear scaling of delay discounting

These laboratory findings are supported by evidence from operational settings. Rusou, Amar and Ayal (Rusou et al., 2020) show that operators under workload preferentially complete small, easy tasks, the operational analogue of under-triage for large-scope supply chain activities. DiwasSingh et al. (DiwasSingh et al., 2020) confirm in a large field study that completing easy tasks first significantly hurts overall performance. Hospital admission control under occupancy uncertainty (Kim et al., 2020) produces over- and under-triage patterns analogous to those in logistics operator decisions under uncertain lead times. The magnitude of these biases in industrial supply chain contexts remains a subject of calibration. To mitigate these time-perception and workload biases, human-centered decision support systems (such as Rosenthal and Hiatt (Rosenthal et al., 2020)) use visual feedback techniques to show the mismatch between human plans and system temporal constraints, helping users calibrate their schedules.

This study treats declared urgency as a noisy signal and computes a model-based correction. The correction mechanism, and its empirical predictive validity, is detailed in Section 4.

Summary.

Known: Urgency over-declaration for near-deadline tasks is documented in controlled settings (M. Zhu et al., 2018) and in operational settings (DiwasSingh et al., 2020; Rusou et al., 2020).

Unknown: Bias magnitude in industrial supply chain contexts.

Assumed: Asymmetric penalty $\lambda_{\text{under}} > \lambda_{\text{over}}$ is the appropriate correction structure.

Open gap: Behavioural research diagnoses the problem. It does not prescribe a real-time operational correction.

3.2.3 Closest Antecedents: Digital Twins of Network Risk and MCDM Supplier Scoring

Role: Closest antecedents.

Why: These are the works most structurally similar to the thesis. Identifying their gap precisely isolates what this work adds.

The works most structurally similar to this thesis combine multi-criteria supplier scoring with disruption-aware network monitoring. Sawik (Sawik, 2013; Sawik, 2014; Sawik, 2016) develops mixed-integer programming models for resilient supplier order assignment integrating cost, service level, and disruption risk. Cohen, Grinshpoun and Stern (Cohen et al., 2019) and Afrasiabi et al. (Afrasiabi et al., 2022) extend this to MCDM supplier prioritisation under disruption, combining criteria weighting with allocation optimisation. Chakraborty et al. (Chakraborty et al., 2020) confirm that rule-based heuristics remain competitive against exact methods in dynamic industrial environments. These works are cited here for the closeness of their subject matter (multi-criteria, disruption-aware supply chain decisions), not because this thesis reuses their optimisation machinery, which addresses allocation and sequencing rather than urgency measurement.

Taken together, this stream models disruption-aware supplier evaluation with multi-criteria weighting. Two gaps distinguish it from the present work. First, none introduces a network-propagated urgency signal computed independently at every node and combined across tiers; risk is typically evaluated once, at the point of decision. Second, none compares a client-stated urgency against a model-derived urgency: the client-declared signal, where it exists, is used directly, without validation against live operational proxies.

3.2.4 Enabling Method: Bayesian and Probabilistic Risk Updating

Role: Enabling method.

Why: Revising a failure probability from a small number of new observations is a component of operational risk state. Bayesian conjugate updating is the borrowed tool.

Bayesian methods have emerged as a well-supported approach for revising risk estimates from new evidence in supply chain research (Pearl, 1988; Koller et al., 2009). Rather than a full Bayesian network with an explicit conditional-probability-table structure, this study draws on the narrower and more tractable case of Beta-Bernoulli conjugate updating: a failure probability p is treated as a Beta-distributed belief with an equivalent prior sample size n_0 , revised from a small number of newly observed equivalent trials (k, n) each time a calibrated operational event is

declared. This is the standard conjugate-prior construction for binomial-outcome belief revision (Pearl, 1988; Koller et al., 2009).

Three convergent findings from the retrieved Bayesian supply-chain literature support this choice of tool. Hosseini and Ivanov (Hosseini; Ivanov, 2020) provide the fullest survey of Bayesian methods applied to supply chain risk and resilience, finding that probabilistic belief-revision methods are often preferred to simple deterministic scoring when new, sparse evidence must update an existing risk estimate. Huang et al. (Huang et al., 2020) and Deng, Zhang and Wei (Deng et al., 2020) develop Bayesian models for disruption and delay probability estimation that serve as a methodological precedent for revising a KPI block's failure probability from a declared event. Lockamy (Lockamy et al., 2011) and Qazi et al. (Qazi et al., 2017) further support Bayesian revision as a practical tool for dynamic supply chain risk quantification under sparse, event-driven evidence, consistent with the small-sample setting of a weekly-cadence digital twin. Plain frequency counts, by contrast, are unstable when only a handful of events are declared per node per project, which motivates the conjugate-prior formulation adopted here over a raw frequentist estimate (Zheng; L. Zhang, 2020b).

The specific conjugate-update parameters used for each event family are developed in Section 4.

Summary.

Known: Bayesian conjugate updating is an established, low-data-regime tool for revising failure probabilities in logistics contexts.

Assumed: A shared, moderately informative prior (equivalent to roughly six months of weekly observations) is sufficient for initial calibration.

Open gap: Bayesian belief revision quantifies disruption exposure but does not, by itself, compare it against client-declared urgency.

3.2.5 Enabling Method: Digital Twin and Real-Time Decision Support

Role: Enabling method and primary gap source.

Why: DT systems are the architectural family closest to DISCO; this review establishes both what they already do and the specific gap this thesis addresses.

The digital twin concept was formalised by Grieves (Grieves, 2014) as a virtualised representation of a physical process. Tao et al. (Tao et al., 2018) extended the concept to manufacturing; Fuller et al. (Fuller et al., 2020) surveyed the enabling technologies behind modern deployments. In supply chain management, Ivanov et al. (Ivanov; Dolgui, et al., 2021) define the Supply Chain (SC) Digital Twin (DT) as integrating simulation (what-if analysis), optimisation (decision support), and

machine learning (pattern detection). Ivanov and Dolgui (Ivanov; Dolgui, 2020) address the use of digital SC twins for disruption risk and resilience management in the Industry 4.0 era. Wang, Wang and Liu (Wang et al., 2020) demonstrate DT-driven planning in a manufacturing context analogous to DISCO's satellite component network. More recently, in the context of Industry 5.0, these systems have evolved towards Human-in-the-Loop (HITL) Cognitive Digital Twins (CDTs) where semantic models and reasoning methods are used to capture human feedback and synchronize mental models with simulated physical operations (Niloofer et al., 2023).

Real-time freight planning work demonstrates that operational urgency proxies (slack, lateness risk, and shortage exposure) are present within these systems (Hrušovský et al., 2021; Perez et al., 2022). Model-predictive control approaches for vehicle routing (Karam et al., 2022) and digital-twin disruption response (Ding et al., 2023; Leung et al., 2022; X. Zhu et al., 2025) confirm the state of the art in live operational replanning.

Taken together, the digital-twin literature supports simulation, optimisation, and disruption-aware decision support. In the retrieved literature, however, no system surfaces a standard module that captures a client-declared urgency signal as a structured input, computes a model-based operational urgency from the live system state, and measures their mismatch as an active, network-propagated indicator.

Summary.

Known: Digital twins support simulation, optimisation, and disruption-aware decision support in supply chains (Ivanov; Dolgui, et al., 2021; Ivanov; Dolgui, 2020), and recent Cognitive Digital Twin work captures human feedback to synchronise mental models with simulated operations (Niloofer et al., 2023).

Open gap: No retrieved system combines a declared urgency input, a computed operational urgency, and their propagated mismatch within one architecture.

Positioning of this work. The three-layer architecture that fills this gap, capturing declared urgency, computing operational urgency, and propagating their mismatch across the network, is the central construction proposed in this thesis, not a summary of an existing system (stated as an observation from the retrieved literature, not as a claim about the totality of published work). Whether it improves the identification of at-risk nodes is the central empirical hypothesis under test (Challenge 3, Section 2.5).

3.2.6 Enabling Methods: Multi-Criteria Decision Methods

Role: Enabling methods.

Why: These multi-criteria decision-making (MCDM) tools provide the computational foundation for criteria weighting, outranking, and network prioritisation.

This research integrates and builds upon two distinct methodological domains within multi-criteria decision-making: pairwise comparison methods for eliciting weights, and outranking aggregation for turning those weights into a defensible ranking of nodes.

Pairwise Elicitation and Pairwise Comparison Methods. To construct a structured representation of operator-declared urgency (U_d), criteria weights must be systematically elicited. The Analytic Hierarchy Process (Analytic Hierarchy Process (AHP)) (Saaty, 1980) derives criteria weights as the principal eigenvector of a pairwise comparison matrix, validating logical consistency via the Consistency Ratio ($ConsistencyRatio(CR) \leq 0.10$). While widely applied in logistics, AHP suffers from rank reversal, where the introduction of a new alternative alters the relative ranking of existing ones, and is sensitive to cognitive noise during expert elicitation (Dyer, 1990; Triantaphyllou; Yanase, 2024).

To minimize operator cognitive load and improve consistency, the Best-Worst Method (Best-Worst Method (BWM)) (Rezaei, 2015) reduces the number of required pairwise comparisons from $n(n-1)/2$ to $2(n-1)$ by anchoring comparisons around chosen "best" and "worst" criteria. This is further extended by the Fuzzy Best-Worst Method (Fuzzy Best-Worst Method (F-BWM)) (S. Guo et al., 2017; Yucesan et al., 2019; Amiri et al., 2020), which uses Triangular Fuzzy Numbers (Triangular Fuzzy Numbers (TFNs)) to absorb human judgment uncertainty. In this study, consistent FBWM weight vectors (Liang et al., 2020) serve as an alternative elicitation gate to crisp AHP, and both feed the same downstream propagation and adequacy machinery.

Outranking Aggregation, Network Criticality, and Composite Indicator Critique.

For prioritising which nodes require attention, direct weighted averages risk full compensatory aggregation, where an excellent score on one criterion (e.g., cost) completely masks a critical failure on another (e.g., capacity). To prevent this, outranking methods such as PROMETHEE II (Brans; Vincke, et al., 1986; Brans; Mareschal, 2005) establish pairwise preference flows (ϕ^+ and ϕ^-), allowing a non-compensatory ordering of alternatives. A survey of PROMETHEE applications across industries (Behzadian et al., 2010) confirms its suitability for supplier and node evaluation. In this work, PROMETHEE II ranks active network nodes directly

on the same weighted KPI blocks that feed the local urgency aggregation, giving a network-wide outranking view that complements the per-node adequacy score. Systematic what-if shock simulation, in which a node’s local risk is forced to saturation and the resulting propagated delta is measured across the network, provides a second, independent criticality signal that is combined with the outranking flow to identify structural bottlenecks.

However, constructing any composite indicator from weighted criteria introduces methodological challenges. Composite-index literature warns that nominal weights assigned to criteria often diverge from their effective importance due to correlation, normalization, and aggregation dynamics (Greco et al., 2018; Becker et al., 2017; Borgonovo et al., 2021), and the choice of normalization method is itself non-neutral: it can lead to rank reversal (Triantaphyllou, 2000). To keep the aggregation weights interpretable and defensible, this study exposes per-block contributions on the operator explainability page rather than only the aggregate score, while validating the risk parameters under uncertainty (Sawik, 2020).

3.2.7 Analogical and Contextual Literature

Role: Analogies and context.

Why: These bodies of literature provide conceptual analogies (medical triage), network contexts (Physical Internet), or regulatory frameworks (carbon accounting) that support the design of the adequacy variables.

While the core scoring and network-propagation methods are mathematically derived, their operational structure draws grounding and context from three non-foundational literature streams.

Medical Triage as a Conceptual Analogy. Medical triage provides a rich conceptual analogy for prioritizing tasks under severe capacity constraints and cognitive pressure. Standard protocols such as the Manchester Triage System (Manchester Triage System (MTS)) (Mackway-Jones et al., 2013) categorize patients into five urgency levels using flowcharts, whereas the Emergency Severity Index (ESI) integrates predicted resource consumption. Comparative clinical studies (Kuligowski et al., 2025) show significant classification discrepancies (weighted $\kappa = 0.51$) between these systems, highlighting that even trained professionals struggle with consistent urgency assessment.

In a supply chain context, this triage process maps directly to node prioritisation. The design of the cognitive adequacy score $A(t)$ is motivated by the clinical principle of asymmetric risk: under-triage (under-declaring urgency, leading to untreated critical conditions) is far more dangerous than over-triage (over-declaring

urgency, leading to resource congestion) (Kim et al., 2020). Consequently, the adequacy score enforces an asymmetric penalty ($\lambda_{\text{under}} > \lambda_{\text{over}}$) to protect the physical supply chain from hidden stockout ruptures. Table 9 maps these triage concepts to their logistics equivalents.

Table 9: Heuristic mapping of medical triage concepts to DISCO logistics prioritisation.

Medical concept	DISCO logistics equivalent
Patient priority P1 (immediate)	Node with $A(t) < 20$ and high network criticality rank
Under-triage (P3 classified as P1)	Hidden risk $H = [U_r - U_d]_+$: silent bottleneck builds unnoticed
ESI resource consumption estimate	DISCO criticality index (propagated impact on final-customer nodes)
Resource-blind triage (MTS)	Local urgency threshold ignoring network position
Resource-aware triage (ESI)	Criticality-weighted prioritisation of coordinator attention

Physical Internet and Supply Chain Sychromodality. The Physical Internet (PI) (Montreuil, 2011) conceptualizes open, mutualized logistics networks where routing, storage, and transport assets are shared. In such hyperconnected systems, urgency ceases to be a localized concern and becomes a network-level Quality of Service variable. A related but distinct concept is sychromodality: the ability to shift a shipment between transport modes, road, rail, or waterway, in real time as conditions change, rather than committing to one mode in advance. Sychromodal networks and the Physical Internet can operate together, but neither presupposes the other, and dynamic routing and order synchronization within a sychromodal network require real-time, responsive priority signals compatible with rolling planning windows (W. Guo et al., 2020). Disruptions also propagate rapidly through hyperconnected networks, creating cascading ripple effects (Li et al., 2020) that validate the need for a synchronized, digital-twin-driven propagation model (Ivanov; Dolgui, 2020).

Regulatory Carbon Accounting. The integration of an environmental criterion is motivated by modern corporate sustainability requirements, such as the Corporate Sustainability Reporting Directive (CSRD Scope 3) and ALTEN's environmental guidelines (Cordova, 2024). Standardized frameworks like the GLEC Framework

v3.2 (Smart Freight Centre, 2023) and Vehicle Energy Consumption calculation TOol (VECTO) (European Commission, 2017) define emissions calculation baselines for heavy duty vehicles. ADEME (ADEME, 2014) reports that emissions factors are non-linearly related to vehicle speed and cargo load, supporting the treatment of carbon overshoot as its own urgency block rather than folding it into a generic cost term.

3.2.8 Validation Methodology: Serious Games

Role: Validation rationale.

Why: It justifies the weekly-cadence serious game used as the validation environment for the whole framework.

The Beer Game (Sternan, 1989) is the foundational supply chain serious game: Sternan shows that players systematically misperceive the feedback in their own supply line, underestimating the orders already placed but not yet received, and over-order with a lag once a shortage appears, amplifying the resulting oscillation as it travels upstream, the mechanism now known as the bullwhip effect (Lee et al., 1997). This decision bias is analogous to the false-urgency pathology in logistics, where an operator over-reacts to a signal only after a delay rather than adjusting smoothly as the signal changes. The finding that game-based observation captures systematic decision biases supports the use of digital twin serious games as an empirical instrument. Prensky (Prensky, 2001) establishes the theoretical basis for digital game-based learning, and Deterding et al. (Deterding et al., 2011) formalise gamification design principles.

Mohaddesi et al. (Mohaddesi et al., 2020) introduce the Gamette methodology, a playful, bite-sized approach to capturing decision-making patterns to build behavioural models, directly applicable to a weekly-review protocol in which operators repeatedly submit urgency assessments over the course of a project. Bellotti et al. (Bellotti et al., 2013) provide an evaluation framework for serious game learning outcomes, which supports the design of the calibration methodology detailed in Section 4. Other logistics serious games study professional decision-making under conditions that do not yet exist in practice: Lafkihi, Pan and Ballot (Lafkihi et al., 2025) use a freight transportation game to study carrier behaviour on online spot-market platforms, a setting connected to the Physical Internet's vision of mutualised, on-demand capacity allocation.

3.3 Summary: Limitations of Existing Solutions

A review of the literature indicates that classical operational monitoring typically evaluates a single criterion (a deadline, a cost, or a risk score) at a time, failing to capture the multi-block nature of supply chain KPIs or the systemic impact of operator decision biases. While pairwise comparison methods (AHP, BWM, and Fuzzy BWM) are widely applied for criteria weighting, they function as static elicitation tools and are rarely compared against a computed operational counterpart. Outranking methods like PROMETHEE II prevent compensatory aggregation but are not usually combined with a network-propagated urgency signal.

While contemporary supply chain digital twins and Bayesian risk models offer strong capabilities in real-time disruption monitoring and belief revision, they typically lack a mechanism to compare human operator inputs with the physical state of the digital twin. This review identifies a key gap: the absence of an integrated architecture that compares a declared client-side urgency signal against a computed, network-propagated operational urgency to evaluate and correct cognitive mismatch, and that validates this correction empirically against recorded operational outcomes. Whether this adequacy-driven prioritisation improves the identification of at-risk nodes compared to standard KPI baselines is the central open question of this R&D operation.

Placed against the five elements this thesis requires (Section 3.4 restates them as a two-layer architecture), the reviewed streams individually cover multi-block risk aggregation, network propagation, digital twin operations, and multi-criteria elicitation; none of them combines those elements with a structured client-side urgency input and an empirically validated adequacy layer. That missing combination, not any single missing method, is what the remainder of this document builds and tests.

3.4 Research Axis and Proposed Adequacy-Propagation Structure

To investigate this gap, this study proposes an architecture centered on a cognitive-physical adequacy check combined with network-wide risk propagation. The main objective is to design and evaluate a two-layer control loop:

1. **Cognitive adequacy check:** This layer captures the operator's declared urgency U_d (elicited through consistency-gated pairwise comparisons) and compares it with a physically-grounded real urgency $U_r(t)$ computed from live digital twin state. The mismatch is processed using an asymmetric penalty to debias operator judgments.
2. **Network propagation and criticality-driven prioritisation:** Local urgency

is propagated across the supply chain DAG (declared urgency descending toward suppliers, real urgency ascending toward the final customer), and the resulting network state feeds both a PROMETHEE II outranking ranking and a systematic shock-simulation criticality index, jointly identifying which nodes warrant prioritised operator attention.

This two-layer architecture is exercised and empirically validated within a persistent, multi-tenant digital twin operated as a weekly-cadence serious game, closing the loop between the adequacy signal and observed operational outcomes (Section 4).

4 Conducted Operations

4.1 Methodology: The Agile Test-and-Learn Program

This R&D operation is structured as an agile Test-and-Learn program designed to investigate cognitive biases in supply chain urgency assessment and formulate a digital-twin-based correction loop. Rather than following a rigid sequential path, the program operates through continuous iterations of literature analysis, mathematical specification, software implementation, and scenario-based simulation.

The R&D pipeline is organized into four core operational areas:

- **Literature Review and Gap Analysis:** Establishing the psychological and mathematical foundations of multi-block urgency aggregation and cognitive bias (detailed in Section 3).
- **Mathematical Framework Specification:** Deriving and formalizing the multi-block real urgency function, its DAG propagation, and the asymmetric cognitive adequacy score from first principles to ensure dimensional and probabilistic consistency.
- **Pairwise comparison elicitation:** Formulating crisp Analytic Hierarchy Process (AHP) and Fuzzy Best-Worst Method (F-BWM) eigenvectors to build a reliable representation of declared operator urgency.
- **Implementation, Simulation and Validation:** Building a modular, tested Python library (65 modules, 1425 automated tests) exercised through a multi-page Dash web application, running controlled shock and scenario simulations on a persistent digital twin, and confronting the hidden-risk signal with recorded outcomes through a dedicated calibration service and a serious-game campaign (Section 4.4).

The remainder of this section follows that same separation explicitly: the mathematical model is presented on its own (Section 4.2), independently of how it is elicited, computed, and exposed (Section 4.3), before turning to the use case built to validate it (Section 4.4).

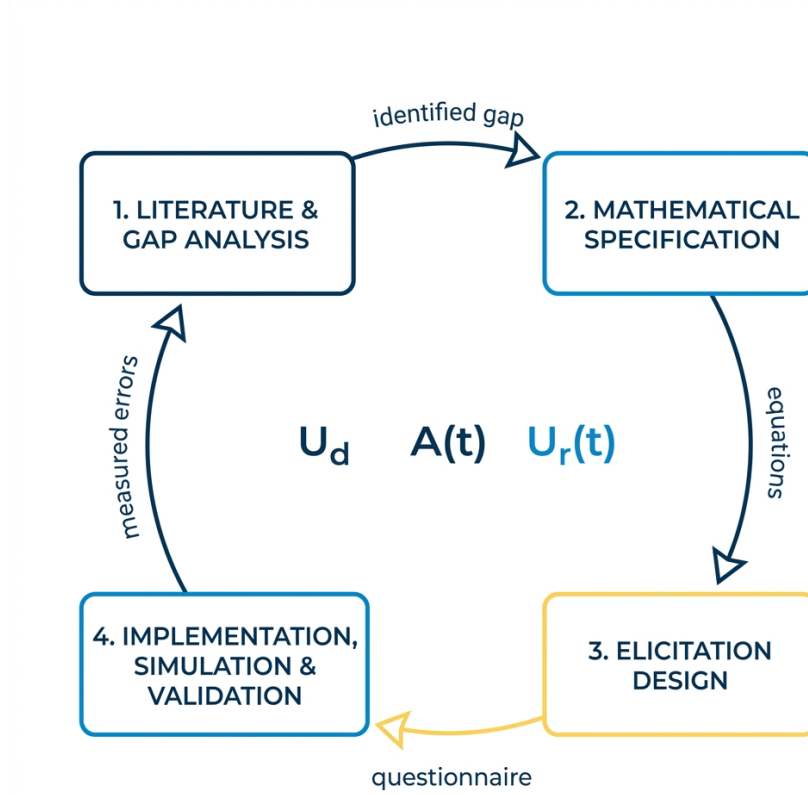


Figure 5: The agile Test-and-Learn iteration loop of the R&D operation.

4.2 Mathematical Model

This section states the model on its own terms: what real urgency and declared urgency are, how each is propagated across the network, and how their mismatch is scored. It intentionally leaves out tooling and interface choices, elicitation mechanics, and operational cadence, which are application decisions addressed separately in Section 4.3.

4.2.1 Core Model: Real and Declared Local Urgency, and Network Propagation

Local Real Urgency ($U_{r,local}$). Real urgency is not read from a single deadline-proximity signal; it is aggregated at each node from six independent KPI blocks, $\{\text{time, cap, perf, risk, cost, co2}\}$, each producing a partial urgency $u_m \in [0, 1]$ or `None` when the underlying KPI data is unavailable. The aggregation is a weighted

probabilistic OR, so that a single saturated block is sufficient to drive the local urgency toward its maximum, regardless of how well the other blocks score:

$$U_{r,\text{local}} = 1 - \prod_m (1 - u_m)^{\omega_m}, \quad (4)$$

where $\omega_m \geq 0$ is the aggregation weight of block m , and the product runs over blocks with available data. Table 10 gives, for each block, its schematic formula and what it measures; the notation $[x]_+ = \max(x, 0)$ denotes the positive part, and the per-term weights (α_i, β_i) are renormalised over the terms whose data is actually available.

Table 10: The six real urgency blocks: formula and explanation.

Block	Formula	Explanation
u_{time}	$P(L > d^* - t) + \kappa_{\text{retard}}[r]_+ - \kappa_{\text{avance}}[-r]_+$	Probability that the lead time L , modelled as $\text{Normal}(\mu, \sigma)$, exceeds the time left before the next active deadline d^* ; forced to 1 once the deadline is passed. The planning term r compares theoretical progress with reported progress: a milestone behind plan raises the block, one ahead of plan lowers it.
u_{cap}	$\alpha_1 \left(1 - \frac{V_{\text{rem}}}{V_{\text{max}}}\right) + \alpha_2 \left(1 - \frac{W_{\text{rem}}}{W_{\text{max}}}\right) + \alpha_3 \left[\frac{D - \Phi}{D}\right]_+$	Saturation of the node's capacity: how little storage volume and weight remain, plus the fraction of demand D not covered by the current flow rate Φ .
u_{perf}	$1 - \text{OEE}$	Complement of Overall Equipment Effectiveness (availability $\times$ performance $\times$ quality): the lower the productive health of the critical resource, the higher the urgency.
u_{risk}	$1 - \exp(-p_{\text{fail}} \times t_{\text{recov}} \times \text{sev}/T_{\text{ref}})$	Expected unavailability: failure probability times recovery time times impact severity, normalised by a reference horizon T_{ref} , then combined by probabilistic OR with the environmental and political exposure terms.
u_{cost}	$\beta_1 \frac{c_{\text{op}} - c_{\text{nom}}}{c_{\text{nom}}} + \beta_2 [\text{tariff} - 1]_+ + \beta_3 \frac{c_{\text{stock}}}{c_{\text{ref}}}$	Financial stress: drift of operating cost above its nominal value, tariff multiplier above neutral, and storage cost relative to a reference cost.
u_{co2}	$\left[\frac{e - e_{\text{target}}}{e_{\text{max}} - e_{\text{target}}}\right]_+$	Normalised overshoot of measured emissions e above the target, saturating at the maximum threshold.

Two conventions make this aggregation unambiguous, since a bare formula does not by itself say what its bounds mean. First, each u_m is engineered to carry the same meaning regardless of which raw KPI feeds it: $u_m \rightarrow 0$ means the block is healthy and contributes no urgency, $u_m \rightarrow 1$ means the block is saturated, that is, it has reached the point past which it cannot make the node any more urgent, whether that point is a missed deadline, a fully depleted buffer, or a near-certain failure. This use of “saturated” is a design choice made for this model, not a term imported as-is from the literature; it is what makes the noisy-OR product

in Equation 4 valid across six structurally different blocks, since the product only requires each factor to share the same $[0, 1]$ urgency scale, not the same physical units or the same statistical behaviour between 0 and 1. Second, a block whose data is `None` is excluded from the product entirely rather than assigned a value: because the product runs only over blocks with available data, a missing block neither raises nor lowers $U_{r,local}$, it is simply absent from the evidence, and the remaining weights are renormalised over what is left. Treating the six blocks as comparable on $[0, 1]$ is therefore a modelling commitment, not a measured property, and Section 5.1 lists it among the framework's stated limitations.

Two lifecycle status rules are applied before aggregation: a node marked `DONE` has its effective local U_r forced to 0.0 (a completed task carries no further physical risk), while a node marked `ABANDONED` has it forced to 1.0 (complete failure). These rules are applied identically whether the local urgency comes from the analytic blocks above or from the Monte Carlo estimate of the time block described below.

Why These Six Blocks: A Rupture-Risk-Driven Selection. The six blocks were not chosen as an arbitrary KPI shortlist. The selection started from the question that defines urgency in a supply chain: *what mechanisms make the risk of a delivery rupture rise?* A rupture, that is the inability to deliver at the agreed time, place, quantity, and quality, is the terminal event the urgency signal must anticipate (Heckmann et al., 2015). Working backwards from that event, and consistently with indicator-based supply risk monitoring (Mokhtar et al., 2019), four *direct causal mechanisms* and two *arbitration variables* were identified:

- **Time erosion (u_{time}):** the remaining slack between deadline and realistic lead time. Time is the only non-compressible resource: whatever the state of the other indicators, a task whose slack reaches zero can no longer be saved, which makes deadline proximity the variable most directly tied to operational urgency (Cohen et al., 2019; Bojko et al., 2021; Sun et al., 2020; Sawik, 2013).
- **System saturation (u_{cap}):** queues, transport capacity, workshop load, and bottlenecks. A saturated system introduces hidden waiting delays on top of the nominal lead time, so capacity stress raises non-delivery risk even when the announced lead time looks safe (Tomlin, 2006; Hosseini; Tajik, et al., 2019; Taghavi et al., 2024).
- **Critical-resource performance (u_{perf}):** drift in the performance of the resource that produces the part. OEE is used as a composite proxy (availability $\times$ performance $\times$ quality) rather than a single machine metric, because precise performance data is rarely accessible on the supplier side (Mokhtar et al., 2019; Hlioui et al., 2017).

- **Adverse-event exposure** (u_{risk}): the expectation of unavailability caused by a disruption, computed as failure probability $\times$ recovery time $\times$ impact severity, composed with environmental and political exposure. The default incident probability (0.02) is revised weekly from declared events through the Bayesian operator of Section 4.3.4 over a $n_0 = 26$ iteration window; this block builds on the DIN supplier risk scorer (Kabouche, 2025) and on probabilistic ripple-effect models (Hosseini; Ivanov; Dolgui, 2020; Liu et al., 2021; Hosseini; Ivanov, 2020).
- **Financial flexibility** (u_{cost}): expediting surcharges, tariff rises, storage costs. Cost drift is not a rupture cause in itself but a stress indicator of the chain (paying premium express transport to avoid a physical rupture is a measurable symptom that the system is straining), and it bounds which corrective actions remain economically viable (Shen et al., 2019; Tomlin, 2006).
- **Carbon overshoot** (u_{co2}): emissions above target. Like cost, it is an arbitration criterion rather than a rupture mechanism: it only becomes decisive when the regulatory or ESG constraint genuinely restricts the set of admissible responses, typically the ship-vs-air trade-off (Hoen et al., 2014; Smart Freight Centre, 2023; International Organization for Standardization, 2023).

Each candidate indicator was screened against four filter questions, summarised in Table 11: only indicators passing all four were retained, which yields a set of four direct causes plus two arbitration variables and explains why the list is six blocks rather than four or eight.

Table 11: Selection filter applied to each candidate KPI block.

Criterion	Question asked of each candidate indicator
Causal link	Does the indicator have a plausible mechanism linking it to delivery rupture?
Observability	Can it be measured or credibly estimated from client, supplier, or planning data?
Actionability	Can an operator act on it, or use it to decide?
Non-redundancy	Does it add information not already carried by a retained indicator?

At this stage the selection remains a *conceptual* construction: it must still be confirmed by the real availability of the underlying data and by sensitivity tests verifying that each block contributes useful information to the aggregate (Section 5.1).

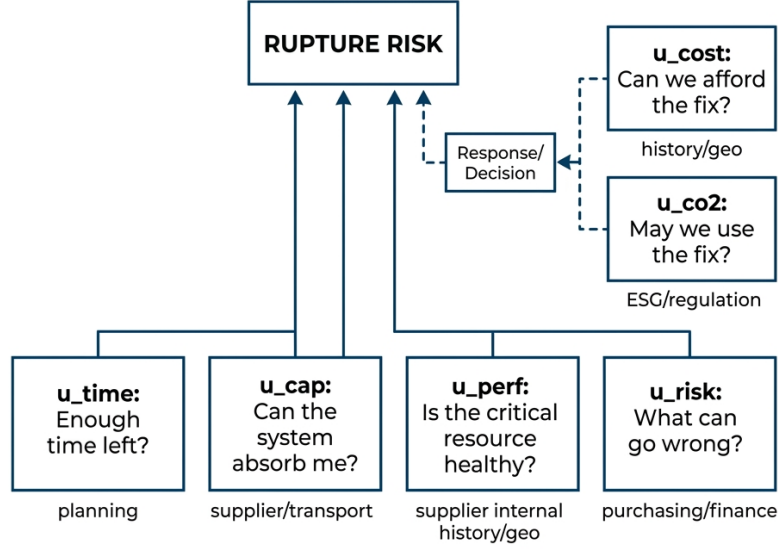


Figure 6: Causal structure behind the six-block selection: four direct rupture mechanisms and two arbitration variables.

Local Declared Urgency ($U_{d,\text{local}}$). Symmetrically to real urgency, declared urgency is not read as a single free-text estimate: it is reconstructed from a structured elicitation, so that it can be scored, propagated, and compared to $U_{r,\text{local}}$ on the same $[0, 1]$ scale. At each node, the operator compares four criteria pairwise through the Analytic Hierarchy Process (AHP), which yields a consistency-checked weight w_j per criterion, then rates each criterion on a severity scale $s_j \in \{1, \dots, 9\}$. Local declared urgency is the weight-averaged severity, rescaled to $[0, 1]$:

$$U_{d,\text{local}} = \sum_{j=1}^4 w_j \frac{s_j - 1}{8}, \quad \sum_{j=1}^4 w_j = 1. \quad (5)$$

The elicitation tool itself, its four criteria, and the consistency gate that rejects an incoherent weight vector before it reaches this formula are application choices, described in Section 4.3.2. What matters at this stage is that $U_{d,\text{local}}$ and $U_{r,\text{local}}$

(Equation 4) are both bounded, node-local quantities in $[0, 1]$ before either one is propagated, since the propagation equations below require both of them already defined.

Network Propagation. Declared urgency U_d propagates *downstream*, from the final customer (rank 0) toward raw-material suppliers, while real urgency U_r propagates *upstream*, from suppliers toward the final customer, using a leaky product rule that keeps both signals bounded in $[0, 1]$ and monotonically increasing as more neighbours become urgent:

$$U_{d,i} = 1 - (1 - U_{d,\text{local},i}) \prod_{k \in \text{Succ}(i)} (1 - \gamma_{ik} U_{d,k}), \quad (6)$$

$$U_{r,i} = 1 - (1 - U_{r,\text{local},i}) \prod_{j \in \text{Pred}(i)} (1 - \beta_{ji} U_{r,j}), \quad (7)$$

where $\gamma_{ik} \in [0, 1]$ is the downstream attenuation coefficient on arc (i, k) and $\beta_{ji} \in [0, 1]$ is the upstream amplification coefficient on arc (j, i) . Backup arcs are purely documentary and inert with respect to both propagations. Table 12 details the propagation variables.

Table 12: Component definitions for the network propagation equations.

Component	Type	Description
$U_{d,i}, U_{r,i}$	Dependent variables	Propagated declared/real urgency at node i
$U_{d,\text{local},i}, U_{r,\text{local},i}$	Input variables	Local urgency at node i after status-rule adjustment
γ_{ik}	System parameter	Downstream (demand) attenuation coefficient on arc $i \rightarrow k$
β_{ji}	System parameter	Upstream (risk) amplification coefficient on arc $j \rightarrow i$
$\text{Succ}(i), \text{Pred}(i)$	Structural sets	Nominal (non-backup) successors and predecessors of node i

Propagation is recomputed incrementally: only the nodes downstream or upstream of a change (a submitted questionnaire, a revised KPI, a structural edit) are recalculated, rather than the full graph, which keeps the digital twin responsive as the network grows. Computation and propagation are triggered per node, not on a single synchronised clock: whenever a node's local urgency changes, whether from a submitted weekly review or a manual KPI edit, the system recomputes

that node's aggregation and propagates the change through its upstream or downstream cone only (Section 4.3.1). Because each node is reviewed on its own weekly cadence rather than all at once, propagation across the network is continuous and asynchronous in practice, not a single global weekly batch; Section 4.3.4 describes the review cadence itself.

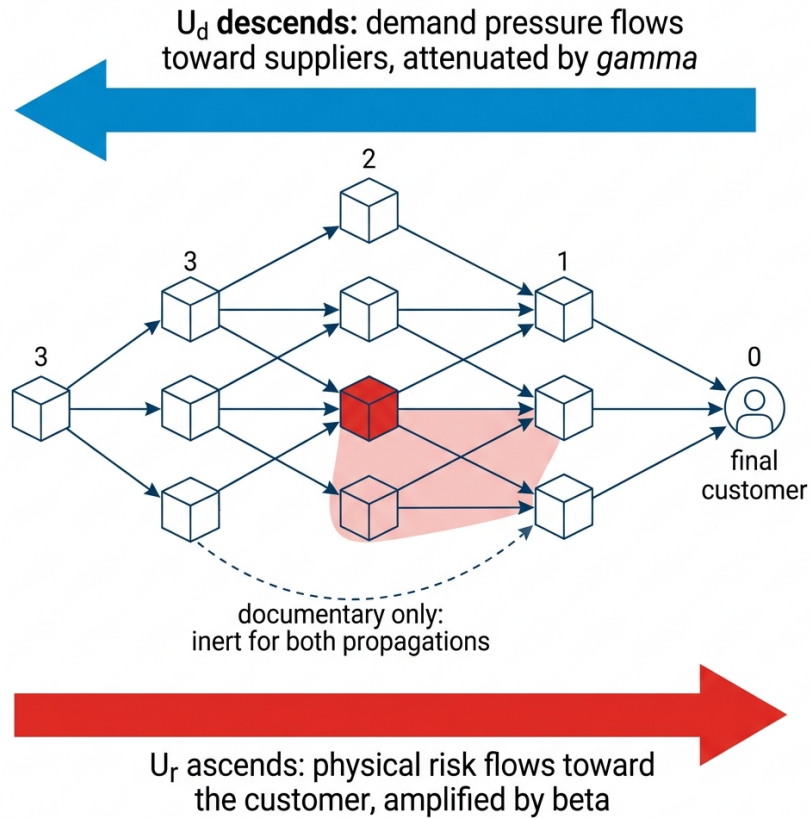


Figure 7: Bidirectional urgency propagation on the supply chain DAG.

4.2.2 Core Model: Asymmetric Cognitive Adequacy (Prospect Theory)

The mismatch between declared and real urgency is split into two named pathologies, false urgency F (over-declaration, panic) and hidden risk H (under-declaration, silent danger):

$$F = \max(U_d - U_r, 0), \quad H = \max(U_r - U_d, 0). \quad (8)$$

Drawing on Kahneman and Tversky's Prospect Theory (Kahneman; Tversky, 1979), the two pathologies are penalised asymmetrically and with diminishing marginal

sensitivity to large errors:

$$\text{penalty} = \lambda_{\text{under}} H^\alpha + \lambda_{\text{over}} F^\alpha, \quad (9)$$

and the penalty is rescaled into an adequacy score $A \in [0, 100]$ through an exponential mapping bounded by the penalty of a maximal one-sided error:

$$A = 100 \times \frac{\exp(-\text{penalty}) - \exp(-\text{penalty}_{\max})}{1 - \exp(-\text{penalty}_{\max})}, \quad \text{penalty}_{\max} = \max(\lambda_{\text{under}}, \lambda_{\text{over}}). \quad (10)$$

Table 13 lists the parameters, with their governing constraint $\lambda_{\text{under}} > \lambda_{\text{over}}$: a coordinator who panics wastes premium transport capacity, but a coordinator who under-reports risk misses the window to prevent a line stoppage, so under-declaration is priced more expensively.

Table 13: Component and parameter definitions for the asymmetric adequacy equations.

Component	Type	Description	Default
F	Dependent variable	False urgency (over-declaration)	n/a
H	Dependent variable	Hidden risk (under-declaration)	n/a
λ_{under}	System parameter	Penalty weight on hidden risk	2.25
λ_{over}	System parameter	Penalty weight on false urgency	1.0
α	System parameter	Psychophysical curvature (diminishing sensitivity)	0.88
A	Dependent variable	Adequacy score, 100 = perfect alignment	[0, 100]

The default $\lambda_{\text{under}}/\lambda_{\text{over}} = 2.25$ ratio reprises the loss-aversion coefficient measured by Kahneman and Tversky (Kahneman; Tversky, 1979; Kahneman, 2011), and $\alpha = 0.88$ reprises their measured curvature of the value function. These defaults carry the status of literature-derived starting points, not measured properties of supply chain coordinators: they have not yet been validated empirically in this context, and their recalibration is one purpose of the validation campaign of Section 4.4. Both the aggregation weights ω_m and the penalty parameters $(\lambda_{\text{under}}, \lambda_{\text{over}}, \alpha)$ are configurable per project rather than hard-coded, so that this recalibration requires no code change.

4.2.3 Monte Carlo Estimate of the Time Block

The analytic u_{time} block assumes a Normal lead time at a single node. On a multi-tier network this assumption breaks: a node can only start when *all* its predecessors

have delivered, and the maximum of several skewed lead times has no convenient closed form. The module `mc/lead_time.py` therefore replaces the analytic estimate with a stochastic one when enabled.

The simulation works in three steps. First, each node draws its lead time L_i from a configurable distribution (lognormal, normal clipped to zero, or triangular). Second, completion times propagate along the DAG in topological order: a node starts when its last predecessor finishes ($S_i = \max_{j \in \text{Pred}(i)} C_j$) and completes after its own lead time ($C_i = S_i + L_i$). Third, the delay probability of each node is estimated as the fraction of draws in which it misses its deadline, $u_{\text{time}} = P(C_i > d_i)$, over $N = 10\,000$ draws by default. A memory guardrail caps the simulation at $N \times n \leq 5 \times 10^7$ across n nodes.

The estimate feeds directly into Equation 4 through a dedicated override point, so the aggregation, status rules, and propagation machinery are reused unchanged.

4.2.4 Network Criticality Index

The criticality service (`services/criticite.py`) answers one question systematically: which node, if it failed completely, would cause the most damage to final-customer delivery? For every active node in turn, the service forces the node's local real urgency to 1.0, re-runs the ascending propagation, and measures two deltas: the rise of U_r at the final-customer nodes ($\Delta U_{r,\text{final}}$) and the largest rise anywhere in the network ($\Delta U_{r,\text{max}}$). Nodes are ranked by $\Delta U_{r,\text{final}}$ first, then $\Delta U_{r,\text{max}}$, then name.

A small example makes the mechanics concrete. Take a three-node chain, raw-material supplier $\rightarrow$ sub-assembly $\rightarrow$ final assembly, with upstream coefficients $\beta = 0.9$ then 0.7 , and a baseline where the sub-assembly node carries $U_r = 0.40$ and the final node $U_r = 0.28$. Forcing the sub-assembly node to $U_r = 1.0$ propagates to the final node as $1 - (1 - 0.7 \times 1.0) = 0.70$, so its shock deltas are $\Delta U_{r,\text{final}} = 0.70 - 0.28 = 0.42$ and $\Delta U_{r,\text{max}} = 1.0 - 0.40 = 0.60$ (Figure 9).

One behaviour of this index matters for interpretation: a node already saturated at $U_r \geq 1.0$ yields a delta of zero and sinks to the bottom of the ranking, since its risk is already realised in the baseline. In a network-wide crisis, where many nodes saturate at once, the index therefore degenerates; this documented behaviour drives the exclusion rule used by hypothesis H3 in the validation campaign (Section 4.4).

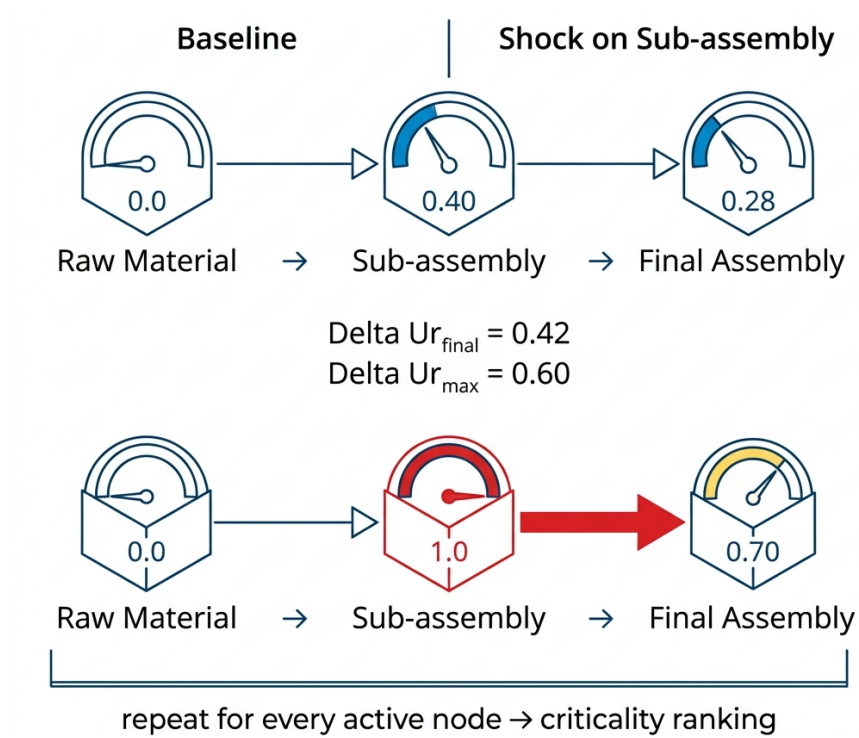


Figure 8: One iteration of the systematic shock simulation behind the criticality index.

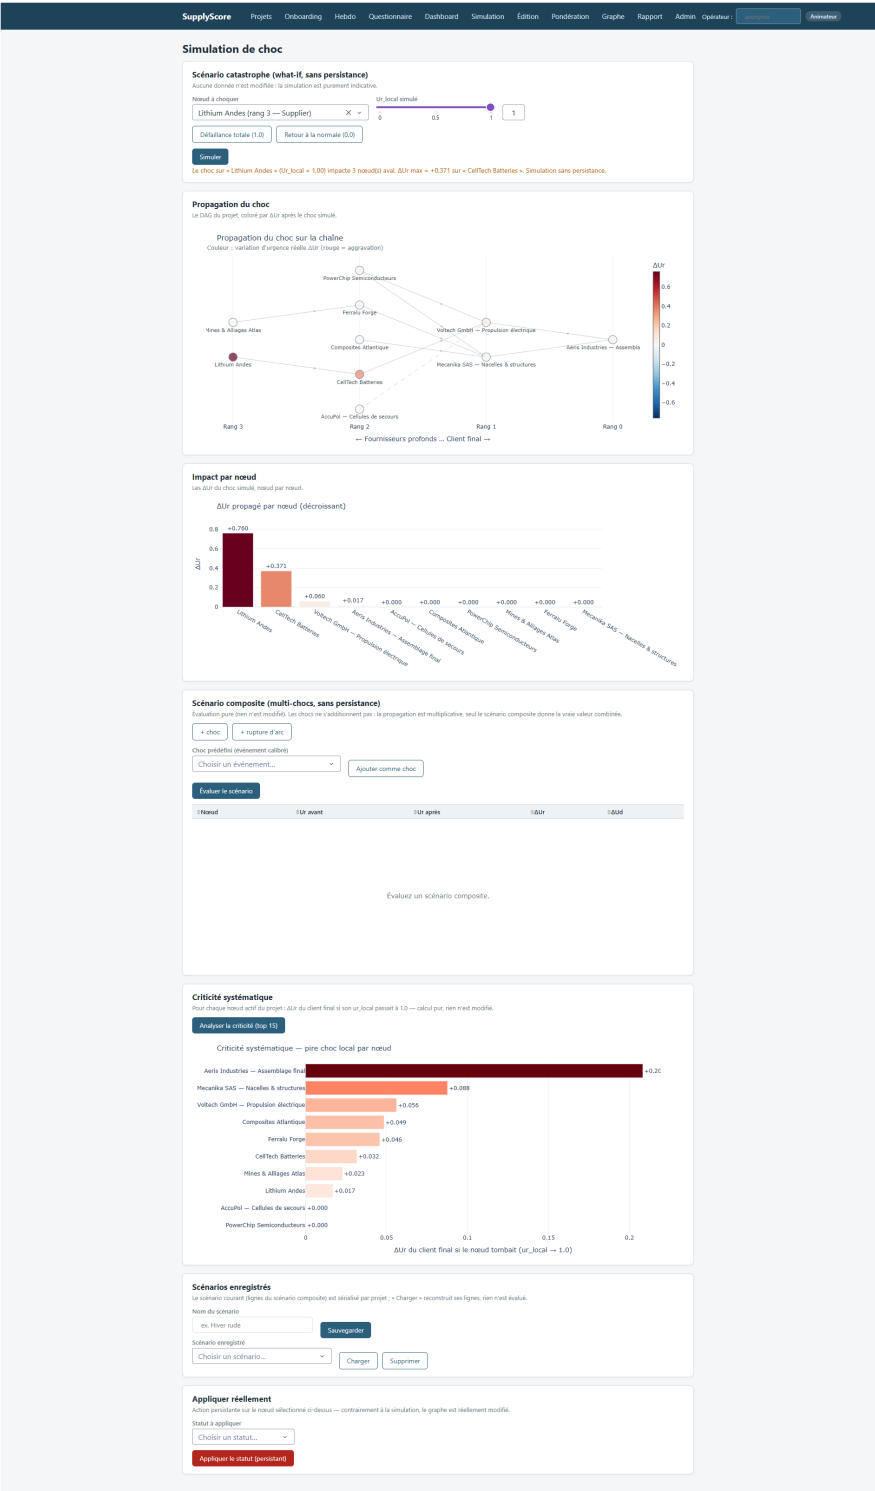


Figure 9: Network criticality ranking after a systematic shock simulation.

4.3 Application and Implementation Choices

The model above is deliberately silent on tooling. This section describes how each mathematical object it defines is actually elicited from an operator, computed inside the running system, exposed back to a coordinator, and kept safe: the system architecture, the AHP and Fuzzy-BWM/PROMETHEE elicitation tools, the weekly operating cycle and event engine that keep KPI data current, and the explainability, export, and hardening capabilities built on top.

4.3.1 System Architecture

Several design guarantees relied upon elsewhere in this document, auditability, multi-tenant isolation, deterministic testing, are architectural choices, not mathematical ones, and are described here on their own.

Single facade. All operations go through one orchestrating facade (`services/orchestrator.py`), which coordinates the mathematical core, the graph repository, and the persistence layers behind a single thread-safe entry point. The web interface, the command-line entry point, and the test suite all drive the system exclusively through this facade; there is no second write path.

Multi-tenant persistence. Data is split between one global registry database (projects, nodes, arcs, milestones, scenarios, registry-level audit) and one SQLite database per node holding its private operational history (KPI snapshots, AHP assessments, weekly reviews, urgency history, declared events, and a per-node audit log). This isolation prevents write contention between operators reviewing different nodes and enforces the boundary that one supplier's raw history is not readable through another supplier's records.

Gatekeeper writes. Every business write crosses a single mutation service that whitelists the touched fields, validates ranges, diffs against the current state (a no-op write leaves no trace), records the change in the audit log with its author and source, and synchronises the in-memory graph. Audit tables are strictly immutable, including from the administration interface: the "no silent data loss" guarantee is enforced here, not by convention.

Graph and propagation. The network lives in an in-memory graph repository (networkx-backed; a Neo4j adapter is scaffolded but out of scope). Propagation is incremental: only the upstream/downstream cones of nodes whose local urgency

actually changed are recomputed, which keeps the twin responsive at the 1000+-node scale validated by dedicated performance benchmarks.

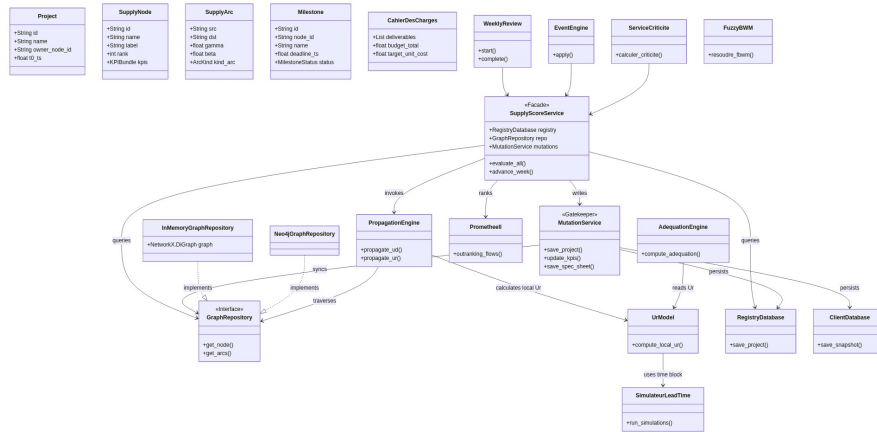


Figure 10: Doxygen-generated UML view of the SupplyScore architecture: single facade, gated writes, multi-tenant persistence.

Decoupled clocks. Each project carries its own clock: a system clock for live use, a game clock advancing week by week for serious-game sessions, and a fixed clock for deterministic tests. The weekly cycle, milestone pressure, and event timestamps all read the project clock, never the wall clock directly, which is what makes a compressed serious-game campaign (Section 4.4) possible without touching the model.

4.3.2 Expert Criteria Weighting: The AHP Elicitation Model

The Python module `core/ahp.py` implements the Saaty eigenvector method for a 4-criterion questionnaire administered through the weekly review page, and is the tool that produces $U_{d,local}$ (Equation 5). The four criteria are:

- C1. Operational impact:** “If this task fails, what is the downstream consequence?”
- C2. Time window:** “How long before the situation becomes irredeemable?”
- C3. Downstream dependencies:** “How many tasks are blocked waiting for this one?”
- C4. Recoverability:** “Can we catch up if this task is late?” (inverted, where lower is worse)

How the Four Criteria Were Chosen. The criteria set follows from two constraints, one cognitive and one structural. Cognitively, Saaty's method degrades as the number of pairwise-compared criteria grows: consistency becomes hard to maintain

beyond roughly seven items, and each additional criterion adds $n - 1$ comparisons to the operator’s weekly workload (Saaty, 1990; Triantaphyllou; Yanase, 2024). Keeping $n = 4$ limits the questionnaire to six comparisons, short enough for a weekly cadence. Structurally, the four questions mirror the way declared urgency is consumed downstream: C1 (impact) and C3 (dependencies) capture the *severity* dimension that propagation amplifies across the network, C2 (time window) captures the *slack* dimension confronted with the u_{time} block of Section 4.2, and C4 (recoverability) captures the *resilience* dimension confronted with the recovery term of u_{risk} . Each declared criterion therefore has an operational counterpart in the computed signal, which is what makes the later U_d vs. U_r comparison meaningful rather than a comparison of unrelated quantities. Larger criteria sets (cost, sustainability) are deliberately routed to the F-BWM weighting of the KPI blocks instead (Section 4.3.3), where the elicitation burden is lower.

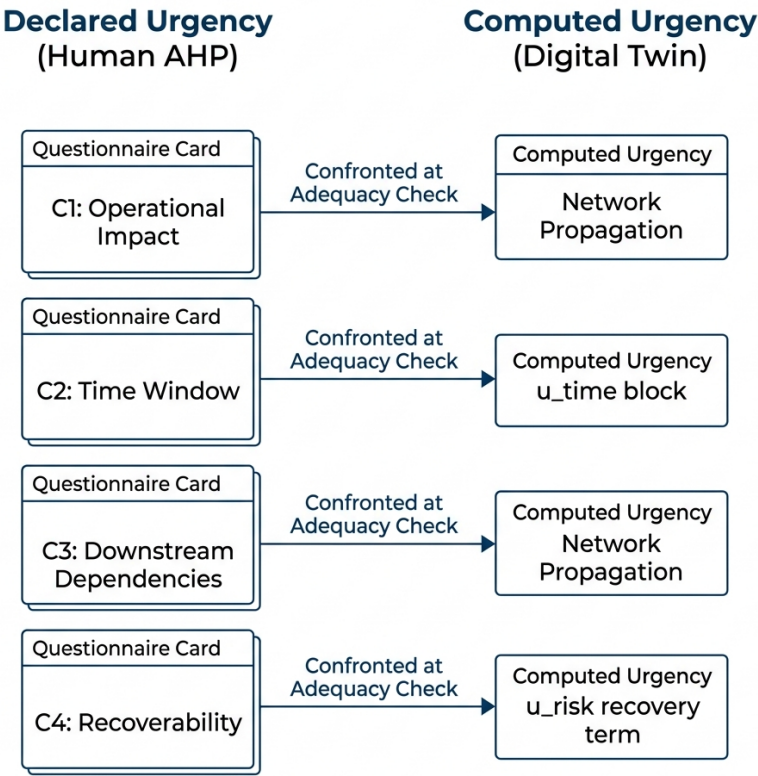


Figure 11: Correspondence between the four AHP criteria and their computed counterparts.

The consistency ratio is computed as:

$$CR = \frac{ConsistencyIndex(CI)}{RandomIndex(Saaty)n(RI)}, \quad (11)$$

where the components of the consistency assessment are defined in Table 14.

Table 14: Component definitions for the AHP consistency ratio equation.

Component	Type	Description
CR	Dependent variable	Consistency Ratio of the pairwise comparison matrix
CI	Dynamic variable	Consistency Index computed from the matrix principal eigenvalue
RIn	System parameter	Random Index for a matrix of size n (from Saaty's reference table)

using Saaty's Random Index table. If $CR > 0.10$, the interface flags the questionnaire as inconsistent and invites the operator to revise the most contradictory pairwise comparison before the assessment is accepted. Declared urgency is further smoothed with an exponential moving average ($\rho = 0.3$) across consecutive weekly assessments, absorbing transient stress spikes without erasing a sustained change in perception.

Figure 12: AHP questionnaire page with the consistency-ratio validation gate.

AHP is viable for U_d encoding in this context. The CR validation gate, exercised by both worked examples and automated tests, detects and rejects cognitively inconsistent expert inputs before they enter the network. Its behaviour with real operators is measured in the campaign of Section 4.4.

4.3.3 Stochastic Judgment Modeling: Fuzzy BWM and PROMETHEE II Network Outranking

The modules `mcda/fbwm.py` and `mcda/promethee.py` implement two complementary methods: F-BWM offers a lower-effort elicitation route than crisp AHP for weighting the six KPI blocks, and PROMETHEE II turns those weights into a network-wide, non-compensatory ranking of nodes that complements the per-node adequacy score.

F-BWM Implementation. The linguistic TFN scale maps five qualitative judgements to triangular fuzzy numbers on a 1–4.5 scale (from *equally important* to *absolutely more important*), each carrying its own maximum consistency index. The consistency ratio of the Fuzzy BWM model is formulated as:

$$CR_{\text{BWM}} = \frac{\xi^*}{\text{CI}}, \quad (12)$$

where ξ^* is the optimal consistency indicator returned by a constrained non-linear optimisation that fits the fuzzy weights to the elicited Best-to-Others and Others-to-Worst comparisons, and CI is read from the strongest linguistic judgement used. If the optimisation fails to converge, the implementation falls back to uniform weights rather than raising an error, so that a single malformed elicitation cannot crash the weekly review. Fuzzy weights are defuzzified with the Graded Mean Integration Representation (GMIR):

$$w_j = \frac{a_j + 4b_j + c_j}{6}, \quad (13)$$

then renormalised so that $\sum_j w_j = 1$. Table 15 defines the terms; the triangular fuzzy bounds are written (a_j, b_j, c_j) here, rather than the more common (l_j, m_j, u_j) notation, to avoid colliding with the risk blocks u_m used throughout Section 4.2.

Table 15: Component definitions for the Fuzzy BWM defuzzification.

Component	Type	Description
w_j	Dependent variable	Crisp normalized weight of criterion j
a_j, b_j, c_j	Dynamic variables	Lower, modal, and upper bounds of the triangular fuzzy weight
ξ^*	Dynamic variable	Optimal consistency indicator from the fitted optimisation

PROMETHEE II Implementation. PROMETHEE II is applied in production to rank the active nodes of a project on the same six weighted urgency blocks (Table 10) that feed local real urgency, giving coordinators a network-wide, non-compensatory view of which nodes are collectively driving urgency. For each pair of active nodes (a, b) and each block m , the preference intensity $P_m(g_m(a) - g_m(b)) \in [0, 1]$ is computed with a Type V linear preference function with an indifference threshold; the aggregated preference index and net outranking flow follow the standard PROMETHEE II formulation:

$$\pi(a, b) = \sum_m \omega_m P_m(g_m(a) - g_m(b)), \quad \phi(a) = \phi^+(a) - \phi^-(a), \quad (14)$$

where $\phi^+(a)$ and $\phi^-(a)$ are the average preference of a over, respectively by, all other active nodes. Table 16 defines the outranking flow terms.

Table 16: Component definitions for the PROMETHEE II net flow.

Component	Type	Description
$\pi(a, b)$	Dependent variable	Aggregated preference index of a over b
$\phi(a)$	Dependent variable	Net outranking flow representing overall preference ($\in [-1, 1]$)
$\phi^+(a), \phi^-(a)$	Dynamic variables	Positive (dominance) and negative (dominated) outranking flows

The same weights ω_m used in the local urgency aggregation (Equation 4) are reused as the PROMETHEE criteria weights, so that the network ranking and the per-node urgency score remain mutually consistent rather than relying on a separately calibrated weight set. Both modules are unit tested and wired into the production ranking path (`services/orchestrator.py`).

4.3.4 Weekly Review Cycle and Calibrated Event Engine

The weekly review cycle is what keeps the digital twin’s KPI state current without free-form manual data entry: every write is structured, signed, and auditable. This section describes the cycle itself, then the calibrated event engine that translates real-world disruptions into KPI updates.

The Weekly Cycle. The system is operated as a serious game advancing in discrete weeks (`core/clock.py` implements `SystemClock` for live use, `GameClock` for week-by-week advancement, and `FixedClock` for deterministic testing). At each new week, every active node is flagged as due for review, and the operator

completes four sequential sections (`services/weekly.py`, served by the `/hebdo` page, Figure 14):

1. the AHP questionnaire, pre-filled with the previous week's values, with a one-click "unchanged" confirmation;
2. the week's KPI values, shown against the previous week for quick differencing;
3. progress on active milestones, compared against theoretical planning progress;
4. any supply chain events that occurred during the week.

All four sections are written as a single signed, atomic transaction, so that no partial or anonymous review can corrupt the node's history.

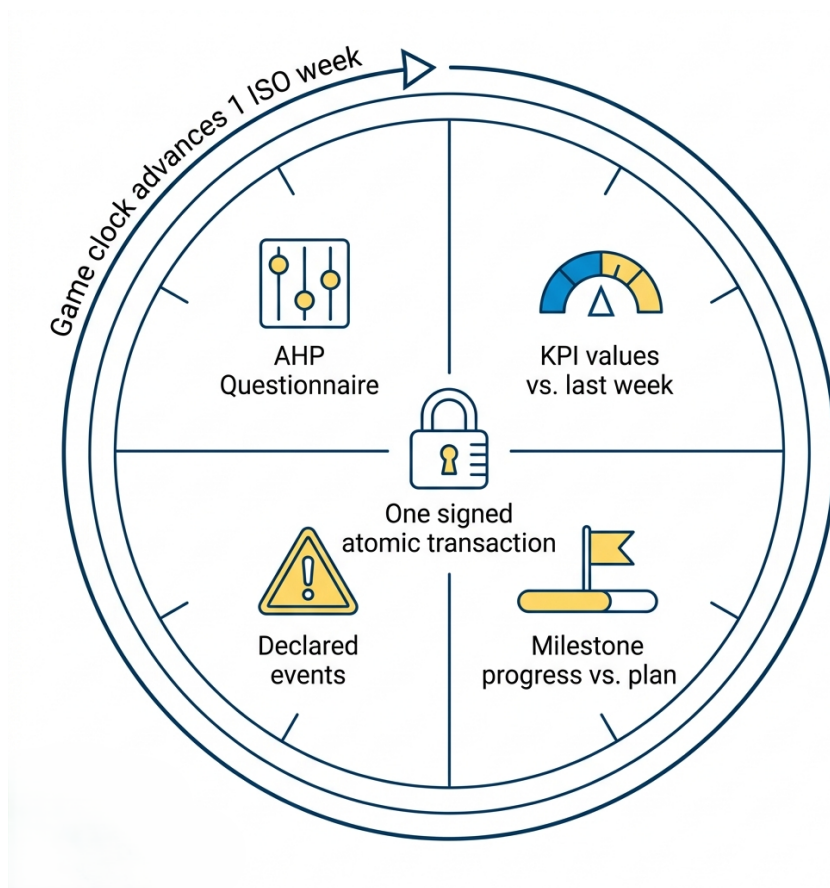


Figure 13: The four-section weekly review cycle committed as a single signed transaction.

Figure 14: Weekly review page: the four sequential sections (AHP, KPIs, milestones, events).

Calibrated Event Engine. Rather than asking operators to manually recompute the KPI impact of a disruption, the system exposes a taxonomy of fourteen calibrated event types (`domain/events.py`), listed in Table 17, each mapped to one or more KPI updates through one of three calibration operators.

Table 17: The fourteen calibrated event types.

panne_machine	retard_fournisseur	greve
hausse_tarif	rupture_matiere	accident
non_conformite_qualite	perturbation_transport	instabilite_politique
cyber_incident	hausse_energie	alerte_financiere_
perte_capacite	pic_demande	fournisseur

The three calibration operators are: a Beta-Bernoulli Bayesian update for failure probabilities,

$$p_{\text{post}} = \text{clip}\left(\frac{p \times n_0 + k}{n_0 + n}, 10^{-4}, 0.99\right), \quad (15)$$

with a prior weight $n_0 = 26$ (roughly six months of weekly observations); an exponential moving average for continuous variables such as OEE, $x_{\text{new}} = (1 - \lambda)x_{\text{old}} + \lambda x_{\text{obs}}$ with $\lambda \in [0.3, 0.5]$; and a ratchet (max) update for severity, $s_{\text{new}} = \max(s_{\text{old}}, s_{\text{event}})$. Table 18 defines the shared terms.

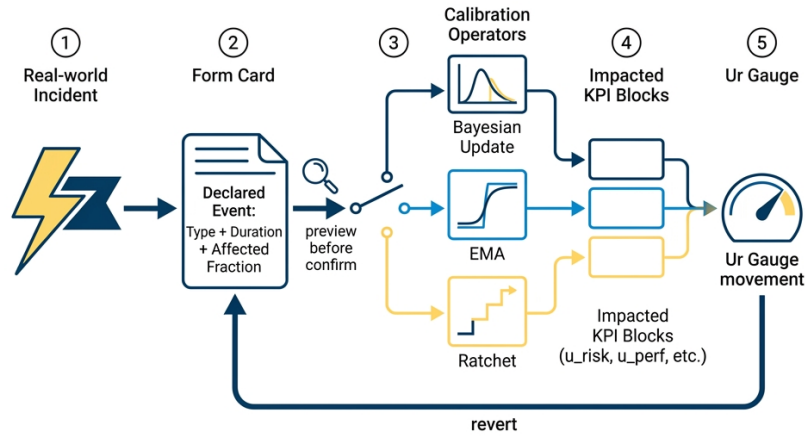


Figure 15: From a declared event to a calibrated KPI update: the three calibration operators.

Table 18: Component definitions for the event calibration operators.

Component	Type	Description
p_{post}	Dependent variable	Revised failure probability after the event
n_0	System parameter	Prior equivalent sample size (26, roughly six months)
(k, n)	Event input	Equivalent failures/trials implied by the declared event's severity
x_{new}	Dependent variable	Revised continuous KPI (e.g., availability)
λ	System parameter	EMA reactivity weight, $\in [0.3, 0.5]$

Numerical Example: Strike Event. A node declares a greve (strike) lasting $d = 24$ hours affecting a fraction $p_f = 0.50$ of the workforce. Equivalent lost hours

over a 168-hour week: $H_{\text{lost}} = d \times p_f = 12$, so $f_{\text{lost}} = 12/168 \approx 0.0714$ and the observed availability is $A_{\text{obs}} = 1 - 0.0714 \approx 0.9286$. Updating a baseline availability $A_{\text{old}} = 0.95$ with EMA reactivity $\lambda = 0.5$: $A_{\text{new}} = 0.5(0.95) + 0.5(0.9286) \approx 0.9393$. Risk severity is updated in ratchet mode: $S_{\text{new}} = \max(0.30, 0.50) = 0.50$. Operators can preview the transient impact of a declared event before confirming it, and reverting an event recomputes the baseline KPIs automatically.

The weekly cycle and event engine let coordinators keep the digital twin's KPI state current from structured, auditable inputs rather than free-form edits. This is what makes the resulting urgency history usable as a calibration dataset for the campaign of Section 4.4.

4.3.5 Explainability, Export, and Operational Hardening

When the digital twin flags a node, the coordinator's first question is *why*. The last three capabilities of the demonstrator exist to answer that question exactly, to let the data leave the tool, and to make sure nothing is ever lost.

Explainability. The per-node explanation page (`services/explain.py`) answers "why this score?" with exact decompositions, not approximations. Declared urgency splits into one additive contribution per AHP criterion, the same construction as Equation 5 read term by term:

$$\kappa_j = w_j \frac{s_j - 1}{8}, \quad \sum_j \kappa_j = U_{d,\text{local}}, \quad (16)$$

so the page can state, for instance, that half of a node's declared urgency comes from the time-window criterion. Local real urgency splits the same way through the logarithm of Equation 4: each block contributes $\ell_m = -\omega_m \ln(1 - u_m)$, the block shares $c_m = \ell_m / \sum_k \ell_k$ sum to one, and $U_{r,\text{local}} = 1 - \exp(-\sum_m \ell_m)$ reconstructs the score exactly. A coordinator can therefore see which block drives a node's urgency, and whether the gap between U_d and U_r is panic or hidden danger (Figure 16).

SupplyScore Projets Onboarding Hebdo Questionnaire Dashboard Simulation Edition Pondération Graphe Rapport Admin Opérateur [Accéder](#) [Ajouter](#)

Pondération des critères d'urgence (FBWM)

À quoi servent ces poids ?
FBWM pondère les 6 BLOCS KPI de l'urgence réelle Ur — l'ANP du questionnaire reste dédié aux 4 critères du besoin déclaré Ud.
 Seulement 2n-3 comparaisons (2 pour 6 blocs) suffisent, à la fois une matrice complète en respectant 15.
 Poids actuels du projet : Temps 0.285 - Capacité 0.091 - Performance (DEI) 0.154 - Risque 0.229 - Coût 0.164 - CO2 0.075.

1. Meilleur et Pire critère
 Le FBWM associe toutes les comparaisons sur ces deux extrêmes.
 Critère le PLUS important Critère le MOINS important

2. Jugements linguistiques (2n-3)
 Choisissez d'abord le critère le PLUS important et le MOINS important.
 0/9 jugements validés

3. Aperçu des poids résolus
 L'aperçu se recalcule à chaque jugement modifié — rien n'est persisté.
 Choisissez deux critères distincts (Meilleur et Pire) puis les 2n-3 jugements.

4. Enregistrement (projet actif)
 L'enregistrement sauvegarde les jugements validés (CR < 0.10), renvoie les poids à la semaine ISO du projet et réévalue tous les scores.

Figure 16: Criteria weighting interface (Fuzzy-BWM) feeding the explainability decomposition.

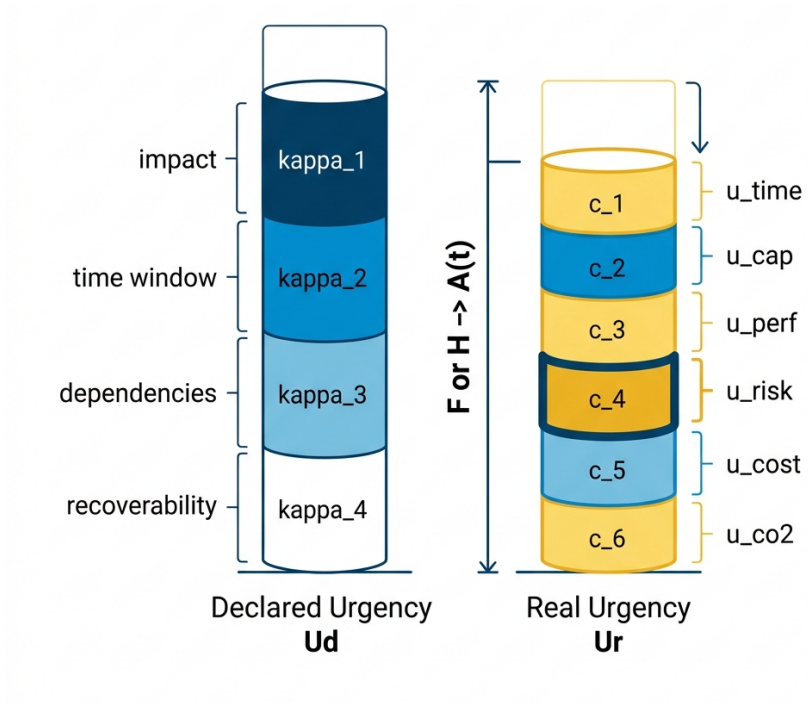


Figure 17: Per-node explainability: exact additive decomposition of U_d and U_r .

Export. One action compiles the global registry and the per-node history into a twelve-worksheet Excel workbook (`services/exports.py`): project metadata, nodes, arcs, milestones, KPI history, urgency history, assessments, weekly reviews, KPI snapshots, events, decisions, and two audit-log sheets. All timestamps are converted to the coordinator’s local timezone. Analyses can therefore run outside the tool, on frozen data, which the validation campaign relies on.

Operational Hardening. The persistence layer is SQLite, so data safety is handled explicitly (`data/backup.py`): Write-Ahead Logging is enabled, an integrity check runs before every backup, and the native `sqlite3.backup` API copies each database without locking the running application. Automatic backups run every 24 hours with a retention cap of 20 archives, and schema migrations are versioned (Figure 18).

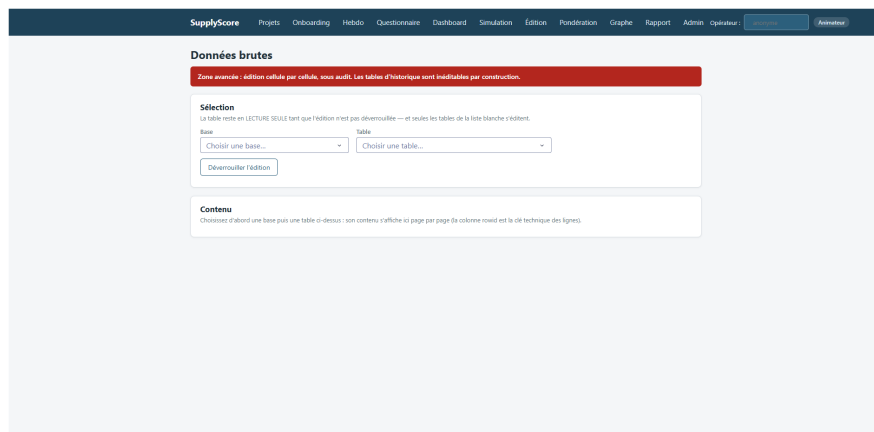


Figure 18: Administration page: database sizes, migration version, and backup controls.

These three capabilities are not peripheral. They are what make the adequacy score auditable and the digital twin operable outside a development environment, and they are preconditions for the campaign below.

4.4 Use Case and Validation: The HÉLIOS Serious-Game Campaign

Everything up to this point establishes that the model is implemented, internally consistent, and tested. What it does not establish is whether the adequacy signal says anything true about real operators facing a real crisis. This section presents the validation campaign built for that purpose: its design, its current state of readiness, and what has been learned from the two checks already run against it.

The campaign replays the 2020–2022 semiconductor shortage as a serious game, designed, calibrated through five dry runs, and pre-registered before execution.

Campaign Design. The campaign, presented to participants under the cover story *Programme HÉLIOS*, replays the semiconductor crisis over 18 weekly game turns compressed into nine working days (two turns per day, morning and late afternoon). The supply network has eight nodes, one per consultant, anchored on the real actor types of the crisis (Table 19).

Table 19: The eight HÉLIOS campaign nodes and their real-world anchors.

Node	Role	Rank	Real-world anchor
Orbitalys	Satellite prime, final customer	0	Fictional satellite programme, contractual delivery pressure
AvioSys Intégration	Avionics integrator	1	Qualified-stock squeeze on avionics tier-1 suppliers
Électis EMS	Electronics manufacturing services	2	EMS plants placed under allocation, 2021
CompoDis Europe	Component distributor	3	Distributor allocation and spot-market premiums
TransGlobal Fret	Freight operator	3	Suez Canal blockage, West Coast port congestion
NovaFab Semiconductors	Foundry, expected bottleneck	4	Composite of the Texas winter-storm freeze and the Taiwan drought
Meridian Semi	Secondary foundry, inert backup arc	4	A foundry sold out through 2023: the announced "plan B" that delivers nothing
SilPure Materials	Raw wafer supplier	5	Roughly 60% of world wafer supply; polysilicon price roughly tripling in 2021

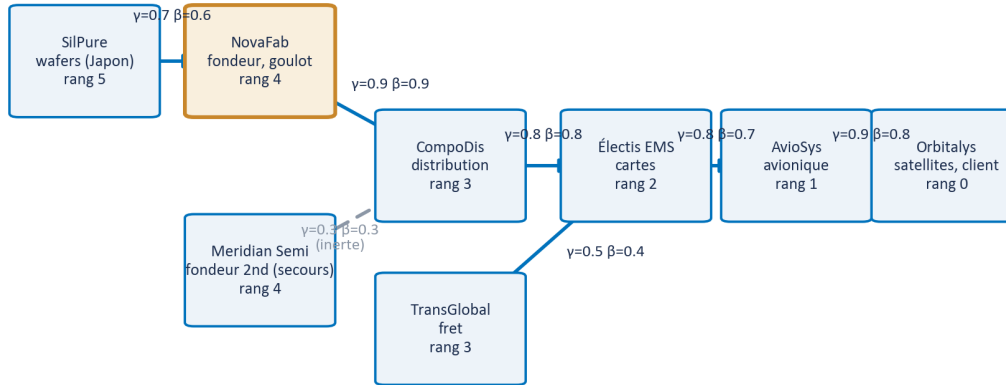


Figure 19: The HÉLIOS network with its per-arc coefficients: γ (downstream demand attenuation) and β (upstream risk amplification), and each node's tier rank. NovaFab (rank 4) is the expected bottleneck; the dashed Meridian arc is the inert backup arc of Section 4.2.

The eighteen turns follow five acts drawn from the documented chronology of the real crisis: a baseline where early warning signs go unheeded (turns 1–4), a tipping point where a foundry freeze, a competitor fire, and a canal blockage make the crisis visible (turns 5–8), a squeeze where spot prices reach ten to forty times catalogue and no node keeps any margin (turns 9–12), a rupture where a major manufacturer cuts production by roughly 40%, the bullwhip effect takes hold, and the announced backup foundry turns out to have no spare capacity (turns 13–16), and a slow, partial recovery that measures whether declared urgency lags the physical de-escalation (turns 17–18).

Each act is driven by calibrated engine events whose type, parameters, and timing are taken from that chronology; the twelve events are listed in Table 20. One game turn represents one real month, from September 2020 to February 2022, with real durations converted into engine hours at a fixed rate (one real week maps to roughly 38.7 engine hours), so that documented figures such as the twelve-to-twenty-five-week supplier lead times or the six-day canal blockage keep their true magnitude inside an eighteen-turn game.

Table 20: The twelve calibrated engine events that drive the HÉLIOS crisis, with their real-world anchors.

Turn	Node	Event type	Real-world anchor
T3	NovaFab	Demand spike	Automotive reorders on sold-out slots, Nov. 2020
T5	CompoDis	Supplier delay	Distributor moves to allocation, Jan. 2021
T6	NovaFab	Accident (critical)	Texas winter-storm fab freeze, Feb. 2021
T7	NovaFab	Demand spike	Reorders after a competitor’s fab fire, Mar. 2021
T7	TransGlobal	Transport disruption	Suez Canal blockage, six days, Mar. 2021
T10	NovaFab	Material shortage	Taiwan drought, ultrapure-water rationing, Jun. 2021
T12	Électis	Capacity loss	Southeast-Asia backend lockdowns, Aug. 2021
T12	TransGlobal	Transport disruption	Shift to saturated air freight, Aug. 2021
T13	CompoDis	Demand spike	Bullwhip over-ordering, Sep. 2021
T14	TransGlobal	Transport disruption	West Coast port congestion, Oct. 2021
T14	CompoDis	Tariff increase	Allocation spot premiums, Oct. 2021
T16	SilPure	Tariff increase	Polysilicon price pass-through, Dec. 2021

One encoding choice is worth noting. The Texas freeze is modelled as a production accident, which triggers the engine’s Bayesian revision of the failure probability, a recovery clock, and a severity ratchet, rather than as a capacity loss, which would only reduce storage volume: a hot shutdown of a fabrication line is mechanically closer to the former.

The two urgency signals come from different worlds, which is the point of the design. Real urgency U_r is fed from public time series of the actual crisis (INSEE production and price indices, WSTS worldwide semiconductor billings, and a Kaggle semiconductor-shortage dataset), frozen in a data pack whose SHA-256 hashes were recorded before the first turn. Every third-party dataset is passed through an authenticity check before use: a Kaggle logistics dataset initially considered for the freight node was rejected at this check for the signatures of synthetic training data (pre-computed target columns, GPS coordinates scattered outside the region the dataset claimed to cover, uniformly flat distributions), while the semiconductor-shortage dataset passed (a secular price decline, a genuine 2021 upturn, sectoral

employment consistent with official series) and was kept. The rejected dataset is retained in the raw-data folder for audit but is never loaded by the analysis pipeline; the freight node's cost block is instead driven entirely by the calibrated events of Section 4.3.4. Declared urgency U_d comes from the consultants through the same weekly AHP questionnaire described in Section 4.3.2, under ten minutes per turn, with the consistency gate active.

Two data-provenance rules protect the real-urgency series. First, a single-pilot rule forbids double counting: each KPI field of each node is driven either by a time series or by a calibrated event, never by both, so when an event writes a field on a given turn the series skips that turn. Second, one honest artefact of the real data is kept rather than smoothed away. The INSEE business-failures series falls through 2020 and 2021, under public support schemes, so a failure-rate indicator taken alone would lower external risk exactly while the sector crisis was worsening. Keeping it is the multi-block argument made concrete: a single macroeconomic indicator can tell the opposite of the ground truth, which is why real urgency aggregates six blocks rather than trusting one.

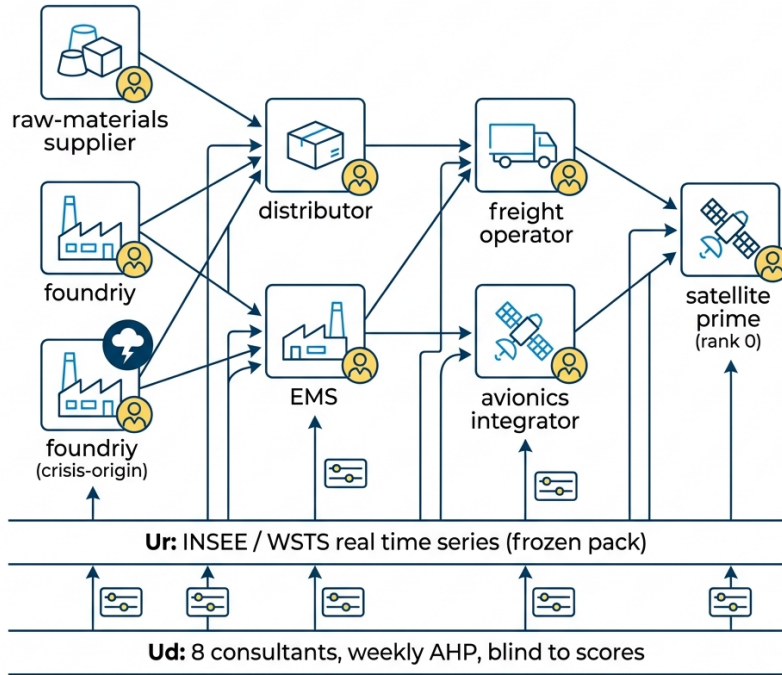


Figure 20: The eight-node HÉLIOS campaign network. The person icon on each node marks the consultant who plays it, feeding U_d through the weekly AHP questionnaire; the lower band shows the public data series that independently feed U_r .

Protocol Integrity. The experimental protocol was frozen before the first real turn. Consultants see only their own role sheet and the questionnaire page; dashboards and scores stay hidden until the debrief. No discussion between roles is allowed during the campaign. Two placebo sheets, a false de-escalation handed to the avionics integrator at turn 9 and to the wafer supplier at turn 15, control for suggestibility: if a consultant's declared urgency follows the placebo, the measure is reflecting the sheet handed to them rather than a memory of the real crisis. A weak-signal turn (turn 4), where the press mentions a tension with no engine event behind it, measures who reacts to media noise alone. Every deviation from protocol is logged the same day in a campaign journal. Two further safeguards protect internal validity: roles are assigned so that the two most information-heavy nodes, the foundry and the distributor, do not go to the most supply-chain-experienced consultants, keeping a role effect from being confounded with an expertise effect, and at the debrief each consultant is asked at which turn they recognised the real crisis, which turns the main threat to a blind declaration into a

measured quantity rather than an uncontrolled one. The hypothesis register is the one of Section 2.4: H1 to H6, each with a numeric threshold fixed in advance, so that success and failure are defined before the data exists. Any analysis outside that register is labelled exploratory.

Empirical Calibration of the Adequacy Signal. The campaign feeds the calibration service (`services/calibration.py`), which addresses Challenge 3 (Section 2.5) and hypotheses H2 and H4: does the hidden-risk signal H have predictive value beyond a raw KPI threshold? For each recorded (node, week) pair, the historical hidden risk is compared against negative outcomes observed in a subsequent four-week window (a missed milestone, an abandoned node, or a critical event). The resulting confusion matrix yields precision, recall, and F1:

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}, \quad F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}. \quad (17)$$

Mechanical Validation: The Synthetic Dry Run. Before any human data exists, the full campaign mechanics were exercised end to end on a disposable database: five iterations (v1 to v5) calibrated the scenario and eliminated engine artefacts, and the final iteration replayed all nineteen turns (T0 to T18) with synthetic respondents standing in for the eight consultants. This dry run is an engineering check, not a scientific one: the synthetic declared urgency is anchored directly on each node's own local real urgency ($1 + 8 \times U_{r,\text{local}}$, perturbed by a fixed panic, under-report, or neutral bias per role), so it cannot lag its own physical signal, cannot under-report it for long, and cannot over-react to a press event it does not read. Run against the pre-registered tests, the dry run therefore returns not confirmed or not applicable on all six hypotheses by construction (Table 21), which is the expected and informative outcome: it confirms that the full pipeline, injection, weekly decay, milestone drift, event calibration, propagation, criticality, snapshotting, and backup, runs correctly under the two-turn-per-day cadence the real campaign requires, and it calibrated the null-baseline detector (a raw threshold $U_{r,\text{local}} \geq 0.5$) at precision 0.16 and recall 0.27: the bar the hidden-risk signal must clear in the real campaign to demonstrate added value beyond a raw KPI reading (Challenge 3, Section 2.5).

Table 21: Dry-run v5 verdicts on synthetic data: an engineering check, not a hypothesis test.

Test	Verdict	Measured value
H1 latency	Not confirmed	1 of 8 nodes at lag ≥ 1 turn (threshold: 5 of 8)
H2 early detection	Not confirmed	$H(\text{NovaFab})$ peaks at 0.070 on turns 6–12 (threshold: 0.3)
H3 criticality	Not confirmed	NovaFab top-1 on 7 of 10 informative turns, 70% (threshold: 80%; saturated turns excluded, Section 4.2.4)
H4 predictive validity	Not confirmed	precision 0.00, recall 0.00 at $H \geq 0.5$ (threshold: 0.5; null baseline: precision 0.16, recall 0.27)
H5 contagion	Not confirmed	non-affected nodes: $\Delta U_d = -0.007$ on press turns vs. $+0.031$ on calm turns
H6 hysteresis	Not applicable	real urgency shows no measurable decline over the recovery turns, so the comparison is not meaningful

Figure 21 shows why: for every one of the eight nodes, real urgency (solid) and synthetic declared urgency (dashed) track each other closely across the full replay, rather than one lagging or decoupling from the other. That closeness is the visual signature of the circularity described above, not a property the real campaign is expected to reproduce. Table 22 traces the same propagated real urgency at eight representative turns, showing the crisis propagate through the network broadly as the scenario intends: a slow build in the baseline, the sharp T6–T8 tipping point at NovaFab, TransGlobal, and CompoDis, a saturated network through the squeeze and rupture, and a partial, uneven recovery where NovaFab barely eases while Meridian, hit late by the worldwide demand wave, saturates only at the very end.

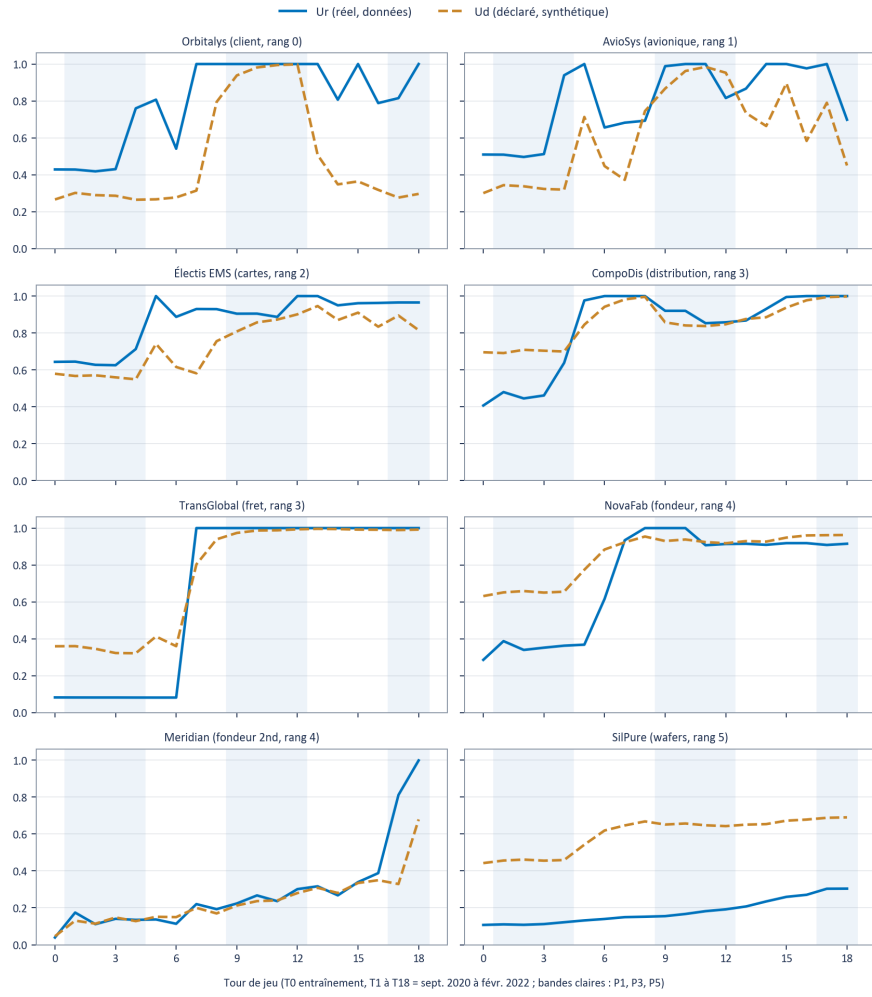


Figure 21: Real urgency U_r (solid) against synthetic declared urgency U_d (dashed) for all eight nodes across the dry-run v5 replay (T0 to T18, shaded bands mark phases P1, P3, P5). The synthetic U_d follows U_r at every node by construction, which is why H1, H2, H5, and H6 are not interpretable on this run.

Table 22: Propagated real urgency U_r by node across the dry-run v5 replay (selected turns).

Node	T0	T4	T6	T8	T10	T13	T16	T18
Orbitalys	0.43	0.76	0.54	1.00	1.00	1.00	0.79	1.00
AvioSys	0.51	0.94	0.66	0.69	1.00	0.87	0.98	0.70
Électis	0.64	0.71	0.89	0.93	0.90	1.00	0.96	0.97
CompoDis	0.41	0.64	1.00	1.00	0.92	0.87	1.00	1.00
TransGlobal	0.08	0.08	0.08	1.00	1.00	1.00	1.00	1.00
NovaFab	0.29	0.36	0.62	1.00	1.00	0.92	0.92	0.92
Meridian	0.04	0.13	0.11	0.19	0.27	0.32	0.39	1.00
SilPure	0.11	0.12	0.14	0.15	0.17	0.21	0.27	0.30

The confusion matrix behind the H4 row of Table 21 is given in full in Table 23, since it is what the real campaign’s confusion matrix will be measured against.

Table 23: Calibration confusion matrix on the dry-run v5 synthetic replay (152 node-turn observations).

Quantity	Value
True positives / false positives (hidden-risk signal, $H \geq 0.5$)	0 / 5
True negatives / false negatives (hidden-risk signal, $H \geq 0.5$)	102 / 37
Precision / recall, hidden-risk signal	0.00 / 0.00
Precision / recall, null baseline ($U_{r,local} \geq 0.5$)	0.16 / 0.27

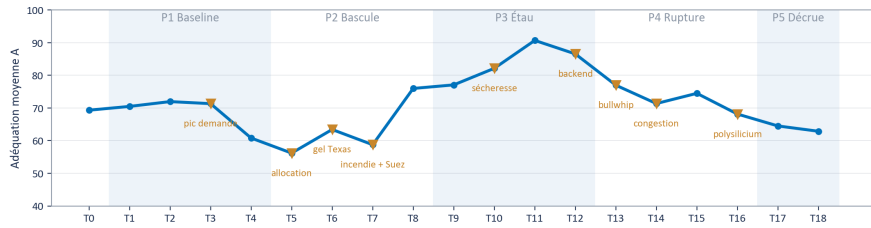


Figure 22: Network-average adequacy score A across the dry-run v5 replay, annotated with the scenario’s calibrated events. The score falls sharply at fast turning points (demand spike, Texas freeze, fire and Suez, drought), then rises through the squeeze (P3) as the saturated network is uniformly tracked by the synthetic U_d , before eroding again as the bullwhip and congestion events land.

Exploratory Pilot: Eight Isolated Language-Model Agents. A second, more informative pilot replaced the fixed-bias synthetic generator with eight isolated large-language-model agents, one per role, each confined to its own role card and turn briefings, with no access to the other roles, the dashboard, or future turns, the same anti-lookahead discipline required of the real consultants. Each agent read its briefing in character and answered the real AHP tool of Section 4.3.2, pairwise comparisons and severity scores through the production consistency-ratio code. Unlike the synthetic dry run, this pilot exercised the full production pipeline rather than a standalone computation: campaign setup, turn injection, and questionnaire submission all ran through the same weekly-cycle service (`services/orchestrator.py`) that a browser-based consultant reaches, closing on the same live dashboard built for the real campaign. This pilot is exploratory: it was not part of the pre-registered protocol, and a single language-model instance per role has no statistical standing as a substitute for eight independent human judges. Table 24 gives its verdicts on the same six pre-registered tests.

Table 24: LLM-pilot verdicts: eight isolated language-model agents replaying the real AHP questionnaire through the production pipeline (exploratory, not pre-registered).

Test	Verdict	Measured value
H1 latency	Not confirmed	3 of 8 nodes at lag ≥ 1 turn (threshold: 5 of 8)
H2 early detection	Not confirmed	$H(\text{NovaFab})$ peaks at 0.017 on turns 6–12 (threshold: 0.3)
H3 criticality	Not confirmed	NovaFab top-1 on 7 of 10 informative turns, 70% (threshold: 80%; identical to the dry run, a U_d -independent test)
H4 predictive validity	Not confirmed	precision 0.00, recall 0.00 at $H \geq 0.5$ ($TP = 0$, $FP = 2$, $TN = 105$, $FN = 37$)
H5 contagion	Confirmed	non-affected nodes: $\Delta U_d = +0.079$ on press turns vs. $+0.005$ on calm turns
H6 hysteresis	Not applicable	real urgency shows no measurable decline over the recovery turns

Two results stand out. The urgency-contagion hypothesis (H5) is confirmed: declared urgency of non-affected nodes rises by 0.079 on press-event turns against 0.005 on calm turns, matching the pre-registered direction and ordering. The latency hypothesis (H1) moves from 1 of 8 nodes on the synthetic dry run to 3 of 8 nodes at a lag of one turn or more, closer to, but still short of, the pre-registered threshold of 5. Both results are a first sign of a non-circular declared-urgency

signal: a reasoning-based judge, unlike the dry run’s formula, can in principle lag or amplify the physical signal instead of tracking it by construction. H4 stays not confirmed, but for a specific and informative reason: these agents read their briefing carefully and declared urgency close to what the briefing in fact described, so their U_d tracks U_r too closely to leave the gap that H is built to catch, an even tighter tracking than the fixed-bias dry run ($H(\text{NovaFab})$ peaks at 0.017 against 0.070). H3 is unchanged from the dry run because it depends only on the network’s real-urgency criticality ranking, not on U_d at all. H6 stays not applicable for the same reason as the dry run: real urgency shows no measurable decline over the recovery turns. Taken together, this pilot motivates, without yet establishing, the idea that a genuinely reasoning-based declared-urgency signal can decouple from the physical signal in the way the model’s hypotheses predict, and it sharpens rather than removes the case for the real campaign: only a human’s imperfect, potentially under-reporting perception of a crisis, not a careful reader’s, can put H4 to a fair test.

The Pilot’s Final Application State. The pilot did not just produce six hypothesis verdicts: it left behind a live SupplyScore project, reloadable in the same dashboard a coordinator would use in production. Read directly from that project after the eighteenth turn, the network-average adequacy is 65.5, four of the eight nodes carry a hidden-risk reading above 0.1, one carries a false-urgency reading above 0.1, and the single worst-scoring node by a wide margin is Orbitalys at $A = 15.3$; Table 25 lists the full per-node snapshot.

Table 25: Final per-node snapshot of the LLM pilot’s SupplyScore project, read directly from the application after turn 18.

Node	U_d	U_r	A	F	H
Orbitalys	0.409	1.000	15.3	0.000	0.591
AvioSys Intégration	0.727	0.698	95.2	0.029	0.000
Électis EMS	0.914	0.965	82.9	0.000	0.052
CompoDis Europe	0.882	1.000	67.5	0.000	0.118
TransGlobal Fret	0.819	1.000	56.0	0.000	0.181
Meridian Semi	0.841	0.998	60.1	0.000	0.157
NovaFab Semiconductors	0.912	0.915	98.5	0.000	0.003
SilPure Materials	0.876	0.304	48.9	0.572	0.000

Orbitalys’ own explanation page traces the chain from a single scripted event to this score without manual computation. Its next milestone, the HÉLIOS-1 delivery, had slipped past its deadline by turn 18, which forces u_{time} to 1 under the rule

of Table 10; with the node's real urgency already saturated by this cause alone, the explanation page attributes the whole of it to the local block, none to the upstream network. The agent playing Orbitalys, meanwhile, declared $U_d = 0.41$. The resulting hidden risk $H = [U_r - U_d]_+ = 0.59$ is penalised through Equation 9 at $2.25 \times 0.59^{0.88} = 1.42$, giving $A = 100 \times (e^{-1.42} - e^{-2.25}) / (1 - e^{-2.25}) = 15.3$, the same number the dashboard reports, reconstructed term by term by the explainability page of Section 4.3.5 rather than by hand. The network-wide PROMETHEE II ranking (Section 4.3.3) places NovaFab and Orbitalys at the top of the priority list ($\phi = +0.220$ and $+0.215$), the expected bottleneck and the node its own missed milestone makes most urgent, though the same ranking carries a caveat the tool raises on its own: on this run, the time and performance blocks are perfectly correlated ($\rho = 1.00$), the composite-indicator risk already flagged in Section 3.2.6.

Current Status and Remaining Steps. At the time of writing, the campaign protocol is frozen, the data pack is hashed and frozen, all six operating scripts are delivered, and both validations above are complete. The human campaign itself, the only run able to decide the six pre-registered hypotheses, has not yet taken place: role assignment, the half-day kickoff, the nine working days of paired turns, and the reveal debrief remain to be executed. Section 5 states this as the immediate next step of the programme, together with what the recorded results will feed back into.

5 Conclusions

This R&D operation delivered *SupplyScore*: a multi-block, noisy-OR real urgency model propagated across the supply chain DAG, an asymmetric cognitive adequacy score confronting declared and computed urgency, and a persistent, multi-tenant digital twin operated as a weekly-cadence serious game, with a pre-registered validation campaign (Section 4.4) sized to measure the predictive validity of the hidden-risk signal. The scope of the delivered framework was bounded by explicit, documented design decisions rather than left implicit. This section states those limitations directly, draws the next steps, and closes with a personal assessment of the internship.

5.1 Limitations

Table 26 lists the limitations judged material enough to state directly, together with their operational implication.

Table 26: Documented limitations and their operational implications.

Limitation	Implication
No probabilistic guarantee	The framework produces probabilities, scores, and confidence intervals, not certainties; a model remains a simplification of the underlying reality.
Score freshness follows input cadence	Without an ERP connector, a node's score is never more current than its last submitted weekly review; manual entry is an accepted trade-off, not an oversight.
AHP is limited to four criteria	Saaty's method degrades past roughly seven pairwise-compared criteria, which is why declared urgency stays on four criteria while the larger set of KPI blocks is weighted through Fuzzy-BWM instead.
Event reversibility is conditional	Reverting a declared event is only guaranteed correct if no other write has touched the same KPI in the meantime; otherwise the system raises an explicit conflict for manual resolution rather than silently rebasing later writes, which would be mathematically dishonest.
Event calibration priors are defensible, not measured	The Beta-Bernoulli prior weight ($n_0 = 26$) and EMA reactivity coefficients follow standard conjugate Bayesian revision but are not empirically fitted; the serious game itself is the intended instrument for recalibrating them.
Empirical validation is designed, not yet executed	The calibration service and its confusion-matrix methodology are implemented and tested, but at the time of writing the HÉLIOS campaign (Section 4.4) has not yet been run with its human participants; no predictive-validity claim is made until it has.
Local time urgency looks only at the next milestone	A distant milestone that is silently drifting stays invisible to u_{time} until it becomes the next active one; the underlying function already receives the full milestone list, so a weighted multi-milestone extension is possible without a model change.
Double counting is possible	An operator can both declare a calibrated event and manually correct the same KPI in the same week; the interface warns but does not forbid this, since the human operator remains the final authority.

5.2 Next Steps

The immediate next step of the programme is to run the human campaign itself (Section 4.4). Role assignment across the eight consultants is not yet finalised and must be recorded in the campaign journal before the first turn, deliberately keeping the two most information-heavy roles, NovaFab and CompoDis, away from the most supply-chain-experienced consultants, so that a role effect is not confused with an expertise effect. A half-day kickoff, presentation, role cards, a signed contract, and a training turn, formally freezes the protocol; the eighteen campaign turns then run over nine working days with a daily journal and the two placebo sheets at turns 9 and 15, and the campaign closes with the pre-registered analysis and a reveal debrief that asks each consultant at which turn they recognised the real crisis.

That analysis will fill the two reserved needs left open by Section 4.4: the measured value of each of H1 to H6 against its pre-registered threshold, a verdict of supported, refuted, or inconclusive for each, and a sensitivity check (exclusion of absent consultants, a $\pm 50\%$ perturbation of the constants fixed by design) confirming whether those verdicts hold. Where the campaign supports a hypothesis but the underlying governance parameters, λ_{under} , α , the block weights ω_m , and the Bayesian prior weight n_0 , fit the measured data poorly, recalibrating them from campaign data rather than from the literature-derived defaults of Section 4.2 is the direct follow-on. Three contingencies are already arbitrated rather than left to be improvised under time pressure: the campaign truncates to turn 16 if it overruns by more than two turns rather than rushing the remaining responses, it drops to one turn per day if the questionnaire exceeds ten minutes at the training turn, and any consultant absence or unresolved consistency failure is carried forward unchanged and excluded from reactivity analyses.

Beyond the campaign itself, three further directions were identified but deliberately left out of the delivered scope: replacing single-operator AHP elicitation with a multi-expert protocol that aggregates several judgments per node, coupling SupplyScore with the DISCO platform it currently runs alongside rather than inside (Section 1), and using the adequacy score as an input to scheduling decisions rather than only as a diagnostic. None of the three is a design commitment yet: each depends on what the human campaign shows about whether the underlying signal is worth building on.

5.3 Bilan Personnel et Compétences Mobilisées

5.3.1 Compétences techniques mobilisées

Cette opération de R&D a mobilisé des compétences dans quatre domaines distincts, dont l'articulation constitue l'essentiel de la difficulté technique du sujet :

- **Aide à la décision multi-critères** : élicitation par comparaisons par paires (AHP, Fuzzy-BWM), validation de cohérence, agrégation par surclassement (PROMETHEE II).
- **Modélisation probabiliste** : agrégation noisy-OR, simulation de Monte Carlo, révision bayésienne conjuguée (Beta-Bernoulli) en régime de données rares.
- **Algorithmique de graphes** : propagation sur graphe orienté acyclique multi-rangs, tri topologique, simulation de choc et calcul de criticité réseau.
- **Génie logiciel** : architecture d'un système multi-tenant en Python, persistance SQLite durcie (WAL, sauvegarde, migrations), interface web Dash, démarche qualité en jalons avec couverture de tests bloquante et tests de bout en bout automatisés.

5.3.2 Apprentissages et difficultés rencontrées

- Quelle a été la principale difficulté technique ou méthodologique rencontrée, et comment a-t-elle été surmontée ?
- Qu'est-ce qui aurait été fait différemment avec le recul ?
- Quelles compétences (techniques ou non) ont le plus progressé au cours du stage ?
- En quoi ce stage éclaire-t-il un projet professionnel ?

Bibliography

- ADEME, 2014. *Impacts des limitations de vitesse sur la qualité de l'air et les émissions de CO₂*. Paris. Rapport technique. Agence de l'Environnement et de la Maîtrise de l'Énergie.
- AFRASIABI, A.; TAVANA, M.; CAPRIO, D., 2022. An extended hybrid fuzzy multi-criteria decision model for sustainable and resilient supplier selection. *Environmental Science and Pollution Research International*.
- AHMETOGLU, Y., 2025. *Debiasing Optimistic Planning in Knowledge Work: A Human-Centered Approach*. PhD thesis. University College London.
- AMIRI, M. et al., 2020. A New Fuzzy Best-Worst Method and Its Application in a Sustainable Supplier Selection Problem. *International Journal of Sustainable Development & World Ecology*. Vol. 27, no. 6, pp. 481–496. Available from DOI: [10.1080/13504509.2020.1731786](https://doi.org/10.1080/13504509.2020.1731786).
- AUCLIN, T., 2025. *BST Développement Applicatif DISCO*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- BAKER, K. R.; BERTRAND, J. W. M., 1982. A Dynamic Priority Rule for Scheduling Against Due-Dates. *Journal of Operations Management*. Vol. 3, no. 1, pp. 37–42. Available from DOI: [10.1016/0272-6963\(82\)90020-1](https://doi.org/10.1016/0272-6963(82)90020-1).
- BAULAIN, D.; DUBARRY, T., 2024. *Internet Physique, Optimisation du Routage et Indice de Performance Globale dans DISCO*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- BAXAS, A., 2025. *Architecture du Jumeau Numérique DISCO et Modélisation de la Résilience*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- BECKER, W. et al., 2017. Weights and Importance in Composite Indicators: Closing the Gap. *Ecological Indicators*. Vol. 80, pp. 12–22. Available from DOI: [10.1016/j.ecolind.2017.03.056](https://doi.org/10.1016/j.ecolind.2017.03.056).
- BEHZADIAN, M. et al., 2010. PROMETHEE: A Comprehensive Literature Review on Methodologies and Applications. *European Journal of Operational Research*. Vol. 200, no. 1, pp. 198–215. Available from DOI: [10.1016/j.ejor.2009.01.021](https://doi.org/10.1016/j.ejor.2009.01.021).
- BELLOTTI, F. et al., 2013. A Framework for Evaluating Serious Games. *Advances in Psychology*. Vol. 196, pp. 503–525. Available from DOI: [10.1016/B978-0-12-415815-5.00020-5](https://doi.org/10.1016/B978-0-12-415815-5.00020-5).
- BEN ZAKOUR, R., 2024. *BST DIN — Smart Green Supply Chain: Contexte et Historique du Programme*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- BOJKO, M. et al., 2021. Due Date-Related Order Prioritization for Scheduling with Decision Support in Dynamic Environments. In: *Advances in Production Management Systems: Artificial Intelligence for Sustainable and Resilient Production*

- Systems (APMS 2021)*. Cham: Springer. Vol. 630. IFIP Advances in Information and Communication Technology. Available from DOI: [10.1007/978-3-030-85874-2_25](https://doi.org/10.1007/978-3-030-85874-2_25).
- BORGONOVO, E.; CILLO, A.; PLISCHKE, E., 2021. Global sensitivity analysis of multi-criteria decision analysis: A review of concepts, methods, and applications. *European Journal of Operational Research*. Vol. 290, no. 3, pp. 805–827. Available from DOI: [10.1016/j.ejor.2020.07.034](https://doi.org/10.1016/j.ejor.2020.07.034).
- BRANS, J.-P.; MARESCHAL, B., 2005. PROMETHEE Methods. In: FIGUEIRA, J.; GRECO, S.; EHRGOTT, M. (eds.). *Multiple Criteria Decision Analysis: State of the Art Surveys*. New York: Springer, pp. 163–195. Available from DOI: [10.1007/0-387-23081-5_5](https://doi.org/10.1007/0-387-23081-5_5).
- BRANS, J.-P.; VINCKE, P.; MARESCHAL, B., 1986. How to Select and How to Rank Projects: The PROMETHEE Method. *European Journal of Operational Research*. Vol. 24, no. 2, pp. 228–238. Available from DOI: [10.1016/0377-2217\(86\)90044-5](https://doi.org/10.1016/0377-2217(86)90044-5).
- BRÉE, D. S.; CHALLET, D.; PEIRANO, P., 2010. Prediction accuracy and sloppiness of log-periodic functions. *Quantitative Finance*. Vol. 10, no. 6, pp. 671–680.
- CHAKRABORTTY, R. K.; RAHMAN, H. U.; RYAN, M. J., 2020. Efficient Priority Rules for Project Scheduling Under Dynamic Environments: A Heuristic Approach. *Computers & Industrial Engineering*. Vol. 147, p. 106627. Available from DOI: [10.1016/j.cie.2020.106627](https://doi.org/10.1016/j.cie.2020.106627).
- CHAUHAN, V. K.; BRINTRUP, A., 2022. Real-time large-scale supplier order assignments across two-tiers of a supply chain with penalty and dual-sourcing.
- CHEN, J. et al., 2025. Theoretical Foundations of Utility Accrual for Real-Time Systems.
- CHEN, K.; MÜHLETHALER, P., 1992. A scheduling algorithm for real-time computer systems using time value functions.
- CHEN, K.; MÜHLETHALER, P., 1996. A scheduling algorithm for tasks described by Time Value Function. *Real-Time Systems*.
- COHEN, L.; GRINSHPOUN, T.; STERN, R., 2019. Assigning Suppliers to Meet a Deadline. In: *Proceedings of the Twelfth International Symposium on Combinatorial Search (SoCS 2019)*. AAAI Press. Vol. 10, pp. 170–171. No. 1. Available from DOI: [10.1609/socs.v10i1.18475](https://doi.org/10.1609/socs.v10i1.18475).
- CORDOVA, A., 2024. *Green Supply Chain et Quantification des Émissions CO₂ dans DISCO*. Paris. Stage Report. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- DENG, J.; ZHANG, H.; WEI, S., 2020. Prediction of Vehicle-Cargo Matching Probability Based on Dynamic Bayesian Network. *International Journal of Production Research*. Vol. 59, no. 5, pp. 1419–1436. Available from DOI: [10.1080/00207543.2020.1714094](https://doi.org/10.1080/00207543.2020.1714094).

- DETERDING, S. et al., 2011. From Game Design Elements to Gamefulness: Defining Gamification. In: *Proceedings of the 15th International Academic MindTrek Conference: Envisioning Future Media Environments*. ACM, pp. 9–15. Available from DOI: [10.1145/2181037.2181040](https://doi.org/10.1145/2181037.2181040).
- DING, Y. et al., 2023. Real-Time Monitoring and Optimal Resource Allocation for Automated Container Terminals: A Digital Twin Application at the Yangshan Port. *Journal of Advanced Transportation*.
- DIWASSINGH, K. et al., 2020. Task Selection and Workload: A Focus on Completing Easy Tasks Hurts Long-Run Performance. *Management Science*. Vol. 66, no. 10, pp. 4397–4416. Available from DOI: [10.1287/mnsc.2019.3366](https://doi.org/10.1287/mnsc.2019.3366).
- DYER, J. S., 1990. Remarks on the Analytic Hierarchy Process. *Management Science*. Vol. 36, no. 3, pp. 249–258.
- EUROPEAN COMMISSION, 2017. *VECTO — Vehicle Energy Consumption Calculation Tool* [online]. [visited on 2026-04-24]. Technical Documentation. European Commission, DG Climate Action. Available from: https://climate.ec.europa.eu/eu-action/transport-decarbonisation/road-transport/vehicle-energy-consumption-calculation-tool-vecto_en.
- FEIGENBAUM, J. A., 2001. A statistical analysis of log-periodic precursors to financial crashes. *Quantitative Finance*. Vol. 1, no. 3, pp. 346–360.
- FULLER, A. et al., 2020. Digital Twin: Enabling Technologies, Challenges and Open Research. *IEEE Access*. Vol. 8, pp. 108952–108971. Available from DOI: [10.1109/ACCESS.2020.2998358](https://doi.org/10.1109/ACCESS.2020.2998358).
- GRECO, S. et al., 2018. On the Methodological Framework of Composite Indices: A Review of the Issues of Weighting, Aggregation, and Robustness. *Social Indicators Research*. Vol. 141, pp. 61–94. Available from DOI: [10.1007/s11205-017-1832-9](https://doi.org/10.1007/s11205-017-1832-9).
- GRIEVES, M., 2014. *Digital Twin: Manufacturing Excellence Through Virtual Factory Replication*. White Paper. Florida Institute of Technology.
- GUO, S.; ZHAO, H., 2017. Fuzzy Best-Worst Multi-Criteria Decision-Making Method and Its Applications. *Knowledge-Based Systems*. Vol. 121, pp. 23–31. Available from DOI: [10.1016/j.knosys.2017.01.010](https://doi.org/10.1016/j.knosys.2017.01.010).
- GUO, W. et al., 2020. A Dynamic Shipment Matching Problem in Hinterland Synchronomodal Transportation. *Decision Support Systems*. Vol. 134, p. 113289. Available from DOI: [10.1016/j.dss.2020.113289](https://doi.org/10.1016/j.dss.2020.113289).
- HECKMANN, I.; COMES, T.; NICKEL, S., 2015. A Critical Review on Supply Chain Risk — Definition, Measure and Modeling. *Omega*. Vol. 52, pp. 119–132. Available from DOI: [10.1016/j.omega.2014.10.004](https://doi.org/10.1016/j.omega.2014.10.004).
- HLIOUI, R.; GHARBI, A.; HAJJI, A., 2017. Joint Supplier Selection, Production and Replenishment of an Unreliable Manufacturing-Oriented Supply Chain. *International Journal of Production Economics*. Vol. 187, pp. 53–67. Available from DOI: [10.1016/j.ijpe.2017.02.004](https://doi.org/10.1016/j.ijpe.2017.02.004).

- HOEN, K. M. R. et al., 2014. Effect of Carbon Emission Regulations on Transport Mode Selection under Stochastic Demand. *Flexible Services and Manufacturing Journal*. Vol. 26, no. 1–2, pp. 170–195. Available from DOI: [10.1007/s10696-012-9151-6](https://doi.org/10.1007/s10696-012-9151-6).
- HOSSEINI, S.; IVANOV, D., 2020. Bayesian Networks for Supply Chain Risk, Resilience and Ripple Effect Analysis: A Literature Review. *Expert Systems with Applications*. Vol. 161, p. 113649. Available from DOI: [10.1016/j.eswa.2020.113649](https://doi.org/10.1016/j.eswa.2020.113649).
- HOSSEINI, S.; IVANOV, D.; DOLGUI, A., 2020. Ripple Effect Modelling of Supplier Disruption: Integrated Markov Chain and Dynamic Bayesian Network Approach. *International Journal of Production Research*. Vol. 58, no. 11, pp. 3284–3303. Available from DOI: [10.1080/00207543.2019.1661538](https://doi.org/10.1080/00207543.2019.1661538).
- HOSSEINI, S.; TAJIK, N., et al., 2019. Resilient Supplier Selection and Optimal Order Allocation under Disruption Risks. *International Journal of Production Economics*. Vol. 213, pp. 124–137. Available from DOI: [10.1016/j.ijpe.2019.03.018](https://doi.org/10.1016/j.ijpe.2019.03.018).
- HRUŠOVSKÝ, M. et al., 2021. Real-time disruption management approach for intermodal freight transportation. *Journal of Cleaner Production*.
- HUANG, P. et al., 2020. A Bayesian Network Model to Predict the Effects of Interruptions on Train Operations. *Transportation Research Part C: Emerging Technologies*. Vol. 114, pp. 338–358. Available from DOI: [10.1016/j.trc.2020.02.015](https://doi.org/10.1016/j.trc.2020.02.015).
- INTERNATIONAL ORGANIZATION FOR STANDARDIZATION, 2023. *ISO 14083:2023 Greenhouse gases — Quantification and Reporting of Greenhouse Gas Emissions Arising from Transport Chain Operations*. Geneva.
- IVANOV, D.; DOLGUI, A., 2019. New Disruption Risk Management Perspectives in Supply Chains: Digital Twins, the Physical Internet, and Smart Contracts. *IFAC-PapersOnLine*. Vol. 52, no. 13, pp. 337–342. Available from DOI: [10.1016/j.ifacol.2019.11.138](https://doi.org/10.1016/j.ifacol.2019.11.138).
- IVANOV, D.; DOLGUI, A., 2020. A Digital Supply Chain Twin for Managing the Disruption Risks and Resilience in the Era of Industry 4.0. *Production Planning & Control*. Vol. 32, no. 9, pp. 775–788. Available from DOI: [10.1080/09537287.2020.1768450](https://doi.org/10.1080/09537287.2020.1768450).
- IVANOV, D.; DOLGUI, A., et al., 2021. Digital Supply Chain Twins: Managing the Ripple Effect, Resilience, and Disruption Risks by Data-Driven Optimization, Simulation, and Visibility. *International Journal of Production Research*. Vol. 59, no. 3, pp. 957–974. Available from DOI: [10.1080/00207543.2020.1841183](https://doi.org/10.1080/00207543.2020.1841183).
- JOHANSEN, A.; SORNETTE, D., 2000. Crashes as Critical Points. *International Journal of Theoretical and Applied Finance*. Vol. 3, no. 2, pp. 219–255. Available from DOI: [10.1142/S0219024900000115](https://doi.org/10.1142/S0219024900000115).
- JOHANSEN, A.; SORNETTE, D., 2001. Finite-time singularity in the dynamics of the world population, economic and financial indices. *Physica A: Statistical*

- Mechanics and its Applications*. Vol. 294, no. 3–4, pp. 465–502. Available from DOI: [10.1016/S0378-4371\(01\)00105-4](https://doi.org/10.1016/S0378-4371(01)00105-4).
- KABOUCHE, I., 2025. *Scoring de Risque et Réseaux Bayésiens pour la Supply Chain Prédictive*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- KAGERMANN, H.; WAHLSTER, W.; HELBIG, J., 2013. *Recommendations for Implementing the Strategic Initiative INDUSTRIE 4.0*. Munich. acatech — National Academy of Science and Engineering.
- KAHNEMAN, D., 2011. *Thinking, Fast and Slow*. New York: Farrar, Straus and Giroux. ISBN 978-0374533557.
- KAHNEMAN, D.; TVERSKY, A., 1979. Prospect Theory: An Analysis of Decision under Risk. *Econometrica*. Vol. 47, no. 2, pp. 263–292. Available from DOI: [10.2307/1914185](https://doi.org/10.2307/1914185).
- KALABOUKAS, K. et al., 2021. Implementation of Cognitive Digital Twins in Connected and Agile Supply Networks—An Operational Model. *Applied Sciences*. Vol. 11, no. 9, p. 4103. ISSN 2076-3417. Available from DOI: [10.3390/app11094103](https://doi.org/10.3390/app11094103).
- KARAM, A.; REINAU, K., 2022. A Real-Time Decision Support Approach for Managing Disruptions in Line-Haul Freight Transport Networks. *IEEE Transactions on Intelligent Transportation Systems*.
- KIM, S.; TONG, J.; PEDEN, C., 2020. Admission Control Biases in Hospital Unit Capacity Management: How Occupancy Information Hurdles and Decision Noise Impact Utilization. *Management Science*. Vol. 66, no. 5, pp. 1803–1824. Available from DOI: [10.1287/mnsc.2019.3303](https://doi.org/10.1287/mnsc.2019.3303).
- KOLLER, D.; FRIEDMAN, N., 2009. *Probabilistic Graphical Models: Principles and Techniques*. Cambridge, MA: MIT Press. ISBN 978-0262013192.
- KULIGOWSKI, T. et al., 2025. Comparative Evaluation of the Manchester Triage System and Emergency Severity Index in Predicting Critical Events in the Emergency Department. *BMC Emergency Medicine*. Vol. 25. Available from DOI: [10.1186/s12873-025-01420-8](https://doi.org/10.1186/s12873-025-01420-8). PMID: 41286624.
- LAFKIHI, M.; PAN, S.; BALLOT, E., 2025. Game: The Freight Transportation Game: Operational Challenges for Carriers in Online Spot Market Platforms. *INFORMS Transactions on Education*. Vol. 26, no. 1, pp. 91–101. Available from DOI: [10.1287/ited.2024.0091](https://doi.org/10.1287/ited.2024.0091).
- LEE, H. L.; PADMANABHAN, V.; WHANG, S., 1997. The Bullwhip Effect in Supply Chains. *MIT Sloan Management Review*. Vol. 38, no. 3, pp. 93–102.
- LEUNG, E.; LEE, C.; OUYANG, Z., 2022. From traditional warehouses to Physical Internet hubs: A digital twin-based inbound synchronization framework for PI-order management. *International Journal of Production Economics*.
- LI, Y.; ZOBEL, C. W., 2020. Exploring Supply Chain Network Resilience in the Presence of the Ripple Effect. *International Journal of Production Economics*. Vol. 228, p. 107664. Available from DOI: [10.1016/j.ijpe.2020.107664](https://doi.org/10.1016/j.ijpe.2020.107664).

- LIANG, F.; BRUNELLI, M.; REZAEI, J., 2020. Consistency Issues in the Best Worst Method: Measurements and Thresholds. *Omega*. Vol. 96, p. 102175. Available from DOI: [10.1016/j.omega.2019.102175](https://doi.org/10.1016/j.omega.2019.102175).
- LIU, M. et al., 2021. A New Robust Dynamic Bayesian Network Approach for Disruption Risk Assessment under the Supply Chain Ripple Effect. *International Journal of Production Research*. Vol. 59, no. 1, pp. 265–285. Available from DOI: [10.1080/00207543.2020.1841318](https://doi.org/10.1080/00207543.2020.1841318).
- LOCKAMY, A.; MCCORMACK, K., 2011. Analysing Risks in Supply Networks to Facilitate Outsourcing Decisions. *International Journal of Production Research*. Vol. 48, no. 2, pp. 593–611. Available from DOI: [10.1080/00207540903175152](https://doi.org/10.1080/00207540903175152).
- MACKWAY-JONES, K.; MARSDEN, J.; WINDLE, J., 2013. *Emergency Triage: Manchester Triage Group*. 3rd ed. Oxford: Wiley-Blackwell. ISBN 978-1118299029.
- MILLS, A.; ANDERSON, J. H., 2010. A Stochastic Framework for Multiprocessor Soft Real-Time Scheduling.
- MOHADDESI, O.; POLARIS, V.; HARTEVELD, C., 2020. Introducing Gamettes: A Playful Approach for Capturing Decision-Making for Informing Behavioral Models. In: *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. ACM, pp. 1–14. Available from DOI: [10.1145/3313831.3376527](https://doi.org/10.1145/3313831.3376527).
- MOKHTAR, S. et al., 2019. Supplier Portfolio Selection Based on the Monitoring of Supply Risk Indicators. *Simulation Modelling Practice and Theory*. Vol. 97, p. 101955. Available from DOI: [10.1016/j.simpat.2019.101955](https://doi.org/10.1016/j.simpat.2019.101955).
- MONDELICE, R., 2025. *BST TMS — Transport Management System et Interface Routage DISCO*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- MONTREUIL, B., 2011. Toward a Physical Internet: Meeting the Global Logistics Sustainability Grand Challenge. *Logistics Research*. Vol. 3, no. 2–3, pp. 71–87. Available from DOI: [10.1007/s12159-011-0045-x](https://doi.org/10.1007/s12159-011-0045-x).
- NILOOFAR, P.; LAZAROVA-MOLNAR, S.; SHAKER, H. R., 2023. A General Framework for Human-in-the-Loop Cognitive Digital Twins. In: *Proceedings of the 2023 Winter Simulation Conference (WSC)*. IEEE, pp. 4155–4166. Available from DOI: [10.1109/WSC57122.2023.10407598](https://doi.org/10.1109/WSC57122.2023.10407598).
- PEARL, J., 1988. *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. San Francisco: Morgan Kaufmann. ISBN 978-1558604797.
- PEREZ, H.; WASSICK, J.; GROSSMANN, I., 2022. A digital twin framework for on-line optimization of supply chain business processes. *Computers & Chemical Engineering*.
- PINEDO, M. L., 2012. *Scheduling: Theory, Algorithms, and Systems*. 4th ed. New York: Springer. Available from DOI: [10.1007/978-1-4614-2361-4](https://doi.org/10.1007/978-1-4614-2361-4).
- PRENSKY, M., 2001. *Digital Game-Based Learning*. New York: McGraw-Hill. ISBN 978-0071454995.

- QAZI, A. et al., 2017. Project Complexity and Risk Management (ProCRiM): Towards Modelling Project Complexity Driven Risk Paths in Construction Projects. *International Journal of Project Management*. Vol. 34, no. 7, pp. 1183–1198. Available from DOI: [10.1016/j.ijproman.2016.05.008](https://doi.org/10.1016/j.ijproman.2016.05.008).
- REZAEI, J., 2015. Best-Worst Multi-Criteria Decision-Making Method. *Omega*. Vol. 53, pp. 49–57. Available from DOI: [10.1016/j.omega.2014.11.009](https://doi.org/10.1016/j.omega.2014.11.009).
- ROSENTHAL, S.; HIATT, L. M., 2020. Human-Centered Decision Support for Agenda Scheduling. *AI & Society*. Available from DOI: [10.1007/s00146-020-00960-y](https://doi.org/10.1007/s00146-020-00960-y).
- RUSOU, Z.; AMAR, M.; AYAL, S., 2020. The Psychology of Task Management: The Smaller Tasks Trap. *Judgment and Decision Making*. Vol. 15, no. 4, pp. 586–599.
- SAATY, T. L., 1980. *The Analytic Hierarchy Process: Planning, Priority Setting, Resource Allocation*. New York: McGraw-Hill. ISBN 978-0070543713.
- SAATY, T. L., 1990. How to Make a Decision: The Analytic Hierarchy Process. *European Journal of Operational Research*. Vol. 48, no. 1, pp. 9–26. Available from DOI: [10.1016/0377-2217\(90\)90057-I](https://doi.org/10.1016/0377-2217(90)90057-I).
- SAWIK, T., 2013. Integrated Selection of Suppliers and Scheduling of Customer Orders in the Presence of Supply Chain Disruption Risks. *International Journal of Production Research*. Vol. 51, no. 23–24, pp. 7006–7022. Available from DOI: [10.1080/00207543.2013.852702](https://doi.org/10.1080/00207543.2013.852702).
- SAWIK, T., 2014. Joint Supplier Selection and Scheduling of Customer Orders under Disruption Risks: Single vs. Dual Sourcing. *Omega*. Vol. 43, pp. 83–95. Available from DOI: [10.1016/j.omega.2013.06.007](https://doi.org/10.1016/j.omega.2013.06.007).
- SAWIK, T., 2016. Integrated Supply, Production and Distribution Scheduling under Disruption Risks. *Omega*. Vol. 62, pp. 131–144. Available from DOI: [10.1016/j.omega.2015.09.005](https://doi.org/10.1016/j.omega.2015.09.005).
- SAWIK, T., 2020. On the Risk-Averse Selection of Resilient Multi-Tier Supply Portfolio. *Omega*. Vol. 98, p. 102102. Available from DOI: [10.1016/j.omega.2019.102102](https://doi.org/10.1016/j.omega.2019.102102).
- SHEN, X. et al., 2019. Managing Supply Chains with Expediting and Multiple Demand Classes. *Production and Operations Management*. Vol. 28, no. 5, pp. 1129–1148. Available from DOI: [10.1111/poms.12974](https://doi.org/10.1111/poms.12974).
- SMART FREIGHT CENTRE, 2023. *GLEC Framework for Logistics Emissions Accounting, Version 3.2* [online]. [visited on 2026-04-24]. Technical Report. Smart Freight Centre. Available from: <https://www.smartfreightcentre.org/en/our-programs/emissions-accounting/global-logistics-emissions-council/calculate-report-glec-framework/>. Aligned with ISO 14083:2023.
- STAM, A.; SILVA, A. P. D., 1997. Stochastic Judgments in the AHP: The Measurement of Rank Reversal Probabilities. *Decision Sciences*. Vol. 28, no. 3, pp. 655–688. ISSN 1540-5915. Available from DOI: [10.1111/j.1540-5915.1997.tb01326.x](https://doi.org/10.1111/j.1540-5915.1997.tb01326.x).

- STEEL, P.; KÖNIG, C. J., 2006. Integrating theories of motivation. *Academy of Management Review*. Vol. 31, no. 4, pp. 889–913. Available from DOI: [10.5465/AMR.2006.22527462](https://doi.org/10.5465/AMR.2006.22527462).
- STERMAN, J. D., 1989. Modeling Managerial Behavior: Misperceptions of Feedback in a Dynamic Decision Making Experiment. *Management Science*. Vol. 35, no. 3, pp. 321–339. Available from DOI: [10.1287/mnsc.35.3.321](https://doi.org/10.1287/mnsc.35.3.321).
- SUN, X. et al., 2020. Combating Lead-Time Uncertainty in Global Supply Chain's Shipment-Assignment: Is It Wise to Be Risk-Averse? *Transportation Research Part B: Methodological*. Vol. 138, pp. 406–434. Available from DOI: [10.1016/j.trb.2020.05.015](https://doi.org/10.1016/j.trb.2020.05.015).
- TAGHAVI, S. M. et al., 2024. Sustainable and Resilient Supplier Selection, Order Allocation, and Production Scheduling Problem under Disruption Utilizing Conditional Value at Risk. *Journal of Modelling in Management*. Vol. 19, no. 2, pp. 658–692. Available from DOI: [10.1108/JM2-10-2022-0250](https://doi.org/10.1108/JM2-10-2022-0250).
- TAO, F. et al., 2018. Digital Twin-Driven Product Design, Manufacturing and Service with Big Data. *The International Journal of Advanced Manufacturing Technology*. Vol. 94, no. 9–12, pp. 3563–3576. Available from DOI: [10.1007/s00170-017-0233-1](https://doi.org/10.1007/s00170-017-0233-1).
- TOMLIN, B., 2006. On the Value of Mitigation and Contingency Strategies for Managing Supply Chain Disruption Risks. *Management Science*. Vol. 52, no. 5, pp. 639–657. Available from DOI: [10.1287/mnsc.1060.0515](https://doi.org/10.1287/mnsc.1060.0515).
- TRIANAPHYLLOU, E.; YANASE, J., 2024. The use of pairwise comparisons for decision making may lead to grossly inaccurate results.
- TRIANAPHYLLOU, E., 2000. *Multi-Criteria Decision Making Methods: A Comparative Study*. Dordrecht: Kluwer Academic Publishers. Available from DOI: [10.1007/978-1-4757-3157-6](https://doi.org/10.1007/978-1-4757-3157-6).
- TURNER, C. J. et al., 2021. Human in the Loop: Industry 4.0 Technologies and Scenarios for Worker Mediation of Automated Manufacturing. *IEEE Access*. Vol. 9, pp. 104484–104502. Available from DOI: [10.1109/ACCESS.2021.3098555](https://doi.org/10.1109/ACCESS.2021.3098555).
- WANG, Y.; WANG, X.; LIU, A., 2020. Digital Twin-Driven Supply Chain Planning. *Procedia CIRP*. Vol. 93, pp. 526–531. Available from DOI: [10.1016/j.procir.2020.04.073](https://doi.org/10.1016/j.procir.2020.04.073).
- YUCESAN, M. et al., 2019. An Integrated Best-Worst and Interval Type-2 Fuzzy TOPSIS Methodology for Green Supplier Selection. *Mathematics*. Vol. 7, no. 2, p. 182. ISSN 2227-7390. Available from DOI: [10.3390/math7020182](https://doi.org/10.3390/math7020182).
- ZHANG, J. et al., 2021. Bi-level dynamic scheduling architecture based on service unit digital twin agents. *Journal of Manufacturing Systems*. Vol. 60, pp. 59–79. ISSN 0278-6125. Available from DOI: [10.1016/j.jmsy.2021.05.007](https://doi.org/10.1016/j.jmsy.2021.05.007).
- ZHANG, Y. et al., 2026. Digital twin driven opti-state control approach for smart warehousing in the synchronous operating environment. *Robotics and Computer-*

- Integrated Manufacturing*. Vol. 97, p. 103099. ISSN 0736-5845. Available from DOI: [10.1016/j.rcim.2025.103099](https://doi.org/10.1016/j.rcim.2025.103099).
- ZHENG, X.; LU, J.; KIRITSIS, D., 2021. The emergence of cognitive digital twin: vision, challenges and opportunities. *International Journal of Production Research*. Vol. 60, no. 24, pp. 7610–7632. ISSN 1366-588X. Available from DOI: [10.1080/00207543.2021.2014591](https://doi.org/10.1080/00207543.2021.2014591).
- ZHENG, X.; ZHANG, L., 2020a. Risk assessment of supply-chain systems: a probabilistic inference method. *Enterprise Information Systems*. Vol. 14, no. 6, pp. 858–877. ISSN 1751-7583. Available from DOI: [10.1080/17517575.2020.1762004](https://doi.org/10.1080/17517575.2020.1762004).
- ZHENG, X.; ZHANG, L., 2020b. Risk Assessment of Supply-Chain Systems: A Probabilistic Inference Method. *Enterprise Information Systems*. Vol. 14, no. 9–10, pp. 1336–1353. Available from DOI: [10.1080/17517575.2019.1684163](https://doi.org/10.1080/17517575.2019.1684163).
- ZHOU, C. et al., 2021. Analytics with digital-twinning: A decision support system for maintaining a resilient port. *Decision Support Systems*. Vol. 143, p. 113496. ISSN 0167-9236. Available from DOI: [10.1016/j.dss.2021.113496](https://doi.org/10.1016/j.dss.2021.113496).
- ZHU, M.; YANG, Y.; HSEE, C. K., 2018. The Mere Urgency Effect. *Journal of Consumer Research*. Vol. 45, no. 3, pp. 673–690. Available from DOI: [10.1093/jcr/ucy008](https://doi.org/10.1093/jcr/ucy008).
- ZHU, X. et al., 2025. Towards Industry 5.0: digital twin-enhanced approach for dynamic supply chain rescheduling with real-time order arrival and acceptance. *International Journal of Production Research*.

Glossary

- $A(t)$ Adequacy score at time t , quantifying the asymmetric Prospect Theory penalty between U_d and $U_r(t)$: $A = 100 \times (\exp(-\text{penalty}) - \text{floor}) / (1 - \text{floor})$ with $\text{penalty} = \lambda_{\text{under}} H^\alpha + \lambda_{\text{over}} F^\alpha, \in [0, 100]$. 88
- F, H False urgency $F = [U_d - U_r]_+$ (unjustified panic) and hidden risk $H = [U_r - U_d]_+$ (silent danger), the two asymmetrically penalised pathologies underlying $A(t)$. 88
- U_d Declared urgency, a scalar $\in [0, 1]$ elicited from the human operator via an AHP or Fuzzy-BWM questionnaire, propagated downstream across the supply chain DAG through the leaky coefficients γ_{ik} . 88
- $U_r(t)$ Real (operationally computable) urgency at time t : a weighted probabilistic OR aggregation of six local KPI blocks (time, capacity, performance, risk, cost, CO2), $U_{r,\text{local}} = 1 - \prod_m (1 - u_m)^{\omega_m}$, propagated upstream across the supply chain DAG through the leaky coefficients β_{ji} . 88
- $\Delta U(t)$ Raw urgency mismatch: $U_d - U_r(t) \in (-1, 1)$; positive = over-triage (false urgency), negative = under-triage (hidden risk). 88
- $\Delta U_{r,\text{final}}$ Network criticality delta: the increase in real urgency propagated to final-customer nodes when a given node's local urgency is forced to saturation in a systematic shock simulation. 88
- ω_m Aggregation weight of KPI block m in the real urgency noisy-OR aggregation, reused unchanged as the PROMETHEE II criteria weight for network-wide node outranking. 88

Acronyms

- AHP** Analytic Hierarchy Process. 32, 44, 88
- BST** Bilan Scientifique et Technique. 88
- BWM** Best-Worst Method. 32, 88
- CI** Consistency Index. 54, 88
- CR** Consistency Ratio. 32, 54, 55, 88

CSRD Corporate Sustainability Reporting Directive. 88

DAG Directed Acyclic Graph. 88

DIN Direction Innovation Numérique. 10, 88

DISCO Supply Chain Illustrator (Digital Twin, DIN ALTEN). 88

DT Digital Twin. 30, 31, 88

F-BWM Fuzzy Best-Worst Method. 32, 88

GLEC Global Logistics Emissions Council. 10, 34, 88

HGV Heavy Goods Vehicle. 12, 88

IoT Internet of Things. 10, 88

IPG Itinerary Performance Index. 12, 88

KPI Key Performance Indicator. 88

LPPL Log-Periodic Power Law. 26, 88

MCDM Multi-Criteria Decision-Making. 22, 29, 32, 88

MTS Manchester Triage System. 33, 88

PI Physical Internet. 22, 34, 88

PROMETHEE Preference Ranking Organisation METHod for Enrichment Evaluation. 88

RI Random Index (Saaty). 54, 88

SC Supply Chain. 30, 31, 88

TFN Triangular Fuzzy Number. 32, 88

TMT Temporal Motivational Theory. 27, 88

TTR Time To Recovery. 12, 88

TTS Time To Service. 12, 88

VECTO Vehicle Energy Consumption calculation TOol. 35, 88

XAI Explainable Artificial Intelligence. 88