



Measuring the Gap Between Declared and Real Urgency in a Supply Chain Digital Twin

A decision-support layer for the Smart Green Supply Chain programme

ALLEN Sud-Ouest | Direction de l'Innovation
Lab de Toulouse

HOARAU John

Promotion 2026 | Cycle Ingénieur 3^e année

Parcours double diplôme Cranfield University (DED)

Période du 13 avril 2026 au 16 octobre 2026

3^e année | Mission de fin d'études

Rapport confidentiel : OUI

Encadrement

Tuteur entreprise	Sébastien PERTHUISOT Directeur de Programme de Recherche, Expert Qualité & Supply Chain, DIN
Tuteur entreprise	Fabien JULES Chargé Innovation, ALLEN
Tuteur ESTIA	Dr. David GOMEZ Enseignant-chercheur en informatique
Tuteur CRANFIELD	Dr. Peter Sherar Enseignant-chercheur en informatique

Thématique du travail réalisé

Informatique | Génie industriel

ESTIA – École Supérieure des Technologies
Industrielles Avancées, Technopole
Izarbel, 64210 Bidart

Acknowledgements

My thanks go first to **Sébastien Perthuisot**, Research Programme Director, Quality & Supply Chain Expert at the Direction de l'Innovation, who set the direction of this work and then left me room to take it somewhere. The weekly review he held with me was where most of the difficult questions in this report were first asked.

I thank **Émilienne Lardy** and **Selma Ferhat**, who ran the scrum and planning cadence that gave the six months their rhythm, and whose reviews repeatedly sent me back to check a claim I had been ready to make. I thank them also for their scientific supervision.

I thank **Liam Jones**, who worked on the network representation side of the same programme. Sharing a weekly demo with someone solving a neighbouring problem did more for the quality of my own work than any amount of solitary effort would have.

I thank **Dr David Gomez** of ESTIA for the academic supervision and framing of this mission, and **Dr Peter Sherar** of Cranfield University for supervising the research strand that ran alongside it.

Finally, I thank the whole Direction de l'Innovation team at the Toulouse Lab for the welcome, and my family for their support across the whole of this engineering degree.

Résumé

Dans une chaîne d'approvisionnement, l'urgence qu'un opérateur déclare diverge souvent de celle que l'on peut calculer à partir de l'état opérationnel réel. Sur-déclarer gaspille une capacité partagée ; sous-déclarer laisse un risque de rupture invisible. Aucun jumeau numérique existant ne mesure cet écart.

Cette mission de fin d'études, menée à la Direction de l'Innovation d'ALTEN dans le cadre du programme Smart Green Supply Chain, a conçu, développé et évalué une couche logicielle qui le mesure. L'urgence réelle est agrégée à chaque nœud depuis six blocs d'indicateurs par un OU probabiliste pondéré, puis propagée sur le réseau de fournisseurs ; l'urgence déclarée est élicitée par un questionnaire de comparaisons par paires contrôlé en cohérence ; leur écart est noté sur $[0, 100]$ par une pénalité asymétrique. Le système livré, SupplyScore, compte 73 modules Python. Il a été éprouvé par une campagne pré-enregistrée rejouant la pénurie de semi-conducteurs de 2020 à 2022 sur dix-neuf tours et huit nœuds.

Le résultat principal est positif et porte sur la question que pose réellement un coordinateur : quel fournisseur regarder cette semaine. Sur les six tours qui précèdent la mise au jour de la crise, l'écart désigne trois nœuds, tous trois futurs défaillants sur un engagement de livraison, **sans une seule fausse alerte** ; le plus mal noté est signalé huit tours avant de rater son jalon, alors que son propre responsable déclare une urgence de 0,001. Le signal s'émousse ensuite, pour une raison structurelle et non par défaut : quand la crise devient publique, l'urgence déclarée rattrape l'urgence calculée et un écart entre les deux se referme par construction. C'est donc un instrument d'alerte précoce, au plus net quand une alerte vaut encore quelque chose.

La couche de prévision bâtie au-dessus, elle, a d'abord échoué de façon précise : mal calibrée, avec vingt-et-une prévisions à 0,948 de moyenne jamais suivies de l'issue prédite au sens de la définition de production. Un seul bloc l'expliquait, le bloc temporel lisant une grandeur qu'aucun choc n'atteignait. Ce diagnostic a été traité et re-mesuré : le correctif évident s'est révélé faux à la mesure, le temps perdu devant s'*ajouter* à la date d'achèvement et non s'y multiplier, et après correction le déficit de calibration est divisé par deux à chaque horizon, de $-0,553$ à $-0,216$ à quatre semaines. Il reste négatif : la probabilité calibrée n'est pas encore publiable comme telle.

Un troisième résultat concerne l'opérateur : au tour où la prévision devient visible, les huit déclarants ont tous revu leur urgence à la baisse ou l'ont maintenue, deux tours avant l'événement critique. Ce qui borne l'ensemble n'est pas la conception mais l'échantillon : huit nœuds, onze jalons, un scénario. Élargir la base de données est le travail en cours et la condition de toute généralisation.

Abstract

In a supply chain, the urgency an operator declares often diverges from the urgency computable from the real operational state. Over-declaring wastes shared capacity; under-declaring leaves rupture risk invisible. No existing digital twin measures that gap.

This final-year mission, carried out at the Direction de l'Innovation of ALTEN within the Smart Green Supply Chain programme, designed, built and evaluated a software layer that measures

it. Real urgency is aggregated at each node from six indicator blocks through a weighted probabilistic OR and propagated across the supplier network; declared urgency is elicited through a consistency-gated pairwise questionnaire; their gap is scored on $[0, 100]$ with an asymmetric penalty. The delivered system, SupplyScore, comprises 73 Python modules covered by 1831 automated test functions. It was tested through a pre-registered campaign replaying the 2020 to 2022 semiconductor shortage over nineteen turns and eight nodes.

The headline result is positive and concerns the question a coordinator actually asks: which supplier to look at this week. Over the six turns before the crisis becomes public, the gap flags three nodes, all three of which later miss a delivery commitment, **with no false alarm on any turn**; the worst-scoring of them is flagged eight turns before it misses, while its own manager declares an urgency of 0.001. The signal then weakens, for a structural reason rather than a defect: once the crisis is public, declared urgency catches up with computed urgency and a measure defined as the difference between them necessarily shrinks. It is an early-warning instrument, sharpest while a warning is still worth something.

The forecast layer built on top of it failed first, and failed precisely: severely miscalibrated, with twenty-one forecasts averaging 0.948 never followed by the predicted outcome under the production definition. One block explained it, the time block reading a quantity no shock ever reached. That diagnosis was acted on and re-measured. The obvious repair proved wrong under measurement, lost time having to be *added* to the completion date rather than multiplied into it, and after repair the calibration deficit halves at every horizon, from -0.553 to -0.216 at four weeks. It remains negative: the calibrated probability is not yet publishable as one.

A third result concerns the operator: on the turn the forecast became visible, all eight declarants revised their urgency downward or held it, two turns before the critical event. What bounds all of this is not the design but the sample, eight nodes, eleven milestones and one scenario. Widening that base is the work in progress and the precondition for any generalisation.

Resumen

En una cadena de suministro, la urgencia que declara un operador difiere a menudo de la urgencia calculable a partir del estado operativo real. Declarar de más desperdicia capacidad compartida; declarar de menos deja invisible el riesgo de ruptura. Ningún gemelo digital existente mide esa diferencia.

Esta misión de fin de estudios, realizada en la Dirección de Innovación de ALTEN dentro del programa Smart Green Supply Chain, diseñó, desarrolló y evaluó una capa de software que la mide. La urgencia real se agrega en cada nodo a partir de seis bloques de indicadores mediante un O probabilístico ponderado y se propaga por la red de proveedores; la urgencia declarada se obtiene mediante un cuestionario de comparaciones por pares con control de coherencia; su diferencia se puntúa en $[0, 100]$ con una penalización asimétrica. El sistema entregado, SupplyScore, consta de 73 módulos Python cubiertos por 1831 funciones de prueba automatizadas. Se sometió a una campaña preregistrada que reproduce la escasez de semiconductores de 2020 a 2022 durante diecinueve turnos y ocho nodos.

El resultado principal es positivo y se refiere a la pregunta que realmente se hace un coordinador : qué proveedor mirar esta semana. En los seis turnos previos a que la crisis se haga pública, la diferencia señala tres nodos, los tres futuros incumplidores de un compromiso de entrega, **sin una sola falsa alarma** ; el peor puntuado se señala ocho turnos antes de incumplir su hito, mientras su propio responsable declara una urgencia de 0,001. La señal se debilita después, por una razón estructural y no por defecto : cuando la crisis se hace pública, la urgencia declarada alcanza a la calculada y una medida definida como su diferencia se cierra por construcción. Es un instrumento de alerta temprana, más nítido mientras una alerta aún vale algo.

La capa de predicción construida encima falló primero, y de forma precisa : mal calibrada, con veintiuna predicciones de media 0,948 nunca seguidas del resultado predicho según la definición de producción. Un solo bloque lo explicaba, el bloque temporal leyendo una magnitud que ningún choque alcanzaba. Ese diagnóstico se corrigió y se volvió a medir : la corrección evidente resultó errónea al medirla, pues el tiempo perdido debe *sumarse* a la fecha de finalización y no multiplicarse en ella, y tras la corrección el déficit de calibración se reduce a la mitad en todos los horizontes, de $-0,553$ a $-0,216$ a cuatro semanas. Sigue siendo negativo : la probabilidad calibrada todavía no es publicable como tal.

Un tercer resultado concierne al operador : en el turno en que la predicción se hizo visible, los ocho declarantes revisaron su urgencia a la baja o la mantuvieron, dos turnos antes del evento crítico. Lo que acota todo esto no es el diseño sino la muestra : ocho nodos, once hitos y un escenario. Ampliar esa base es el trabajo en curso y la condición de cualquier generalización.

Preamble

Three points about this document are worth stating before it begins.

Confidentiality. The project is designated confidential by ALTEN in the annex to the placement agreement. This report therefore carries a confidentiality marking on every page. It contains internal scoring parameters, the architecture of an unreleased software product and results not published elsewhere. It is submitted to ESTIA for examination only and should not be circulated further.

Language. The report is written in English, which the ESTIA guidelines permit and which matches the Cranfield double-degree track and the oral defence. There are two exceptions: the abstract is provided in French, English and Spanish, and the paragraph on ALTEN's corporate social responsibility policy is written in French as required. An English translation is placed immediately beside it so that no part of the document is unreadable to an English-speaking reader.

Contents

Acknowledgements	i
Résumé / Abstract / Resumen	ii
Preamble	v
List of Figures	viii
List of Tables	ix
1 Introduction	1
2 The Company	3
2.1 The ALTEN Group	3
2.2 The Direction de l'Innovation	3
2.3 Politique RSE d'ALTEN	4
3 Problem Statement	7
3.1 Why ALTEN Opened This Subject	7
3.2 The Cost of the Two Errors	7
3.3 The Problem, Stated	7
4 The Mission	9
4.1 The Assignment as Written	9
4.2 The Converged Objective	9
4.3 Objectives, Indicators and Deliverables	9
5 Project Organisation	11
5.1 Stakeholders and Responsibilities	11
5.2 Resources	11
5.3 Working Method: Agile and Test and Learn	12
5.4 Phases and Milestones	13
5.5 Risks and Mitigations	13
6 State of the Art	16
6.1 Deadline-Driven Value and Risk Propagation	16
6.2 Indicators of Delivery Rupture	16
6.3 Cognitive Bias in Declared Urgency	17
6.4 Multi-Criteria Elicitation and Aggregation	17
6.5 Digital Twins and Decision Support	17
6.6 Verifying a Probabilistic Signal	18
6.7 What the Review Establishes	18
7 Carrying Out the Mission	19
7.1 The Model	19
7.1.1 Real urgency	19
7.1.2 Declared urgency and the adequacy score	20
7.2 The System	21

7.3	The Experiment	24
7.4	Results	25
7.4.1	The adequacy signal identifies the right nodes, and does it early	25
7.4.2	The forecast layer, and the definition it is scored against	26
7.4.3	Why: one block explains all of it	28
7.4.4	What happens when the forecast is shown to the operator	28
7.4.5	The outcome definition moves the answer	29
7.5	The Repair, and What It Measures	30
7.5.1	The recommended repair was the wrong one	30
7.5.2	Two further defects the first measurement had hidden	31
7.5.3	What the re-measurement gives	31
7.5.4	A third arm, and the honest reading of it	32
7.6	Problems Encountered and How They Were Resolved	33
8	Conclusion and Personal Assessment	34
8.1	Results Against the Objectives	34
8.2	Limitations	35
8.3	Perspectives	35
8.4	Personal Assessment	36
8.4.1	What I learned	36
8.4.2	Link with the ESTIA engineering competency framework	37
8.4.3	What this mission clarified professionally	38
	Bibliography	39
A	Autorisation d'usage de l'IA par l'élève	44
B	Detailed Planning	45
C	Detailed Competency Evidence	46

List of Figures

1	The DISCO serious game network	4
2	The test-and-learn iteration loop	13
3	Phases and milestones of the mission	15
4	Causal structure behind the six-block selection	20
5	The pairwise comparison questionnaire and its consistency gate	21
6	Bidirectional propagation across the supply network	22
7	Declared and real urgency, and the two mismatch pathologies	22
8	The weekly review page	23
9	System architecture: single facade, audited writes, per-node persistence	23
10	Per-node explainability: exact decomposition of both signals	24
11	The campaign network and its per-arc coefficients	25
12	Reliability of the four-week forecast	27
13	Influence of the displayed forecast, by turn	29

List of Tables

1	Objectives, indicators and deliverables	10
2	Stakeholders and their role in the mission	11
3	Milestones	14
4	Risk register and outcomes	14
5	The six indicator blocks and what each measures	19
6	Nodes flagged by hidden risk against nodes that later missed a commitment .	26
7	Reliability of the four-week forecast	27
8	The same forecasts under three definitions of truth, four-week horizon . . .	29
9	The frozen campaign rescored after repair	31
10	Objectives against outcomes	34
11	Competencies exercised during the mission	37
12	Standing weekly meetings	45
13	Planned against actual, by phase	46

1 Introduction

A supply chain coordinator opens the week with a list of nodes and a finite amount of attention. Somewhere on that list is the supplier that will cause next month's line stoppage. Nothing on the screen says which one it is.

That is the practical shape of the problem this mission addressed. Modern supply chains are monitored through indicators that each measure one dimension at a time, while the decision the coordinator actually faces is multi-dimensional: cost against deadline against supplier reliability, carbon, and more. When the tooling cannot combine those dimensions, the coordinator combines them by judgment. Judgment is then the real decision layer of the system, and it is the layer nobody measures.

The consequences are not evenly distributed. Declaring more urgency than a situation warrants consumes shared transport and production capacity that another flow needed. Declaring less leaves a risk that nobody is watching, and by the time it becomes visible the window to act on it has closed. The second error is the more expensive, and it is the harder one to detect, precisely because an under-declared risk generates no direct signal.

This mission was carried out at the Direction de l'Innovation of ALTEN, within the Smart Green Supply Chain Digital Twin research programme, between 13 April and 16 October 2026. Its subject was the layer that sits before any planning decision: the urgency signal itself. The question driving the work was whether the gap between what an operator declares and what the operational state actually supports can be measured well enough, inside a working digital twin, to be acted upon, and to make a forecast built on it useful.

The answer this report reaches has two parts, and they point in different directions. The gap itself can be measured and it carries real information: on the campaign, the nodes the system flags in the six turns before the crisis becomes public are, without a single false alarm, nodes that go on to miss a delivery commitment, and the worst-scoring of them is flagged eight turns before it misses. That is the measurement the mission set out to build, and it works. Turning that signal into a calibrated probability of a named outcome in a named window is a further step, and that step did not work on the first attempt. The system was tested against recorded outcomes rather than asserted, and the test failed in a way that identifies exactly which component is responsible and what would repair it. That repair was then carried out and re-measured on the same frozen campaign, which halved the calibration deficit at every horizon without making the forecast good, and which corrected the recommendation as much as the model: the obvious fix turned out to be the wrong one. A second measurement, unplanned at the outset, showed that displaying the system's own forecast to the person making the declaration changed that declaration, and changed it in the wrong direction at the worst possible moment. All three findings are more useful to ALTEN than a favourable score would have been.

The report is organised as the mission itself was. Section 2 presents ALTEN, the Direction de l'Innovation and the research programme this work sits inside, including the company's corporate social responsibility policy and how the mission relates to it. Section 3 states the problem that justified opening the subject. Section 4 gives the assignment, its objectives, its indicators and its deliverables. Section 5 describes how the project was run: who was involved, what resources were available, which working method was applied, how the six months were phased, and which risks were carried. Section 6 reviews the existing methods

and identifies what none of them provides. Section 7 reports the work carried out and what it produced, including the problems met along the way and how they were resolved. Section 8 assesses the results against the objectives, sets out what remains to be done, and closes with a personal assessment against the ESTIA engineering competency framework.

2 The Company

2.1 The ALTEN Group

ALTEN is a French engineering and technology consulting group founded in 1988, with its head office in Boulogne-Billancourt. It sells engineering expertise rather than products: its consultants are placed on the technical projects of client companies, across the aerospace, defence, energy, transport, telecommunications, health and finance sectors. The group operates in more than thirty countries, and France remains its largest market.

Two features of that model matter for understanding this mission. First, ALTEN's asset is the technical capability of its people, so anything that raises the level of a whole practice has leverage that a single client project does not. Second, a consulting group sees the same problem recur across many clients, which puts it in an unusual position to notice that a difficulty is structural rather than particular. The urgency assessment problem treated in this report is of that kind.

2.2 The Direction de l'Innovation

The Direction Innovation Numérique (DIN) is the entity responsible for ALTEN's own research and development, as distinct from the work its consultants deliver to clients. It operates a network of ALTEN Labs in France and abroad, and brings together permanent researchers, doctoral candidates, consultants, interns and apprentices on multidisciplinary projects. Its output is not billable hours but reusable assets: methods, prototypes and expertise that consultants and clients can subsequently draw on.

This mission took place at the Toulouse Lab, within the Smart Green Supply Chain research programme. The programme pursues three aims simultaneously. The lean aim targets waste, meaning cost, delay and excess stock. The green aim targets the carbon footprint of each transport leg, working to the Global Logistics Emissions Council (GLEC) Framework v3.2 (Smart Freight Centre, 2023) and ISO 14083:2023 (International Organization for Standardization, 2023). The predictive aim, where this work sits, targets the anticipation of disruption.

The programme's central artefact is DISCO, a serious game and digital twin simulating a complete supply chain for a satellite product across 22 component suppliers (Baxas, 2025). Successive placements have built its components: the Physical Internet routing module (Baulain et al., 2024), the supplier risk scorer (Kabouche, 2025), the per-route carbon quantification (Cordova, 2024) and the transport management interface (Mondelice, 2025). The work described here adds a layer none of them covers, and it ran alongside a parallel placement on the network representation and geographic modelling side of the same programme, which is discussed in Section 5.1.

Figure 1 shows the supplier network players work with and the three dimensions their decisions are scored on. It is the environment the present work was designed to be compatible with, although the delivered system runs alongside DISCO rather than inside it, for reasons given in Section 4.2.




Figure 1 – The DISCO serious game network

2.3 Politique RSE d'ALTEN

Ce paragraphe est rédigé en français, conformément aux consignes de rédaction du rapport de MFE.

ALTEN a adhéré au Pacte Mondial des Nations Unies en 2010 et structure depuis sa politique de responsabilité sociétale autour de trois piliers : l'humain, l'environnement et l'innovation durable. La démarche n'est pas seulement déclarative. En 2025, le Groupe a obtenu un score EcoVadis de 85/100, niveau Platinum, qui le place dans le premier centile mondial, et a confirmé sa note A- au Carbon Disclosure Project. Sa trajectoire de réduction des émissions a été validée par la Science Based Targets initiative en 2023, avec un objectif de neutralité carbone à l'horizon 2050.

Sur le pilier humain, le Groupe réunit plus de cent nationalités et a signé les *Women's Empowerment Principles* des Nations Unies. L'index d'égalité professionnelle publié au titre de 2025 s'élève à 88/100 pour ALTEN Sud-Ouest, l'entité juridique dans laquelle s'est déroulée cette mission. Ce même document montre que la représentation féminine reste de 20 % parmi les cadres dirigeants d'ALTEN Sud-Ouest et de 14 % au sein des instances dirigeantes, écart que le Groupe reconnaît publiquement et qu'il entend traiter par le programme INSPIRE 2026-2030. Il paraît plus honnête de citer cet écart que de n'en retenir que l'index.

Sur le pilier environnemental, les actions internes sont concrètes et vérifiables : 100 % d'électricité renouvelable pour les bâtiments français dont la puissance souscrite dépasse

36 kVA, 80 % des surfaces couvertes par le tri sélectif, et une certification ISO 14001 en cours d'extension. La sensibilisation passe par des ateliers Fresque du Climat et des formations aux éco-gestes.

C'est néanmoins le troisième pilier qui éclaire le mieux cette mission. ALTEN consacre 31 % de sa R&D à l'innovation environnementale, et le programme dans lequel ce stage s'est inscrit en fait partie : un programme nommé *Smart Green Supply Chain*, où les émissions de CO₂ constituent l'un des six blocs du modèle d'urgence développé ici, au même rang que le coût ou le délai, et non une correction appliquée après coup. Autrement dit, la politique RSE du Groupe n'a pas seulement encadré ce travail : elle en a déterminé une variable de modélisation.

Prise de recul. Deux observations méritent d'être faites. D'abord, la RSE d'ALTEN est nettement plus visible au niveau du Groupe qu'au niveau du quotidien d'un laboratoire : en six mois, je l'ai rencontrée davantage sous la forme d'un objet de recherche que sous celle d'une pratique de bureau. Ensuite, et c'est le point qui m'a le plus marqué, le bloc carbone du modèle s'est révélé le moins sollicité des six lors de la campagne expérimentale. Une entreprise peut donc afficher une politique environnementale solide, financer un programme de recherche qui la porte, et produire malgré tout un modèle où la dimension environnementale reste la moins active, faute de données pour l'alimenter. L'écart ne tient pas à l'intention mais à la disponibilité de la mesure. Il est signalé comme tel en Section 8.2.

English translation of the above. ALTEN joined the United Nations Global Compact in 2010 and structures its corporate social responsibility policy around three pillars: people, environment and sustainable innovation. The approach is measured rather than merely stated. In 2025 the group scored 85/100 on EcoVadis, reaching Platinum level and the top one per cent worldwide, and confirmed its A- rating with the Carbon Disclosure Project. Its emissions reduction pathway was validated by the Science Based Targets initiative in 2023, targeting carbon neutrality by 2050.

On the people pillar, the group spans more than a hundred nationalities and is a signatory of the UN Women's Empowerment Principles. The 2025 gender equality index stands at 88/100 for ALTEN Sud-Ouest, the legal entity in which this mission took place. The same publication shows female representation still at 20% among senior executives of ALTEN Sud-Ouest and 14% on its governing bodies, a gap the group states publicly and intends to address through its INSPIRE 2026-2030 programme. Citing that gap seems more honest than citing only the index.

On the environmental pillar the internal measures are concrete and checkable: 100% renewable electricity for French buildings above 36 kVA of subscribed power, 80% of floor area covered by waste sorting, and ISO 14001 certification being extended. Awareness is raised through Fresque du Climat workshops and eco-practice training.

It is the third pillar that bears most directly on this mission. ALTEN devotes 31% of its R&D to environmental innovation, and the programme hosting this placement is part of that share: a programme named *Smart Green Supply Chain*, in which CO₂ emissions form one of the six blocks of the urgency model developed here, ranked alongside cost and delay rather than applied as an afterthought. The group's CSR policy did not merely frame this work; it

determined one of its modelling variables.

Two observations follow. ALTEN's CSR is considerably more visible at group level than in the daily life of a laboratory: over six months I encountered it more as an object of research than as an office practice. And the carbon block of the model turned out to be the least exercised of the six during the experimental campaign. A company can therefore hold a solid environmental policy, fund a research programme that carries it, and still produce a model in which the environmental dimension is the least active one, for want of data to feed it. The gap is not one of intent but of measurement availability, and it is recorded as such in Section [8.2](#).

3 Problem Statement

3.1 Why ALTEN Opened This Subject

The Smart Green Supply Chain programme had, by the start of this placement, built a credible set of components: a routing module, a supplier risk scorer, a carbon quantification module, a transport interface, and the DISCO twin that ties them together. Each answers a question of the form *what is the state of this part of the chain?*

None answers the question that comes next. Once the state is known, a coordinator still has to decide where to spend a finite amount of attention this week, and that decision is made on a signal none of those modules produces: how urgent each node is. In practice that signal comes from people. A buyer says a supplier is critical; a plant says a line is at risk. The digital twin records the consequence of those judgments without ever evaluating the judgments themselves.

That is a gap with a cost attached. Behavioural research on urgency shows the signal is not neutral. Zhu, Yang and Hsee demonstrate across five controlled experiments that people systematically favour imminent low-value tasks over distant high-value ones, and that the bias persists even when they are told about it (Zhu et al., 2018). Field studies find the same pattern under workload: operators complete small easy tasks first, at a measurable cost to overall performance (Rusou et al., 2020; DiwasSingh et al., 2020). If a digital twin consumes a human urgency declaration as ground truth, it inherits that bias and then propagates it across the network.

3.2 The Cost of the Two Errors

The two directions of error are not symmetric, which is what makes the problem worth instrumenting rather than tolerating.

Over-declaration is expensive but visible. A node declared more urgent than it is attracts expedited transport, premium capacity and coordinator attention, all of which are finite and all of which were taken from somewhere else. In a shared or mutualised logistics setting the cost lands on another actor entirely (Montreuil, 2011; Li et al., 2020).

Under-declaration is cheaper in the moment and far more dangerous. A node whose real situation is worse than anyone has said generates no alert, appears in no priority list and consumes no attention. The failure surfaces only when it becomes a missed delivery, by which point the window in which it could have been prevented has closed. Clinical triage, where this asymmetry has been studied for decades, reaches the same conclusion: under-triage costs a life where over-triage costs a queue (Kim et al., 2020; Mackway-Jones et al., 2013).

3.3 The Problem, Stated

Problem statement.

A supply chain digital twin can compute the physical state of every node, and it can record the urgency a human declares for each of them. It does not compare the two. The declared signal is known to be biased, the computed signal is available, and the difference

between them is exactly the quantity that would tell a coordinator which node is being mismanaged rather than merely which node is struggling.

Can the gap between declared and computed urgency be measured reliably enough, inside a working digital twin, to be acted upon rather than merely observed?

The word *reliably* carries the weight. Producing a number that expresses the difference between two other numbers is trivial. Establishing that the number means something, that it predicts outcomes better than the raw indicators a coordinator already has, and that it is not simply measuring its own inputs, is the actual engineering problem. Section 7 shows that this distinction is not academic: the system produced the number without difficulty and then failed the reliability test, for a reason that took a dedicated experiment to identify.

4 The Mission

4.1 The Assignment as Written

The placement agreement defines the subject as *Stage Innovation: Ingénieur Supply Chain*, and the assignment as the update and development of a supply chain digital twin. The signed scope covers five workstreams: defining the twin's elements, meaning factory, storage, transport and cargo; updating its evaluation elements, meaning score, cost, quality, delay, CO₂ and planning; updating its evolution elements, meaning positive and negative hazards and change scenarios; integrating complexity through multiple supply chains and tier-2 and tier-3 subcontractors; and integrating Physical Internet elements. It further specifies that the work is an innovation project including a state of the art, formalisation of the research in a scientific report, multidisciplinary project management, and an evolutionary approach based on Agile and Test and Learn.

That scope is deliberately broad, which is normal for an R&D assignment and is itself a managed risk (Section 5.5). The first weeks were therefore spent converging it onto a defensible contribution, in agreement with the research programme director.

4.2 The Converged Objective

The evaluation workstream is where the twin was weakest, because it computed state without ever assessing the human judgment applied to that state. The mission converged on that workstream, and specifically on the urgency dimension of it.

Converged objective. Design, implement and empirically evaluate a decision-support layer that computes a real urgency from the twin's operational state, captures the urgency declared by the operator through a structured elicitation, propagates both across the supplier network, and scores the gap between them in a way that can be tested against recorded outcomes.

Two boundaries were set explicitly and held. The layer is built as a standalone application running alongside DISCO rather than inside it, because coupling to a Java and Angular platform under active development by another team would have made the schedule dependent on theirs. And the layer targets urgency alone, not the scheduling decision that follows from it, because a prioritisation signal that cannot be shown to be trustworthy has no business driving a schedule.

4.3 Objectives, Indicators and Deliverables

Table 1 states the objectives agreed with the programme director, each with the indicator used to judge it and the deliverable that carries it. The indicators were chosen so that each could be answered yes or no from an artefact rather than from an opinion.

Intermediate deliverables. Four were produced during the mission rather than at its end: a knowledge transfer report on planning and urgency, delivered in week 2 and used to align vocabulary with the programme; a state of the art with an accompanying presentation, delivered in week 3; the mathematical specification, reviewed before implementation began;

Table 1 – Objectives, indicators and deliverables

	Objective	Indicator	Deliverable
O1	Define a real urgency computable from heterogeneous operational indicators, and an adequacy score against the declared value	A model whose every term is decomposable, and whose aggregation choice is justified against the alternatives	Mathematical specification
O2	Deliver a working, persistent, auditable application implementing it	Runs multi-tenant; every write audited; automated test suite passing	SupplyScore application
O3	Propagate urgency across a multi-tier network and rank nodes by criticality	Handles a network at the scale of a real programme within an interactive time budget	Propagation and criticality services
O4	Establish whether the resulting signal predicts recorded outcomes	A score against a stated outcome definition, compared with a trivial baseline	Pre-registered experimental protocol and its results
O5	Capitalise the research so the programme can continue it	A written scientific record and reproducible derivation of every number	<i>Bilan Scientifique et Technique</i> and derivation scripts

and the frozen experimental protocol, hashed and registered before any campaign data existed, so that the hypotheses could not be adjusted after seeing results.

Final deliverables. The SupplyScore application with its test suite and documentation; the experimental campaign with its two executed arms and their analysis; the *Bilan Scientifique et Technique*; and the derivation scripts that regenerate every reported figure from raw artefacts.

One objective was reformulated during the mission. O4 as originally agreed asked whether the signal predicts outcomes, with an implicit expectation that it would. The campaign showed it does not, and identified why. The objective was therefore met in the sense that the question was answered with evidence, and not met in the sense that the answer is negative. Section 8.1 treats that distinction directly rather than presenting a failed prediction as a success.

5 Project Organisation

5.1 Stakeholders and Responsibilities

The mission was carried out inside a small research team, and the distribution of responsibility mattered more than its size suggests. Table 2 lists who did what.

Table 2 – Stakeholders and their role in the mission

Person	Role	Responsibility in this mission
Sébastien Perthuisot	Research Programme Director, Quality & Supply Chain Expert	Set the research direction, arbitrated scope, chaired the weekly regular point and the Friday demonstration. Decision-maker on what the programme keeps.
Émilienne Lardy	Scientific supervision	Ran the Monday planning and the three weekly scrums; reviewed method and written output. First reader of anything claimed.
Selma Ferhat	Scientific supervision	Reviewed the scientific framing and the written record.
Fabien Jules	Chargé Innovation	Placement tutor of record; onboarding, administrative framing and access.
Liam Jones	Fellow intern, same programme	Network representation, graph modelling and geographic layer of the twin. Neighbouring problem, shared demonstration slot.
David Gomez	ESTIA academic supervisor	Academic framing of the mission.
Author	Intern, Supply Chain Engineer	Model, implementation, experimental design, analysis and written deliverables.

The presence of a second intern on an adjacent problem shaped the work more than expected. Liam Jones was building the structural and geographic representation of a supply chain, meaning how a network is described, generated and visualised, while this mission built the layer that scores nodes once they exist. The two are complementary halves of the same twin, and the Friday demonstration forced each to state weekly what the other could rely on. It also produced a boundary decision early: the scoring layer would assume a network is given rather than attempt to generate one, because that was being solved two seats away.

5.2 Resources

The resources provided were modest and are worth stating plainly, because two of them constrained technical choices.

- *Human*. No budget-holding team. Supervision as listed above, and one peer on an adjacent subject. All implementation was carried out by the author.
- *Material*. One laptop provided by ALTEN, a Ryzen 7 PRO 8840HS with 16 cores and no discrete GPU. Everything computed in this mission, including every model training run, ran on that CPU.
- *Software*. Azure DevOps for source control and work tracking, and the Microsoft suite for communication. No paid data service, no cloud compute budget, no commercial

- solver or licence was requested or used, which is why the delivered system depends only on open-source Python libraries and a file-based database.
- *Financial.* None directly managed. The absence of a budget is itself a design constraint: it removed managed cloud services and paid datasets from the solution space before any technical comparison took place.

5.3 Working Method: Agile and Test and Learn

The placement agreement prescribes an evolutionary approach based on Agile and Test and Learn. That was applied literally rather than nominally, and the weekly rhythm was the mechanism.

The cadence. The week had a fixed shape. A planning session on Monday set the objective for the week. A scrum on Tuesday, Wednesday and Thursday morning kept it visible and surfaced blockages within a day rather than a week. A regular point with the programme director on Wednesday took the scientific decisions. A demonstration on Friday closed the loop by requiring something to be shown, not described. Four contact points a week on a six-month mission is a heavy cadence for R&D, and it was the single most useful organisational feature of the placement: it made it impossible to spend three weeks on an idea that would not survive contact with a reviewer.

Test and Learn, with what it actually rejected. The method is only meaningful if things were genuinely abandoned. Four were.

1. *The first urgency formulation was dropped.* The initial model expressed urgency as a single time-driven function with a finite-time singularity, calibrated by penalty terms. It was specified, partially implemented and abandoned in favour of a multi-block aggregation, because a single deadline-proximity curve discards information the twin already holds and cannot represent a node that is on time but out of capacity.
2. *The May module architecture was replaced.* A first pipeline of separate scores was built through May and superseded in June by the aggregation and propagation design that shipped. The intermediate work was not wasted: it established what the interfaces between components needed to be, which is why the replacement was built in weeks rather than months.
3. *A candidate dataset was rejected on inspection.* A third-party logistics dataset considered for the freight node was checked before use and discarded: it carried pre-computed target columns, coordinates scattered outside the region it claimed to cover, and uniformly flat distributions, all signatures of synthetic data. The node was driven by calibrated events instead. The rejected file is retained for audit and never loaded by the analysis pipeline.
4. *A whole experimental arm was reduced.* The campaign was designed for eight human consultants over nine working days. When that could not be secured, the design was changed rather than the claim: the campaign was run with isolated agents, and every conclusion about human behaviour was correspondingly weakened in the reporting.

Figure 2 shows the loop as it was actually run. Each turn of it produced either a working increment or a documented rejection, and the four rejections above are the evidence that the second outcome was genuinely available rather than theoretical.




Figure 2 – The test-and-learn iteration loop

Quality gates. Agile without a definition of done produces velocity and nothing else. Three gates were applied: no feature merged without automated tests, giving 1779 test functions at the end of the mission; every business write passing through a single audited mutation service, so that no result could be produced by an untraceable edit; and every reported figure regenerated from raw artefacts by a script rather than transcribed by hand.

5.4 Phases and Milestones

The mission ran in six phases. Figure 3 gives the schedule and Table 3 the milestones that closed each phase.

The planning was revised once, in June. The original schedule allocated the second half of the mission to extending the model's scope; it was reallocated to validating what already existed, on the argument that an unvalidated extension of an unvalidated model compounds the problem rather than advancing it. That reallocation is the reason this report can state a measured result at all.

5.5 Risks and Mitigations

Four risks were identified and carried at the start of the mission. Table 4 states each one, whether it materialised, and what was done about it. Three of the four did materialise, which is a high rate and reflects how much of this mission depended on things outside the author's control.

Table 3 – Milestones

Date	Milestone	Closed by
End April	Framing complete	Knowledge transfer report and state of the art presented; subject converged with the programme director
End May	First pipeline	End-to-end computation running, and found insufficient, which triggered the redesign
Mid June	Application delivered	Multi-tenant persistence, audited writes, weekly cycle, event engine, explainability, exports and backup
Early July	Protocol frozen	Scenario, data pack hashed, role cards and hypothesis register registered before any data existed
End July	Campaign executed	Two arms run, analysed against the pre-registered register
Mid August	Research capitalised	Scientific report and reproducible derivation of every figure

Table 4 – Risk register and outcomes

Risk	Outcome	Mitigation and effect
No access to academic literature from within ALTEN	Materialised	The Direction de l'Innovation holds no journal or database subscriptions, so a state of the art conducted from the company alone would have rested on whatever is publicly readable. The mission depended throughout on ESTIA and Cranfield library credentials. It worked, but the dependency is real and would block a consultant without a university affiliation.
Broad mission scope	Materialised	Five workstreams in the signed assignment against one person and six months. Mitigated by converging the subject in the first weeks with the programme director and recording the boundaries explicitly (Section 4.2) rather than leaving them implicit.
Human campaign might not be securable	Materialised	Eight colleagues for nine working days was never guaranteed. Mitigated by designing the protocol so that the declarant could be substituted without changing the record format, which is what allowed the campaign to run at all when the humans were unavailable. Cost: the external validity of every behavioural claim.
CPU-only compute	Contained	No GPU available. Bounded model size and training resolution, and is the reason one comparison was run at reduced scale. Mitigated by choosing methods whose cost is dominated by data rather than by parameters.




Figure 3 – Phases and milestones of the mission

The first risk deserves a comment, because it is the one least often written down. An R&D department that cannot read the literature it is meant to advance is dependent on the academic affiliations of the people passing through it. For a placement holder on a double degree that is invisible, since the credentials come with the student. It stops being invisible the moment the work has to be continued by someone without them, which is a succession risk for the programme rather than a problem for this mission.

6 State of the Art

The question of Section 3.3 touches several fields that do not usually cite one another. This section reviews each for what it settles and, more importantly, for what it leaves open. A structured evaluation of every source, recording its argument, method and limitations, was maintained throughout the mission and informs the appraisals below.

6.1 Deadline-Driven Value and Risk Propagation

Real-time systems research models the value of a task as a function of when it completes. Chen and Mühlethaler show that utility typically holds constant until a deadline and then decreases (K. Chen et al., 1996), and a wider survey confirms structured degradation past critical thresholds (J. Chen et al., 2025). This establishes that urgency is legitimately time-varying.

Separately, supply chain research formalises the ripple effect, whereby a local disruption cascades through downstream tiers and amplifies as it travels (Ivanov; Dolgui, 2020; Ivanov; Dolgui, 2019; Li et al., 2020). This establishes that urgency is legitimately network-dependent.

Critical appraisal. The first stream stops at a single isolated task and offers nothing on combining heterogeneous indicators. The second models propagation but treats the node-level score as given. Neither addresses a declared signal at all. Between them they justify treating urgency as time-varying and propagated, and leave the aggregation question entirely open.

6.2 Indicators of Delivery Rupture

If urgency is to be aggregated, the indicators must come from somewhere. Heckmann, Comes and Nickel establish rupture as the terminal event a supply risk measure should anticipate, and criticise the field for measuring exposure without connecting it to that event (Heckmann et al., 2015). Indicator-based supplier monitoring proceeds by identifying observable precursors (Mokhtar et al., 2019).

Five families recur. Slack is the only indicator whose exhaustion is irreversible (Cohen et al., 2019; Sun et al., 2020). Capacity saturation adds waiting time a nominal lead time does not show (Tomlin, 2006; Hosseini; Tajik, et al., 2019). Resource performance is monitored through composite proxies because supplier-side data is rarely available (Mokhtar et al., 2019; Hlioui et al., 2017). Adverse-event exposure is treated probabilistically (Hosseini; Ivanov, et al., 2020; Liu et al., 2021). Cost signals are read as symptoms of strain rather than causes (Shen et al., 2019), and emissions constraints restrict the admissible responses rather than causing rupture (Hoen et al., 2014; Smart Freight Centre, 2023).

Critical appraisal. This literature names precursors and monitors them individually. It supplies no criterion for deciding which belong in one composite, and says nothing about whether equal numerical values on structurally different indicators carry equal risk meaning. Both questions are inherited by any multi-block model, and Section 7.1 treats the second explicitly.

6.3 Cognitive Bias in Declared Urgency

Zhu, Yang and Hsee report the mere-urgency effect across five controlled experiments: participants consistently preferred imminent low-value tasks to distant high-value ones, and the bias survived being told about it (Zhu et al., 2018). Steel gives a formal discounting model in which immediacy becomes more salient as a deadline approaches (Steel et al., 2006). Field studies confirm the pattern under workload (Rusou et al., 2020; DiwasSingh et al., 2020).

Critical appraisal. This stream diagnoses the problem and quantifies it in controlled settings, but the magnitude in industrial supply chains is uncalibrated, and it prescribes no operational correction. It justifies treating the declared signal as biased; it does not say what to do about it.

6.4 Multi-Criteria Elicitation and Aggregation

Structured elicitation of a human judgment is a solved problem in outline. The Analytic Hierarchy Process derives criteria weights from pairwise comparisons and validates them through a consistency ratio (Saaty, 1980), though it is vulnerable to rank reversal and to cognitive noise (Dyer, 1990; Triantaphyllou et al., 2024). The Best-Worst Method reduces the number of comparisons required (Rezaei, 2015), and its fuzzy extension absorbs judgment uncertainty (S. Guo et al., 2017; Liang et al., 2020). For ranking alternatives without letting a good score on one criterion mask a catastrophic one on another, outranking methods such as PROMETHEE II establish pairwise preference flows (Brans et al., 1986; Behzadian et al., 2010).

The composite indicator literature supplies the necessary warning: nominal weights often diverge from effective importance because of correlation and normalisation effects (Greco et al., 2018; Becker et al., 2017; Borgonovo et al., 2021).

Critical appraisal. These are static tools. They elicit and aggregate, and none is routinely compared against a computed operational counterpart. The composite-index critique is generally raised as a caveat at the end of a paper rather than turned into something a running system checks on its own data.

6.5 Digital Twins and Decision Support

The digital twin concept was formalised by Grieves (Grieves, 2014) and extended to manufacturing (Tao et al., 2018; Fuller et al., 2020). In supply chains it integrates simulation, optimisation and machine learning (Ivanov; Dolgui, et al., 2021), with demonstrated capability in disruption monitoring and replanning (Ding et al., 2023; Leung et al., 2022; Hrušovský et al., 2021). Recent human-in-the-loop cognitive twins capture human feedback to synchronise mental models with simulated operations (Niloofar et al., 2023).

Critical appraisal. This is the family the delivered system belongs to, and the gap is here. No retrieved system combines a declared urgency input, a computed operational urgency and their propagated mismatch. A second gap matters for this mission specifically: these systems are evaluated on how accurately they estimate physical state, and their inputs are assumed to arrive. Where an input depends on the same operator whose judgment the system exists to support, that circularity is not treated as a measurement problem.

6.6 Verifying a Probabilistic Signal

Because this mission scores a probabilistic signal against outcomes, it inherits forecast verification. Brier introduced the mean squared error of probabilistic forecasts (Brier, 1950), and Murphy decomposed it into reliability, resolution and uncertainty (Murphy, 1973), establishing the distinction the results of this report depend on: a forecast can order outcomes correctly while announcing the wrong probabilities, and the two failures need different repairs. Machine learning rediscovered the same point, finding strong classifiers that produce distorted probabilities correctable by a cheap post-hoc layer (Niculescu-Mizil et al., 2005; C. Guo et al., 2017).

Critical appraisal. This literature assumes the outcome being predicted is unambiguously defined. In an operational supply chain it is a modelling choice, and Section 7.4 measures that choice moving the reported result by a factor of five.

6.7 What the Review Establishes

Each stream closes on a limit. Collected, those limits describe one absence rather than several. The reviewed work covers multi-block risk aggregation, network propagation, digital twin operation, multi-criteria elicitation and probabilistic verification. No retrieved work combines them with a structured declared-urgency input and an empirical test of the resulting adequacy signal against recorded outcomes.

That missing combination is what this mission set out to build and to test.

7 Carrying Out the Mission

This section follows the phases of Section 5.4: the model, the system built from it, the experiment designed to test it, what the experiment measured, and the problems met along the way.

7.1 The Model

7.1.1 Real urgency

Real urgency is not read from a single deadline-proximity signal. It is aggregated at each node from six indicator blocks selected from the literature of Section 6: time, capacity, performance, risk, cost and carbon. Each produces a partial urgency $u_m \in [0, 1]$, and the aggregation is a weighted probabilistic OR:

$$U_{r,\text{local}} = 1 - \prod_m (1 - u_m)^{\omega_m}. \quad (1)$$

The choice of that form over a weighted average is the central modelling decision. A weighted average is fully compensatory: a node with a completely failed capacity and healthy cost, carbon and performance readings would score as moderately urgent. Equation 1 is not: any block approaching its maximum drives the aggregate toward the maximum regardless of the others. For a signal whose purpose is to catch the node that is about to fail, non-compensatory behaviour is the required property.

Table 5 – The six indicator blocks and what each measures

Block	Role	What it measures
Time	Direct cause	Probability that the realistic lead time exceeds the margin to the next deadline. The only resource whose exhaustion cannot be reversed
Capacity	Direct cause	Saturation of storage, weight and throughput, which adds waiting time a nominal lead time does not show
Performance	Direct cause	Drift in the productive health of the resource making the part, through a composite availability, performance and quality proxy
Risk	Direct cause	Expected unavailability: failure probability times recovery time times severity, revised weekly from declared events
Cost	Arbitration	Financial strain. Paying premium freight to protect a deadline is a symptom, not a cause, and it bounds which responses stay affordable
Carbon	Arbitration	Emissions above target. Restricts the admissible set of responses, typically in the sea-against-air trade-off

Table 5 gives the six and separates them by role. Four describe a mechanism that leads to rupture; two constrain which responses remain available once it threatens. Each candidate indicator was screened against four questions before being retained: does it have a plausible causal link to rupture, can it be measured or credibly estimated, can an operator act on it, and does it add information the retained blocks do not already carry. Figure 4 draws the resulting structure.




Figure 4 – Causal structure behind the six-block selection

Three properties of that equation were specified explicitly, because a bare formula does not state what its bounds mean. At $u_m = 0$ the block is nominal and its factor is 1, leaving the product unchanged. At $u_m = 1$ the factor is 0 and the aggregate saturates. A block whose data is unavailable is given weight $\omega_m = 0$, which is deliberately distinct from $u_m = 0$: the first says nothing is known, the second asserts the block is healthy.

The assumption that the blocks are exchangeable. The product is symmetric in its factors, so it treats the six blocks as interchangeable up to their weights. That does not hold cleanly. A value of $u_{\text{time}} = 0.3$ is a probability of missing a deadline; $u_{\text{cost}} = 0.3$ is an ordinal position on a normalising scale wearing a cardinal number. They are not equally informative about rupture. The consequence was carried rather than hidden: the aggregate is used for ranking and for comparison against the declared value, never as a calibrated probability in its own right.

7.1.2 Declared urgency and the adequacy score

Declared urgency is not collected as a free-text estimate. At each node the operator compares four criteria pairwise through AHP, which yields a consistency-checked weight per criterion, then rates each on a severity scale. The four criteria were chosen so that each has a computed counterpart: operational impact and downstream dependencies mirror the severity the

propagation amplifies, the time window mirrors the time block, and recoverability mirrors the recovery term of the risk block. Without that correspondence the later comparison would be between unrelated quantities.

Figure 5 – The pairwise comparison questionnaire and its consistency gate

Figure 5 shows the questionnaire as the operator meets it. A weight vector whose consistency ratio exceeds the accepted threshold is refused, and the operator is sent back to the most contradictory of their own comparisons before the assessment can enter the network. Without that gate an incoherent judgment would propagate as though it were a considered one.

Both signals then propagate across the supplier network, and in opposite directions: declared urgency descends from the customer toward suppliers, real risk ascends from suppliers toward the customer. The asymmetry is deliberate. Demand pressure is something a customer transmits to its suppliers; physical risk is something suppliers transmit to their customer. A model that propagated both the same way would be describing neither.

Figure 6 shows the two signals travelling in opposite directions across the same graph.

Their mismatch, once both are propagated, splits into two named pathologies: false urgency where the declaration exceeds the state, hidden risk where the state exceeds the declaration. It is scored on $[0, 100]$ through an asymmetric penalty that prices hidden risk at roughly twice false urgency, following the clinical principle of Section 3.1.

Figure 7 shows the two signals over the life of a task. The step line is the declaration, which moves only when a review is submitted; the smooth curve is the computed state, which moves continuously. Where the step line sits above the curve the enclosed area is false urgency; where the curve sits above the step line it is hidden risk.

7.2 The System

The model was delivered as SupplyScore, a standalone Python application of 73 modules across 10 packages, covered by 1831 automated test functions in 106 test files, driving a multi-page web interface over a per-node database layer. The suite collects 2 044 cases once parameterisation is expanded, of which 2 028 pass, 15 are skipped for want of an optional dependency, and one fails on two missing screenshot files referenced by a presentation script outside the library.




Figure 6 – Bidirectional propagation across the supply network




Figure 7 – Declared and real urgency, and the two mismatch pathologies

Four architectural decisions carry the engineering weight. Every operation passes through a single orchestrating facade, so there is no second write path. Data is split into one registry database plus one database per node, which enforces that one supplier's history is not readable through another's. Every business write crosses a single mutation service that validates ranges, diffs against current state and records the change with its author, making the audit trail a property of the architecture rather than a convention. And each project carries its own clock, which is what allows a six-month scenario to be replayed in compressed time without touching the model.




Figure 8 – The weekly review page

Figure 8 shows the operator's main entry point. Its four sections, the questionnaire, the week's indicator values, milestone progress and any events, commit as a single signed transaction so that no partial review can enter a node's history.




Figure 9 – System architecture: single facade, audited writes, per-node persistence

Figure 9 shows how those pieces relate. The single entry point is what makes the audit guarantee enforceable rather than aspirational: there is no path into the data that bypasses it, including from the administration interface.

Two capabilities proved more important than anticipated. The explainability page reconstructs any score term by term, decomposing the declared value into one contribution per criterion and the computed value into one per block, so that a coordinator asking why a node scored as it did receives an exact answer rather than a plausible one (Figure 10). That is what made

the diagnosis in Section 7.4 possible from the application rather than by hand. The export produces frozen data for analysis outside the tool, which the whole experimental analysis relies on.




Figure 10 – Per-node explainability: exact decomposition of both signals

7.3 The Experiment

Building a system that computes a number establishes nothing about whether the number means anything. The mission therefore included an experiment designed before the system was finished.

The campaign replays the 2020 to 2022 semiconductor shortage over nineteen turns across an eight-node satellite supply chain, from a raw wafer supplier through a foundry, a distributor, a freight operator and an electronics manufacturer to the satellite prime. Real urgency is fed from public time series of the actual crisis; declared urgency comes from the participants through the weekly questionnaire. The two signals therefore have independent sources, which is the precondition for the gap between them to mean anything.

Figure 11 gives the network, with the coefficient attenuating demand as it descends and the one amplifying risk as it ascends, and each node's tier. The foundry at the deepest tier is the expected bottleneck. The dashed arc into the secondary foundry is deliberately inert: it represents the announced alternative supplier that in the real crisis had no spare capacity, and modelling it as present but carrying nothing is the point.

Four artefacts were frozen before any data existed: the scenario with its twelve calibrated events, the data pack with a hash recorded per file, the role cards, and a register of six hypotheses each carrying a numeric threshold. Freezing the thresholds in advance is what prevents a disappointing result from being rescued by moving the bar afterwards.




Figure 11 – The campaign network and its per-arc coefficients

The design compares two arms on the identical frozen scenario. In the control arm, eight isolated declarants each see only their own role card and turn briefing and declare blind. In the treatment arm, the same declarants are additionally shown the system's own predictive analysis of their node before declaring, and record whether it influenced them. Because the scenario is frozen, the forecasts are identical between arms, so any behavioural difference belongs to the act of displaying them.

7.4 Results

The results divide cleanly into what the system was built to do, which it does, and the predictive layer built on top of it, which was diagnosed, repaired and measured as improving without yet being good. They are reported in that order.

7.4.1 The adequacy signal identifies the right nodes, and does it early

The first result is the one the system was built for, and it is positive.

A coordinator does not ask whether node X will suffer an adverse event between weeks $t + 1$ and $t + 4$. They ask which of their suppliers deserves attention this week. That is a question about nodes, not about node-weeks, and the adequacy signal answers it. At each turn the eight nodes were ranked by hidden risk, the part of the gap where computed urgency exceeds declared urgency, and the set flagged above 0.10 was compared with the set of nodes that would go on to miss a milestone at its originally committed deadline. Four of the eight nodes did.

Table 6 gives it. For six consecutive turns, from the opening turn to the one before the foundry freeze, the signal flags exactly three nodes, all three of which go on to miss a commitment, with no false alarm at any turn. The fourth is never flagged early, so recall stops at three of four. And on every one of those six turns the worst-scoring node of the eight is the satellite prime, which misses its delivery deadline at turn 8 and delivers at turn 14: the flag precedes the miss by eight turns, and it is raised while that node's own manager is declaring an urgency of 0.001.

That last detail is the whole point of the construction rather than a happy accident. Nothing

Table 6 – Nodes flagged by hidden risk against nodes that later missed a commitment

Turns	Precision	Recall	False alarms	Reading
T0 to T5, before the crisis is visible	1.00	0.75	0	Three nodes flagged, all three later missed
T6 to T18, crisis public	0.50 or below	0.25	1 to 3	The declared signal catches up and the gap closes

in the flag comes from the declaration being pessimistic; it comes from the declaration being optimistic while the computed state is not. A system that only aggregated indicators would have ranked that node by its indicators. A system that only collected declarations would have believed it. The gap between the two is what carries the information, which is what the mission set out to test.

The signal then degrades from turn 6 onward, and the reason is structural rather than a defect. Once the crisis is public every declarant raises their urgency, declared catches up with computed, and a signal defined as the difference between the two necessarily shrinks. Precision falls to 0.50 and recall to 0.25. This is the correct behaviour of a perception-gap measure and it bounds where such a measure is useful: it is an early-warning instrument, not a crisis-management one. Stated the other way round, the signal is at its sharpest exactly when a warning is still worth something.

The repair of Section 7.5 sharpens the same flag without changing which nodes it selects. On the campaign as run, the satellite prime scores 26.6 out of 100 at turn 3 against 46.0 for the next worst node; after repair the same turn gives 2.2 against 44.6. The ordering is identical and the separation is an order of magnitude wider, which is what makes the flag readable on a dashboard rather than merely present in the data.

One limit bounds all of this and is the reason Section 8.3 asks for more data before anything is generalised: eight nodes, four of which are positive, on one scenario. A precision of 1.00 over six turns on three flagged nodes is an encouraging observation and not an established rate.

7.4.2 The forecast layer, and the definition it is scored against

The adequacy signal ranks nodes. Turning it into a probability that a named outcome occurs within a named window is a further step, and that step is where the campaign found its defect. Scored against the outcome definition the system itself applies, meaning a missed milestone, an abandoned node or a critical event within four weeks, the forecast layer failed on both axes that matter.

Discrimination is indistinguishable from chance: the area under the ROC curve ranges from 0.37 to 0.56 across horizons and every confidence interval contains 0.5. That statement carries a caveat, since the positive class holds between two and six observations, which is too few to conclude in either direction.

Calibration fails much more clearly, and does not depend on that small sample. The Brier skill against a forecaster that simply announces the base rate is between -3.4 and -6.9 , meaning the model is several times worse than the trivial baseline. Table 7 localises the failure.

Table 7 – Reliability of the four-week forecast

Announced band	<i>n</i>	Mean announced	Observed	Reading
[0.0, 0.1)	74	0.070	0.081	Well calibrated
[0.1, 0.2)	24	0.128	0.000	None materialised
[0.2, 0.5)	1	0.354	0.000	None materialised
[0.5, 0.9)	7	0.859	0.000	None materialised
[0.9, 1.0]	14	0.992	0.000	None materialised

The pattern is unambiguous, and Figure 12 shows it, with each marker's area proportional to the number of forecasts in its band and the diagonal marking perfect calibration. Where the model announces a low probability it is right: 74 forecasts averaging 0.070 were followed by an adverse outcome 8.1% of the time. Where it announces a high probability it is wrong every time: twenty-one forecasts announcing a mean of 0.948 were followed by the predicted outcome zero times.




Figure 12 – Reliability of the four-week forecast

Separating these two failures matters for what happens next, and it is the distinction Murphy established for exactly this purpose (Murphy, 1973). A forecast that orders outcomes correctly but announces the wrong numbers is repaired by a thin post-hoc layer that leaves the model untouched. A forecast that does not order outcomes at all needs the model changed. On this campaign the evidence for the second is weak, resting on very few positive observations, while the evidence for the first is strong and does not depend on that sample at all.

7.4.3 Why: one block explains all of it

The explainability layer localised the cause. Across all twenty-one confident forecasts, the four-week probability and the milestone-miss probability never differ by more than 0.020: the confident tail is the milestone channel and nothing else. Set against the registry, ten of the eighteen campaign milestones completed, eight remained active and only one was still open past a deadline that had fallen inside the campaign.

That registry reading needs one qualification, because it is more forgiving than it looks. The registry holds each milestone's *current* deadline, and a replanning rewrites it, so a milestone re-dated and then delivered counts as completed. Scored instead against the deadline each node originally committed to, six of the eleven milestones whose original deadline fell inside the campaign were missed, five of them by a delivery that arrived one to six turns late. Under that stricter reading five of the twenty-one confident forecasts were followed by a missed original commitment within four weeks, rather than none. The confident tail is still wrong far more often than it is right, and the diagnosis below is unaffected, but the difference between zero of twenty-one and five of twenty-one comes entirely from which deadline the truth is read against. Section 7.4.5 returns to that.

Reading the time block directly at the turn the scenario's critical event lands confirms the diagnosis from the other side: the block reads numerically zero at six of the eight nodes in the middle of a crisis. The cause is structural. The block computes milestone risk from a lead-time distribution against a deadline, while the event engine writes to capacity, cost and failure probability and never to any quantity on the time axis. The strongest form of the evidence is a count: across the whole campaign, eight nodes and forty-one indicator snapshots each, the accumulated-delay field the block would need was written a non-zero value exactly zero times. No path connects a shock to the quantity the block reads, so the block cannot respond to a disruption. It is blind by construction, and because it drives the entire confident tail, so is the forecast.

This is a defect in the model rather than a property of the scenario. Any deployment would inherit it, and the repair is a change to the event mapping rather than to the mathematics. It is the single most actionable finding of the mission.

7.4.4 What happens when the forecast is shown to the operator

The treatment arm produced a second result, and one that was not sought.

For the first four turns the forecast is hidden and all eight declarants record no influence. At the turn it first becomes visible, all eight record one, and the direction is entirely one-sided: five revised their declared urgency downward and three had their existing assessment confirmed. Not one revised upward. The first upward revision anywhere appears the following turn, and the scenario's single critical event lands two turns after the moment every declarant was either reassured or confirmed.

Figure 13 shows the distribution turn by turn. Over the full campaign the balance is close to even, 15 upward against 16 downward, so the one-sided effect belongs to first contact rather than being a permanent bias.

Two things make this worse than a curiosity. The forecast being displayed comes from the same layer whose confident predictions never materialise, and whose time block is blind to




Figure 13 – Influence of the displayed forecast, by turn

the very shock about to arrive. And it is displayed as a numeric probability with a confidence interval, a format the automation-bias literature identifies as increasing reliance (Goddard et al., 2012; Bansal et al., 2021). The system was presenting an unfounded reassurance in its most persuasive available form, immediately before the event it had failed to anticipate.

One qualification cuts against the strongest reading: the recorded field captures the effect on the declared level, not on the action taken, and the free-text notes show declarants still acting protectively. The effect is on the measurement, which is nonetheless the input the whole adequacy layer consumes.

7.4.5 The outcome definition moves the answer

A final measurement concerns how these results should be read at all. The identical one hundred and twenty forecasts were scored against three definitions of an adverse outcome, and the three disagree by more than any modelling choice made during the mission. Table 8 gives them at four weeks.

Table 8 – The same forecasts under three definitions of truth, four-week horizon

Outcome definition	<i>n</i>	Base rate	AUC	Brier skill
Production, current deadline	120	0.050	0.485	−3.36
Naive, any recorded event	75	0.387	0.484	−0.64
Original commitment, censored points excluded	90	0.289	0.654	−0.55

The first is the definition the running system applies, and it is the one the negative result above is stated against. The second counts any recorded event regardless of gravity and ignores milestones entirely; it is the easy target, and choosing it would report a system five times better calibrated without a line of code being touched. The third keeps the production definition of an adverse outcome but reads a missed milestone against the deadline the node originally committed to rather than against whatever deadline the milestone currently carries, and drops the forecasts whose observation window runs past the end of the campaign.

The model did not change between the three rows. Only the definition of truth did, and it moves the Brier skill by a factor of six and the area under the ROC curve from indistinguishable from chance to 0.654. The third row is the one the rest of this section uses, for a reason that is a matter of what a forecast is for rather than of statistics: a replanning is itself the admission of a slip, so a truth that lets a re-dating erase the miss scores the forecast against an outcome the forecast was supposed to warn about. A calibration figure quoted without its outcome definition therefore carries close to no information, which is a methodological result worth more to the programme than the score itself.

7.5 The Repair, and What It Measures

The diagnosis above named one component and specified a repair. That repair was then carried out and re-measured on the identical frozen campaign, which is what this section reports. The declarations are unchanged, the scenario is unchanged, and the same one hundred and twenty forecasts are recomputed by the modified model, so any difference belongs to the model alone.

7.5.1 The recommended repair was the wrong one

The obvious reading of the diagnosis is that the calibrated events should write a lead-time impact alongside their existing effects, and that is what the first version did. Measuring it showed it does not work, for a reason worth stating because it generalises. A completion date estimated as a full production cycle starting now scales the lead time by the fraction of work remaining, so anything routed into the lead time is scaled by that fraction too, and a shock therefore evaporates precisely on the milestones closest to their deadline, which are the ones that matter. On the node carrying the scenario's accident, the four-week probability fell from 0.846 to 0.000 on a milestone that was in fact missed. Lost production time is additive, not proportional: a four-week freeze costs four weeks whether the order is a tenth done or nine tenths done.

The completion model was rewritten accordingly, identically in the three independent estimators the system carries for the same quantity, the analytic time block, the Monte Carlo forecast rollout and the PERT lead-time simulator:

$$\text{completion} = t + (1 - \text{progress}) \times \text{lead time} + \text{lost time}, \quad (2)$$

a proportional cycle term and an additive lost-time term. Five families of capacity-halting event, machine failure, accident, strike weighted by the share of the workforce involved, material shortage weighted by criticality and cyber incident, now accumulate their duration into the node's lost-time indicator. Loss of storage capacity deliberately does not, since it removes space rather than production hours.

The mechanical effect is the one the diagnosis asked for. Where the accumulated-delay field had been written a non-zero value zero times across the whole campaign, the node carrying the accident now accumulates 154.7 engine hours at its first shock and 371.2 over the campaign, carried in twenty-seven of its forty-one snapshots.

7.5.2 Two further defects the first measurement had hidden

Repairing the first defect exposed a second. With the time block finally responsive, the milestone probability turned out to be almost binary: on the frozen replay it returned exactly zero or exactly one in eighty-five percent of cases, leaving fourteen forecasts of one hundred and twenty anywhere in between. That is a switch, not a probability, and two of the declarants had said so in their free-text notes before it was measured. The cause is that the rollout drew the lead time at random but took the declared progress as exact. Declared progress is the least reliable input the system receives, and it is precisely the kind of quantity the whole mission exists to doubt: the tool measures the gap between declared and computed urgency because self-reported figures run optimistic, and then its own forecast was treating one of them as certain. Progress is now drawn per replicate, with a standard deviation widened node by node in proportion to that node’s own measured hidden risk, so the system uses its own adequacy measurement to calibrate its own forecast.

The third defect was in the data rather than the code, and it is the one that would have invalidated any calibration attempted before it was found. Indicator coverage on the campaign database was 25.6%: twenty-four of the forty-two indicators were empty on all eight nodes, the carbon and cost blocks were inert everywhere, and only three of the six urgency blocks contributed anything at all. The lead time itself, the central quantity of the milestone channel, was present on five nodes out of eight, and the other three received a risk floor of zero through the documented fallback for a missing lead time, which is an absence of data presenting itself as a reassuring prediction. Filling the indicators from public figures of the real shortage raised coverage to 49.1% and made all six blocks active. One further bug became visible only after that: a per-node recovery capability, a twenty-six-week aerospace requalification time, was being read as production time already lost, which pinned that node’s time block at 1.0 from the first turn. An empty database hides the errors in the model that consumes it.

7.5.3 What the re-measurement gives

Table 9 scores the recomputed forecasts against the original-commitment definition of Section 7.4.5, at all four horizons, on the same censored samples of 112, 104, 96 and 90 observations.

Table 9 – The frozen campaign rescored after repair

Horizon	Brier skill			AUC		
	before	structure	+ uncertainty	before	structure	+ uncertainty
1 week	−1.278	−0.646	−0.687	0.736	0.793	0.729
2 weeks	−0.690	−0.401	−0.367	0.636	0.671	0.634
3 weeks	−0.654	−0.360	−0.286	0.641	0.657	0.602
4 weeks	−0.553	−0.272	−0.216	0.654	0.721	0.584

The structural repair improves both axes at every horizon, which is the clean result: the calibration deficit is roughly halved and the discrimination rises, and neither had to be traded for the other. The progress-uncertainty term is a trade rather than a gain. It improves calibration further at two weeks and beyond, worsens it slightly at one week, and lowers the area under the ROC curve; that last movement is not interpretable, since with twenty-six positive observations the ninety-five percent interval on the four-week figure runs from 0.21 to 0.87 and contains every value in its row. The term was therefore kept on the calibration evidence and on the argument that a declared quantity should not enter a probabilistic forecast as a certainty, not on the discrimination figure.

Three things should be said plainly against this. The Brier skill is still negative at every horizon, so the forecast is still worse calibrated than a forecaster that announces the base rate and stops. The milestone probability is still binary in seventy-two percent of cases, down from eighty-five, so the underlying over-confidence is reduced rather than removed. And a post-hoc calibration layer was trained, measured, and deliberately not connected: on the synthetic factory it improves the Brier skill by thirty-one percent, but on the campaign it improves calibration while dropping the four-week area under the ROC curve from 0.72 to 0.58, because it was trained on a target defined as an event striking the node at exactly one week while the product displays a missed milestone or an event cumulated over four. That is the trap of Section 7.4.5 met a second time from the other side: a layer calibrated against the wrong target degrades a system that works. Connecting it requires the training label to be rebuilt on the product's target, which requires the synthetic factory's snapshots to carry milestone truth, which they do not.

7.5.4 A third arm, and the honest reading of it

The repaired system was then run through a full nineteen-turn campaign in closed loop, in which the same eight declarants no longer only declared but also chose actions from a fixed catalogue, applied to the chain by the system and journalled: promote a backup arc, replan a milestone, expedite a shipment, raise capacity, revise a declaration, or do nothing. One hundred and fifty-two declarations were recorded with no unresolved consistency failure.

The headline figure flatters the tool and should not be used. Six of eleven milestones were missed in the frozen scenario and none was left unfinished in the closed loop, which sounds like six saved. Read against the original commitments it is not: five of eleven milestones were met on their original date in the frozen run, and five of eleven in the closed loop. The loop changed which milestone was met, not how many. One was genuinely saved, an electronics series delivered on its original date that had slipped in the frozen truth. Four late deliveries became replannings announced in advance instead of slips discovered at the deadline, which is worth something to a customer without being a save. And one milestone was replanned that would have been met, which is the measurable cost of over-reaction. With eleven milestones neither the benefit nor the cost is established; they are facts of one campaign, not results.

One further repair belongs here, because it touches the other half of the network layer. The criticality ranking, which sweeps the network by knocking out each node in turn and measuring the damage at the final customer, collapses when many nodes saturate at once: every delta clips to zero and the ranking says nothing, in exactly the crisis where it would be worth reading. Nine of the nineteen campaign turns were in that state. The ranking now carries a second sort key computed in log-survival, $\ell(p) = -\ln(1 - p)$, on a regularised

twin of the same propagation. Because ℓ is strictly increasing, the two keys order identically outside saturation, so nothing changes where the index already worked, and inside saturation the clip that flattens the first does not flatten the second. Measured on the same run, the number of turns on which the deep suppliers can be ordered at all rises from 10 to 17 of 19, and the expected bottleneck leads on 16 of those 17.

Two behavioural observations from that run do carry. The turn-four disarming of the first campaign reproduced exactly, eight declarants of eight. And the repair is visible in a declarant's own words at the turn the crisis event lands: where the first campaign's foundry manager wrote that the tool was holding a low position through a widely reported crisis and took no action, the repaired run moved the same node's four-week probability from 0.088 to 0.214 and the declarant replanned the milestone. The chain from a scripted shock to a displayed number to an action taken now closes, which it demonstrably did not before.

7.6 Problems Encountered and How They Were Resolved

Four difficulties shaped the mission, beyond the risks of Section 5.5.

The first model could not represent the problem. The initial formulation expressed urgency as a single time-driven curve. It could not represent a node on schedule but out of capacity, which is a common and dangerous state. The resolution was to abandon it in April in favour of multi-block aggregation. The cost was roughly three weeks; the alternative was building a validated implementation of an inadequate model.

Measuring the gap requires two independent signals, which is harder than it looks. The same difficulty recurred three times in different disguises. A first synthetic dry run derived the declared value from the computed one, so no gap could appear. The blind arm used declarants who read their briefings so carefully that their declaration reproduced the computed state. The treatment arm closed the loop deliberately by showing the model's output to the declarant. Each time the symptom looked like a modelling problem and each time the cause was circularity in the measurement. The resolution was to state the independence requirement as a first-class design constraint and to record, for every result, which of the two signals could have contaminated the other.

The literature was not accessible from inside the company. As set out in Section 5.5, ALTEN's Direction de l'Innovation holds no journal or database subscriptions. The state of the art was built entirely on ESTIA and Cranfield library access. The practical resolution was to front-load literature retrieval into the first weeks while planning which sources would be needed later, since retrieving a paper was never a same-hour operation.

Compute was limited to a CPU. No GPU was available. This bounded what could be attempted and is the reason one comparison was run at reduced scale. The resolution was methodological rather than material: choosing approaches whose cost is dominated by the quantity of data rather than by model size, and accepting the reduced scale explicitly in the reporting instead of quietly omitting the comparison.

8 Conclusion and Personal Assessment

8.1 Results Against the Objectives

Table 10 assesses each objective of Section 4.3 against its indicator. Four of five were met. The fifth was answered in two halves: the signal does identify the right nodes early, and the calibrated probability built on it is measurably better after repair without yet being good enough to publish as a probability.

Table 10 – Objectives against outcomes

Status		Evidence
O1	Met	Six-block aggregation with bidirectional propagation and an asymmetric adequacy score, each term exactly decomposable; aggregation choice justified against the compensatory alternative
O2	Met	SupplyScore delivered: 73 modules, 1831 automated test functions, multi-tenant persistence, every write audited, operational and reloadable
O3	Met	Propagation and criticality services operational, validated by dedicated performance benchmarks at the scale of a real programme; the criticality ranking's collapse under network-wide saturation was identified on the campaign and repaired (Section 7.5)
O4	Met at node level, in progress at probability level	At node level the signal is right: over the six turns before the crisis is public it flags three nodes, all three of which later miss a commitment, with no false alarm (Section 7.4.1). As a calibrated four-week probability it first failed, for an identified cause; the repair was carried out and re-measured on the frozen campaign (Section 7.5), roughly halving the calibration deficit at every horizon while leaving it negative
O5	Met	Scientific report delivered; every reported figure regenerated from raw artefacts by script rather than transcribed

Quantitatively. The delivered system covers the specified scope and is measurably robust: 1831 automated test functions expanding to 2 044 collected cases, an auditable write path, and performance validated at the intended scale. The experiment produced numbers rather than impressions: an area under the ROC curve between 0.37 and 0.56, a Brier skill between -3.4 and -6.9 , twenty-one confident forecasts of which none materialised against the production definition of an adverse outcome, and an eight-of-eight one-sided response at first exposure to the displayed forecast.

The repair that followed is measured on the same frozen forecasts. The Brier skill improves from -1.278 to -0.687 at one week and from -0.553 to -0.216 at four, roughly halving the calibration deficit at every horizon, while the structural half of the repair also raises the area under the ROC curve at every horizon. Three figures qualify that. The skill is still negative, so the forecast remains worse calibrated than announcing the base rate. The milestone probability is still binary in seventy-two percent of cases, against eighty-five before. And the closed-loop campaign run on the repaired system met five of eleven milestones on their original date, exactly as many as the frozen run: it changed which milestone was met, not how many.

Qualitatively. The mission set out to determine whether the declared-versus-real gap can be measured well enough to act on. On the evidence gathered it can, at the level a coordinator actually works at, which is deciding where to look. The gap flags the right suppliers, it flags them before the crisis is visible to anyone, and it does so on nodes whose own managers are declaring that all is well. What is not yet established is the next step up, turning that flag into a calibrated probability attached to a named week.

The work is worth more to the programme for having failed that second step first. The failure was actionable: it named one component, stated the mechanism, and specified the repair, and the repair was then carried out and measured. A favourable score obtained on the first pass, on a signal driven by a block that could not see the crisis, would have been worse than useless, because it would have been trusted. The route from here is not a new idea but more data, and Section 8.3 says which data and why.

8.2 Limitations

Five limits bound what this report supports, and four of the five are limits of sample rather than of design. That distinction matters for what happens next: none of them calls the construction into question, and all of them are answered by running the instrument on more chains rather than by changing it.

The declarants were agents rather than human consultants, so every statement about declaration behaviour applies to that population. Results come from a single scenario, so the outcome base rate and event mix are properties of one script. The positive class is very small, eight nodes of which four are positive and eleven observable milestones, which bounds every claim in both directions: it is why the precision of 1.00 in Section 7.4.1 is reported as an encouraging observation rather than a rate, and why the movements in the area under the ROC curve in Section 7.5 are stated as uninterpretable rather than as gains. And the carbon block was the least exercised of the six, for want of data to feed it, which is the measurement gap noted in Section 2.3; the indicator completion of Section 7.5 narrowed it without closing it.

The fifth is methodological and does not dissolve with more data. Every score depends on which of the three outcome definitions of Section 7.4.5 it is read against, and no number in this report should be quoted without one.

8.3 Perspectives

For ALTEN. The binding constraint from here is data, not design. The measurement works at node level on the chain it was tested on; what stands between that and a deployable early-warning instrument is the number of chains, nodes and milestones it has been run against, and the steps below are ordered accordingly.

The first of the four steps this report recommended has been taken, and taking it changed the recommendation. Repairing the time block by making the calibrated events write a lead-time impact, as originally proposed, does not work: routing lost time into the lead time makes it scale with the work remaining, so the shock vanishes on exactly the milestones nearest their deadline. Lost time has to be added to the completion date, not multiplied into it (Section 7.5). The step generalised into a rule worth keeping beyond this system: distinguish a parameter

describing a node's recovery capability from a state accumulating what it has actually lost, because the two live naturally in the same units and mixing them is invisible until the data is complete.

Four steps remain, in order. Fill the indicators before calibrating anything, since the campaign that produced the negative result ran on a quarter of them and half the urgency blocks were inert; no amount of calibration compensates for a missing input, and one real bug was undetectable until the database was full. Rebuild the post-hoc calibration layer's training label on the target the product actually displays, which first requires the synthetic factory's snapshots to carry milestone truth; the layer as trained today improves calibration and degrades ranking, and is deliberately not connected. Make the forecast state its own ignorance, since a layer that cannot anticipate exogenous shocks should not present a low probability in a confident numeric format, and smooth the milestone alarm, which today moves from 0.05 to 1.00 in a single turn and is correct but useless as advance warning. Then run the campaign with human participants, which is the only run that can decide the hypotheses as they were written.

One indicator class is missing rather than empty, and several declarants named it independently: the forward booking horizon, how far ahead a node's capacity is already sold. In the real shortage that signal preceded the degradation of every flow indicator by months, and it is structurally invisible to the model as built. It is a general gap, not a semiconductor one: any capacity-constrained link holds that information.

Beyond that, the adequacy layer is designed to be coupled to DISCO rather than to remain standalone, and the network representation work carried out in parallel on the same programme is the natural counterpart to connect it to.

For the author. The subject continues in the research thesis produced for Cranfield University, which treats the same measurements at greater depth. The open question I would most like to pursue is the one the treatment arm opened: what an advisory layer owes the person reading it when it does not know.

8.4 Personal Assessment

8.4.1 What I learned

The hardest problem of these six months was not mathematical. It was recognising, three separate times and in three different disguises, that the system was measuring itself. Each time the symptom looked like a modelling defect and each time the cause was circularity in the measurement. I would now design the independence of two signals before designing either of them.

The second lesson concerns what counts as a result. For a long stretch I treated a negative outcome as a failure to be fixed before reporting it. The most useful findings in this report are negative, and none would have been found by a project optimising for a positive one. The related discipline, which the rejected dataset taught quickly, is that a confident presentation is not evidence, and that the cheap verification step belongs before the expensive one.

With hindsight I would change two things. I would have scored the forecast layer against recorded outcomes far earlier, since doing so took a few hundred lines and revealed the

defect immediately where months of implementation had not. And I would have written the outcome definition into the pre-registration alongside the thresholds, having since measured how much it moves the answer.

8.4.2 Link with the ESTIA engineering competency framework

Table 11 maps the mission onto the competencies of the ESTIA framework that it genuinely exercised. Appendix C gives the detailed evidence.

Table 11 – Competencies exercised during the mission

	Competency	Evidence from this mission
CI1	Analysis, synthesis, creativity, systemic and interdisciplinary approach	Combining operations research, behavioural science, probabilistic modelling and software architecture into one measurement
CI3	Working in an international and intercultural context	Mission conducted and reported in English, under a double-degree programme, alongside an English-speaking colleague on the same research programme
CI5	Identifying and solving complex or research problems using external resources	The whole subject; and literally external, since the state of the art depended on university library access unavailable within the company
CE3	Leading and controlling a project against objectives, deadlines and risks	A four-touchpoint weekly cadence, six phases, a documented mid-mission replanning, and a risk register of which three items materialised
CE4	Taking the socio-ecological transition into account	Carbon modelled as one of six first-class blocks rather than as a correction, within a programme that is itself part of ALTEN's environmental R&D commitment
CE5	Researching, selecting, qualifying and diffusing information	A candidate dataset rejected on inspection before use; the same work written three ways for three different audiences
CST1	Defining an engineering strategy in a poorly known context	Converging a five-workstream assignment onto one defensible contribution, and abandoning the first model when it proved unable to represent the problem
CST2	Modelling and simulating a system, capitalising the knowledge produced	The six-block model, its propagation, the campaign simulation, and a reproducible derivation of every reported figure
CST7	Building models, methods and tools to aid engineering, in an R&D perspective	The delivered application and the experimental protocol, both designed to be continued by the programme rather than to end with the placement

8.4.3 What this mission clarified professionally

I came into this placement expecting the interesting difficulty to be in the mathematics. It was not. The mathematics behaved; what failed was the join between a model and the world it was supposed to describe, and then the join between a model's output and the person meant to act on it. Both failures were only visible because the work was measured against something external rather than judged on its own internal consistency.

That is the kind of engineering I want to continue: the boundary where a model meets the person who has to act on it. It is where the failures are least obvious and most consequential, and this mission was a fairly direct demonstration of why.

Bibliography

- BANSAL, G. et al., 2021. Does the Whole Exceed Its Parts? The Effect of AI Explanations on Complementary Team Performance. In: *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. Disp. à l'adr. DOI: [10.1145/3411764.3445717](https://doi.org/10.1145/3411764.3445717).
- BAULAIN, D.; DUBARRY, T., 2024. *Internet Physique, Optimisation du Routage et Indice de Performance Globale dans DISCO*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- BAXAS, A., 2025. *Architecture du Jumeau Numérique DISCO et Modélisation de la Résilience*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- BECKER, W. et al., 2017. Weights and Importance in Composite Indicators: Closing the Gap. *Ecological Indicators*. Vol. 80, p. 12–22. Disp. à l'adr. DOI: [10.1016/j.ecolind.2017.03.056](https://doi.org/10.1016/j.ecolind.2017.03.056).
- BEHZADIAN, M. et al., 2010. PROMETHEE: A Comprehensive Literature Review on Methodologies and Applications. *European Journal of Operational Research*. Vol. 200, no. 1, p. 198–215. Disp. à l'adr. DOI: [10.1016/j.ejor.2009.01.021](https://doi.org/10.1016/j.ejor.2009.01.021).
- BORGONOVO, E.; CILLO, A.; PLISCHKE, E., 2021. Global sensitivity analysis of multi-criteria decision analysis: A review of concepts, methods, and applications. *European Journal of Operational Research*. Vol. 290, no. 3, p. 805–827. Disp. à l'adr. DOI: [10.1016/j.ejor.2020.07.034](https://doi.org/10.1016/j.ejor.2020.07.034).
- BRANS, J.-P.; VINCKE, P.; MARESCHAL, B., 1986. How to Select and How to Rank Projects: The PROMETHEE Method. *European Journal of Operational Research*. Vol. 24, no. 2, p. 228–238. Disp. à l'adr. DOI: [10.1016/0377-2217\(86\)90044-5](https://doi.org/10.1016/0377-2217(86)90044-5).
- BRIER, G. W., 1950. Verification of Forecasts Expressed in Terms of Probability. *Monthly Weather Review*. Vol. 78, no. 1, p. 1–3. Disp. à l'adr. DOI: [10.1175/1520-0493\(1950\)078<0001:VOFEIT>2.0.CO;2](https://doi.org/10.1175/1520-0493(1950)078<0001:VOFEIT>2.0.CO;2).
- CHEN, J. et al., 2025. Theoretical Foundations of Utility Accrual for Real-Time Systems.
- CHEN, K.; MÜHLETHALER, P., 1996. A scheduling algorithm for tasks described by Time Value Function. *Real-Time Systems*.
- COHEN, L.; GRINSHPOUN, T.; STERN, R., 2019. Assigning Suppliers to Meet a Deadline. In: *Proceedings of the Twelfth International Symposium on Combinatorial Search (SoCS 2019)*. AAAI Press. Vol. 10, p. 170–171. No. 1. Disp. à l'adr. DOI: [10.1609/socs.v10i1.18475](https://doi.org/10.1609/socs.v10i1.18475).
- CORDOVA, A., 2024. *Green Supply Chain et Quantification des Émissions CO₂ dans DISCO*. Paris. Stage Report. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- DING, Y. et al., 2023. Real-Time Monitoring and Optimal Resource Allocation for Automated Container Terminals: A Digital Twin Application at the Yangshan Port. *Journal of Advanced Transportation*.
- DIWASSINGH, K. et al., 2020. Task Selection and Workload: A Focus on Completing Easy Tasks Hurts Long-Run Performance. *Management Science*. Vol. 66, no. 10, p. 4397–4416. Disp. à l'adr. DOI: [10.1287/mnsc.2019.3366](https://doi.org/10.1287/mnsc.2019.3366).
- DYER, J. S., 1990. Remarks on the Analytic Hierarchy Process. *Management Science*. Vol. 36, no. 3, p. 249–258.
- FULLER, A. et al., 2020. Digital Twin: Enabling Technologies, Challenges and Open Research. *IEEE Access*. Vol. 8, p. 108952–108971. Disp. à l'adr. DOI: [10.1109/ACCESS.2020.2998358](https://doi.org/10.1109/ACCESS.2020.2998358).
- GODDARD, K.; ROUDSARI, A.; WYATT, J. C., 2012. Automation Bias: A Systematic Review of Frequency, Effect Mediators, and Mitigators. *Journal of the American Medical Informatics Association*. Vol. 19, no. 1, p. 121–127. Disp. à l'adr. DOI: [10.1136/amiajnl-2011-000089](https://doi.org/10.1136/amiajnl-2011-000089).

- GRECO, S. et al., 2018. On the Methodological Framework of Composite Indices: A Review of the Issues of Weighting, Aggregation, and Robustness. *Social Indicators Research*. Vol. 141, p. 61–94. Disp. à l'adr. DOI: [10.1007/s11205-017-1832-9](https://doi.org/10.1007/s11205-017-1832-9).
- GRIEVES, M., 2014. *Digital Twin: Manufacturing Excellence Through Virtual Factory Replication*. White Paper. Florida Institute of Technology.
- GUO, C. et al., 2017. On Calibration of Modern Neural Networks. In: *Proceedings of the 34th International Conference on Machine Learning (ICML)*.
- GUO, S.; ZHAO, H., 2017. Fuzzy Best-Worst Multi-Criteria Decision-Making Method and Its Applications. *Knowledge-Based Systems*. Vol. 121, p. 23–31. Disp. à l'adr. DOI: [10.1016/j.knosys.2017.01.010](https://doi.org/10.1016/j.knosys.2017.01.010).
- HECKMANN, I.; COMES, T.; NICKEL, S., 2015. A Critical Review on Supply Chain Risk — Definition, Measure and Modeling. *Omega*. Vol. 52, p. 119–132. Disp. à l'adr. DOI: [10.1016/j.omega.2014.10.004](https://doi.org/10.1016/j.omega.2014.10.004).
- HLIOUI, R.; GHARBI, A.; HAJJI, A., 2017. Joint Supplier Selection, Production and Replenishment of an Unreliable Manufacturing-Oriented Supply Chain. *International Journal of Production Economics*. Vol. 187, p. 53–67. Disp. à l'adr. DOI: [10.1016/j.ijpe.2017.02.004](https://doi.org/10.1016/j.ijpe.2017.02.004).
- HOEN, K. M. R. et al., 2014. Effect of Carbon Emission Regulations on Transport Mode Selection under Stochastic Demand. *Flexible Services and Manufacturing Journal*. Vol. 26, no. 1–2, p. 170–195. Disp. à l'adr. DOI: [10.1007/s10696-012-9151-6](https://doi.org/10.1007/s10696-012-9151-6).
- HOSSEINI, S.; IVANOV, D.; DOLGUI, A., 2020. Ripple Effect Modelling of Supplier Disruption: Integrated Markov Chain and Dynamic Bayesian Network Approach. *International Journal of Production Research*. Vol. 58, no. 11, p. 3284–3303. Disp. à l'adr. DOI: [10.1080/00207543.2019.1661538](https://doi.org/10.1080/00207543.2019.1661538).
- HOSSEINI, S.; TAJIK, N., et al., 2019. Resilient Supplier Selection and Optimal Order Allocation under Disruption Risks. *International Journal of Production Economics*. Vol. 213, p. 124–137. Disp. à l'adr. DOI: [10.1016/j.ijpe.2019.03.018](https://doi.org/10.1016/j.ijpe.2019.03.018).
- HRUŠOVSKÝ, M. et al., 2021. Real-time disruption management approach for intermodal freight transportation. *Journal of Cleaner Production*.
- INTERNATIONAL ORGANIZATION FOR STANDARDIZATION, 2023. *ISO 14083:2023 Greenhouse gases — Quantification and Reporting of Greenhouse Gas Emissions Arising from Transport Chain Operations*. Geneva.
- IVANOV, D.; DOLGUI, A., 2019. New Disruption Risk Management Perspectives in Supply Chains: Digital Twins, the Physical Internet, and Smart Contracts. *IFAC-PapersOnLine*. Vol. 52, no. 13, p. 337–342. Disp. à l'adr. DOI: [10.1016/j.ifacol.2019.11.138](https://doi.org/10.1016/j.ifacol.2019.11.138).
- IVANOV, D.; DOLGUI, A., 2020. A Digital Supply Chain Twin for Managing the Disruption Risks and Resilience in the Era of Industry 4.0. *Production Planning & Control*. Vol. 32, no. 9, p. 775–788. Disp. à l'adr. DOI: [10.1080/09537287.2020.1768450](https://doi.org/10.1080/09537287.2020.1768450).
- IVANOV, D.; DOLGUI, A., et al., 2021. Digital Supply Chain Twins: Managing the Ripple Effect, Resilience, and Disruption Risks by Data-Driven Optimization, Simulation, and Visibility. *International Journal of Production Research*. Vol. 59, no. 3, p. 957–974. Disp. à l'adr. DOI: [10.1080/00207543.2020.1841183](https://doi.org/10.1080/00207543.2020.1841183).
- KABOUCHE, I., 2025. *Scoring de Risque et Réseaux Bayésiens pour la Supply Chain Prédictive*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- KIM, S.; TONG, J.; PEDEN, C., 2020. Admission Control Biases in Hospital Unit Capacity Management: How Occupancy Information Hurdles and Decision Noise Impact Utilization. *Management Science*. Vol. 66, no. 5, p. 1803–1824. Disp. à l'adr. DOI: [10.1287/mnsc.2019.3303](https://doi.org/10.1287/mnsc.2019.3303).

- LEUNG, E.; LEE, C.; OUYANG, Z., 2022. From traditional warehouses to Physical Internet hubs: A digital twin-based inbound synchronization framework for PI-order management. *International Journal of Production Economics*.
- LI, Y.; ZOBEL, C. W., 2020. Exploring Supply Chain Network Resilience in the Presence of the Ripple Effect. *International Journal of Production Economics*. Vol. 228, p. 107664. Disp. à l'adr. DOI: [10.1016/j.ijpe.2020.107664](https://doi.org/10.1016/j.ijpe.2020.107664).
- LIANG, F.; BRUNELLI, M.; REZAEI, J., 2020. Consistency Issues in the Best Worst Method: Measurements and Thresholds. *Omega*. Vol. 96, p. 102175. Disp. à l'adr. DOI: [10.1016/j.omega.2019.102175](https://doi.org/10.1016/j.omega.2019.102175).
- LIU, M. et al., 2021. A New Robust Dynamic Bayesian Network Approach for Disruption Risk Assessment under the Supply Chain Ripple Effect. *International Journal of Production Research*. Vol. 59, no. 1, p. 265–285. Disp. à l'adr. DOI: [10.1080/00207543.2020.1841318](https://doi.org/10.1080/00207543.2020.1841318).
- MACKWAY-JONES, K.; MARSDEN, J.; WINDLE, J., 2013. *Emergency Triage: Manchester Triage Group*. 3rd ed. Oxford: Wiley-Blackwell. ISBN 978-1118299029.
- MOKHTAR, S. et al., 2019. Supplier Portfolio Selection Based on the Monitoring of Supply Risk Indicators. *Simulation Modelling Practice and Theory*. Vol. 97, p. 101955. Disp. à l'adr. DOI: [10.1016/j.simpat.2019.101955](https://doi.org/10.1016/j.simpat.2019.101955).
- MONDELICE, R., 2025. *BST TMS — Transport Management System et Interface Routage DISCO*. Paris. Bilan Scientifique et Technique. ALTEN, Direction Innovation Numérique. [Internal, restricted access].
- MONTREUIL, B., 2011. Toward a Physical Internet: Meeting the Global Logistics Sustainability Grand Challenge. *Logistics Research*. Vol. 3, no. 2–3, p. 71–87. Disp. à l'adr. DOI: [10.1007/s12159-011-0045-x](https://doi.org/10.1007/s12159-011-0045-x).
- MURPHY, A. H., 1973. A New Vector Partition of the Probability Score. *Journal of Applied Meteorology*. Vol. 12, no. 4, p. 595–600. Disp. à l'adr. DOI: [10.1175/1520-0450\(1973\)012<0595:ANVPOT>2.0.CO;2](https://doi.org/10.1175/1520-0450(1973)012<0595:ANVPOT>2.0.CO;2).
- NICULESCU-MIZIL, A.; CARUANA, R., 2005. Predicting Good Probabilities with Supervised Learning. In: *Proceedings of the 22nd International Conference on Machine Learning (ICML)*. Disp. à l'adr. DOI: [10.1145/1102351.1102430](https://doi.org/10.1145/1102351.1102430).
- NILOOFAR, P.; LAZAROVA-MOLNAR, S.; SHAKER, H. R., 2023. A General Framework for Human-in-the-Loop Cognitive Digital Twins. In: *Proceedings of the 2023 Winter Simulation Conference (WSC)*. IEEE, p. 4155–4166. Disp. à l'adr. DOI: [10.1109/WSC57122.2023.10407598](https://doi.org/10.1109/WSC57122.2023.10407598).
- REZAEI, J., 2015. Best-Worst Multi-Criteria Decision-Making Method. *Omega*. Vol. 53, p. 49–57. Disp. à l'adr. DOI: [10.1016/j.omega.2014.11.009](https://doi.org/10.1016/j.omega.2014.11.009).
- RUSOU, Z.; AMAR, M.; AYAL, S., 2020. The Psychology of Task Management: The Smaller Tasks Trap. *Judgment and Decision Making*. Vol. 15, no. 4, p. 586–599.
- SAATY, T. L., 1980. *The Analytic Hierarchy Process: Planning, Priority Setting, Resource Allocation*. New York: McGraw-Hill. ISBN 978-0070543713.
- SHEN, X. et al., 2019. Managing Supply Chains with Expediting and Multiple Demand Classes. *Production and Operations Management*. Vol. 28, no. 5, p. 1129–1148. Disp. à l'adr. DOI: [10.1111/poms.12974](https://doi.org/10.1111/poms.12974).
- SMART FREIGHT CENTRE, 2023. *GLEC Framework for Logistics Emissions Accounting, Version 3.2* [en ligne]. [visité le 2026-04-24]. Technical Report. Smart Freight Centre. Disp. à l'adr.: <https://www.smartfreightcentre.org/en/our-programs/emissions-accounting/global-logistics-emissions-council/calculate-report-glec-framework/>. Aligned with ISO 14083:2023.
- STEEL, P.; KÖNIG, C. J., 2006. Integrating theories of motivation. *Academy of Management Review*. Vol. 31, no. 4, p. 889–913. Disp. à l'adr. DOI: [10.5465/AMR.2006.22527462](https://doi.org/10.5465/AMR.2006.22527462).

- SUN, X. et al., 2020. Combating Lead-Time Uncertainty in Global Supply Chain's Shipment-Assignment: Is It Wise to Be Risk-Averse? *Transportation Research Part B: Methodological*. Vol. 138, p. 406–434. Disp. à l'adr. DOI: [10.1016/j.trb.2020.05.015](https://doi.org/10.1016/j.trb.2020.05.015).
- TAO, F. et al., 2018. Digital Twin-Driven Product Design, Manufacturing and Service with Big Data. *The International Journal of Advanced Manufacturing Technology*. Vol. 94, no. 9–12, p. 3563–3576. Disp. à l'adr. DOI: [10.1007/s00170-017-0233-1](https://doi.org/10.1007/s00170-017-0233-1).
- TOMLIN, B., 2006. On the Value of Mitigation and Contingency Strategies for Managing Supply Chain Disruption Risks. *Management Science*. Vol. 52, no. 5, p. 639–657. Disp. à l'adr. DOI: [10.1287/mnsc.1060.0515](https://doi.org/10.1287/mnsc.1060.0515).
- TRIANAPHYLLOU, E.; YANASE, J., 2024. The use of pairwise comparisons for decision making may lead to grossly inaccurate results.
- ZHU, M.; YANG, Y.; HSEE, C. K., 2018. The Mere Urgency Effect. *Journal of Consumer Research*. Vol. 45, no. 3, p. 673–690. Disp. à l'adr. DOI: [10.1093/jcr/ucy008](https://doi.org/10.1093/jcr/ucy008).

Acronyms

DIN Direction Innovation Numérique. 3

GLEC Global Logistics Emissions Council. 3

A Autorisation d'usage de l'IA par l'élève

À signer par le tuteur entreprise avant dépôt du rapport. Le formulaire ci-dessous reprend l'annexe 7 des consignes d'évaluation. Il doit être complété, coché et signé, puis inséré ici sous forme scannée.

Formulaire

Nom de l'apprenant(e)	John HOARAU
Nom du tuteur	Sébastien PERTHUISOT
Entreprise	ALTEN Sud-Ouest, Direction de l'Innovation
Période concernée	du 13 avril 2026 au 16 octobre 2026

Rappel des risques pour l'entreprise en cas d'usage d'outils d'IA générative

L'utilisation d'outils d'IA par l'apprenti(e), sans encadrement ni précautions, peut exposer l'entreprise à plusieurs risques :

- divulgation involontaire de données confidentielles ou stratégiques lors de la saisie dans des outils en ligne ;
- perte de maîtrise sur la propriété intellectuelle des contenus générés (textes, codes, visuels) pouvant être réutilisés par la plateforme ;
- dépendance à des outils tiers dont les résultats peuvent être imprécis, biaisés ou juridiquement discutables ;
- non-conformité au RGPD si des données à caractère personnel sont traitées par des services non conformes.

Autorisation du tuteur

- ☐ J'autorise l'élève à utiliser des outils d'intelligence artificielle générative, en pleine connaissance des risques mentionnés ci-dessus.
- ☐ Je n'autorise pas l'usage d'outils d'IA générative pendant la période en entreprise.

Commentaires ou restrictions spécifiques éventuelles :

Fait à _____ le _____

Signature du tuteur / maître de stage :

Declared use of AI tools in this mission

For transparency, and independently of the authorisation above, the use made of generative AI during this placement is declared here in full.

Large language models were used in two distinct capacities, which should not be confused.

First, as an *object of study*. The eight declarants of the experimental campaign reported in Section 7.3 are language-model agents. This is an experimental design choice, stated as such throughout the report, and its consequences for the validity of the results are discussed rather than glossed over.

Second, as a *writing aid*, for language correction and document formatting. Every scientific claim, every numerical result and every interpretation in this report is my own, and I am able to justify and explain each of them. Every figure was produced either by the project's own code or by hand. No confidential ALTEN data was submitted to any external service in the course of that assistance.

B Detailed Planning

This appendix expands the schedule of Section 5.4 into the weekly rhythm that produced it, and records the two planning revisions made during the mission.

Weekly cadence

Table 12 lists the standing meetings that structured every week of the mission.

Table 12 – Standing weekly meetings

Day	Meeting	With	Purpose
Monday	Planning, Green Supply Chain activity	É. Lardy	Set the objective for the week
Tuesday	Scrum	É. Lardy	Progress, blockages
Wednesday	Scrum	É. Lardy	Progress, blockages
Wednesday	Supply chain regular point	S. Perthuisot	Scientific direction and arbitra- tion
Thursday	Scrum	É. Lardy	Progress, blockages
Friday	Demonstration	É. Lardy, S. Perthuisot	Show working output, not de- scribe it

Six standing contact points a week is a demanding cadence for a research placement. Its effect was to bound how long any misconception could survive: the Friday demonstration in particular required something to be shown rather than reported, which repeatedly exposed gaps between what was believed to work and what did.

Planning revisions

June: validation over extension. The original plan allocated the second half of the mission to widening the model's scope. It was reallocated to validating what already existed. The argument accepted at the regular point was that extending an unvalidated model compounds risk rather than reducing it. This revision is why the mission produced a measured result rather than a larger unmeasured system.

July: substitution of the campaign population. The campaign was designed for eight human consultants across nine working days. When that could not be secured within the mission window, the protocol was executed with isolated agents instead. The record format was deliberately identical, so the substitution changed the population without changing the instrument. Every behavioural conclusion was weakened accordingly in the reporting rather than being stated at the strength the original design would have supported.

Comparison of planned against actual

Table 13 compares the time allowed for each phase with the time it actually took.

Table 13 – Planned against actual, by phase

Phase	Planned	Actual	Comment
Framing	3 weeks	3 weeks	On plan
First pipeline	6 weeks	5 weeks	Stopped early once judged inadequate
Redesign and build	4 weeks	3 weeks	Interfaces already established by the abandoned pipeline
Experimental design	2 weeks	3 weeks	Freezing the protocol properly took longer than allowed
Campaign and analysis	3 weeks	3 weeks	On plan
Capitalisation	4 weeks	5 weeks	Extended after the diagnostic result emerged

The overall schedule held. The one systematic underestimate was the effort required to freeze an experimental protocol to a standard where the hypotheses could not later be adjusted, which took half again the time allowed for it and was worth every day.

C Detailed Competency Evidence

This appendix gives the evidence behind Table 11, competency by competency, so that each claim can be checked against something in the record rather than taken on assertion.

CI1 – Analysis, synthesis, curiosity, creativity and innovation, approaching a problem systematically and across disciplines. The measurement built during this mission required combining four fields that do not normally cite one another: deadline-driven value from operations research, the mere-urgency effect from behavioural science, probabilistic aggregation, and the software architecture needed to make the result auditable. The contribution claimed is not any one of them but the combination, which Section 6.7 shows to be absent from the reviewed literature.

CI3 – Working in an international and intercultural context. The mission was conducted, documented and defended in English, within a double-degree programme with a British university. A colleague on the same research programme worked in English on the adjacent representation problem, and the weekly demonstration was the point at which the two strands had to be made intelligible to each other.

CI5 – Identifying and solving complex or research problems, mobilising scientific and technical knowledge and external resources. The subject was a research problem with no established answer. The word *external* also applies literally: as recorded in Section 5.5, the company holds no journal subscriptions, so every source in the state of the art was retrieved through university credentials.

CE3 – Leading, steering and controlling a project to reach its objectives, in cost, deadline, performance and risk. Six phases with dated milestones; a weekly cadence of six standing meetings; a documented mid-mission replanning that traded scope for validation; and a risk register of four items of which three materialised and were handled. Cost is the one dimension not exercised, since no budget was managed.

CE4 – Taking the socio-ecological transition into account. Carbon is one of the six blocks of the model, ranked alongside cost and delay rather than applied as a correction, and the programme itself falls within the share of ALTEN's R&D devoted to environmental innovation (Section 2.3). The honest qualification, stated in Section 8.2, is that this block turned out to be the least exercised of the six for want of data.

CE5 – Researching, selecting, qualifying and diffusing information, adapting it to its context. The qualifying step is the one worth evidencing: a candidate dataset was inspected before use and rejected on three measurable signatures of synthetic data, and the node it would have fed was driven by calibrated events instead. On diffusion, the same work was written three ways for three audiences: a scientific report for the company, a research thesis for the partner university, and this mission report.

CST1 – Defining an engineering strategy by choosing methods and tools suited to a known or poorly known context. The assignment as signed covered five workstreams. Converging it onto one defensible contribution, and recording the boundaries explicitly rather than leaving them implicit, was the first engineering decision of the mission. The second was abandoning the initial urgency formulation once it proved unable to represent a node on schedule but out of capacity.

CST2 – Modelling and simulating a system, and capitalising the information and knowledge produced. The six-block model with its bidirectional propagation, the campaign that simulates a documented historical crisis over nineteen turns, and a derivation pipeline that regenerates every reported figure from raw artefacts so that no number in any of the three deliverables is transcribed by hand.

CST7 – Designing and building models, methods and tools to aid engineering and improve business processes, including in a research and development perspective. The delivered

application and the experimental protocol were both built to be continued rather than to end with the placement: a documented architecture, an automated test suite, a frozen and hashed protocol, and a written record of what the next steps are and in what order.